



# **Introduction to 3D Vision**

## **: A Tutorial for Everyone**

**Sunglok Choi, Senior Researcher**

[sunglok@etri.re.kr](mailto:sunglok@etri.re.kr) | <http://sites.google.com/site/sunglok>

Electronics and Telecommunication Research Institute (ETRI)



An Invitation  
~~Introduction~~ to 3D Vision  
: A Tutorial for Everyone

Sunglok Choi, Senior Researcher

[sunglok@etri.re.kr](mailto:sunglok@etri.re.kr) | <http://sites.google.com/site/sunglok>

Electronics and Telecommunication Research Institute (ETRI)

# What is 3D Vision?

## Computer Vision

**What** is it?

- Label (e.g. Tower of Pisa)
- Shape (e.g. )



**Where** is it?

- Place (e.g. Piazza del Duomo, Pisa, Italy)
- Location (e.g. )



# What is 3D Vision?

## Computer Vision

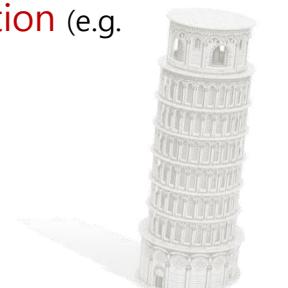
**What** is it?

- **Label** (e.g. Tower of Pisa)
- **Shape** (e.g. )



**Where** is it?

- **Place** (e.g. Piazza del Duomo, Pisa, Italy)
- **Location** (e.g. )



(84, 10, 18) [m]



**Recognition Problems v.s. Reconstruction Problems**

# What is 3D Vision?

Visual Geometry (Multiple View Geometry)

Geometric Vision

## Computer Vision

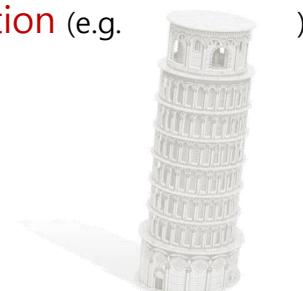
**What** is it?

- **Label** (e.g. Tower of Pisa)
- **Shape** (e.g. )



**Where** is it?

- **Place** (e.g. Piazza del Duomo, Pisa, Italy)
- **Location** (e.g. )



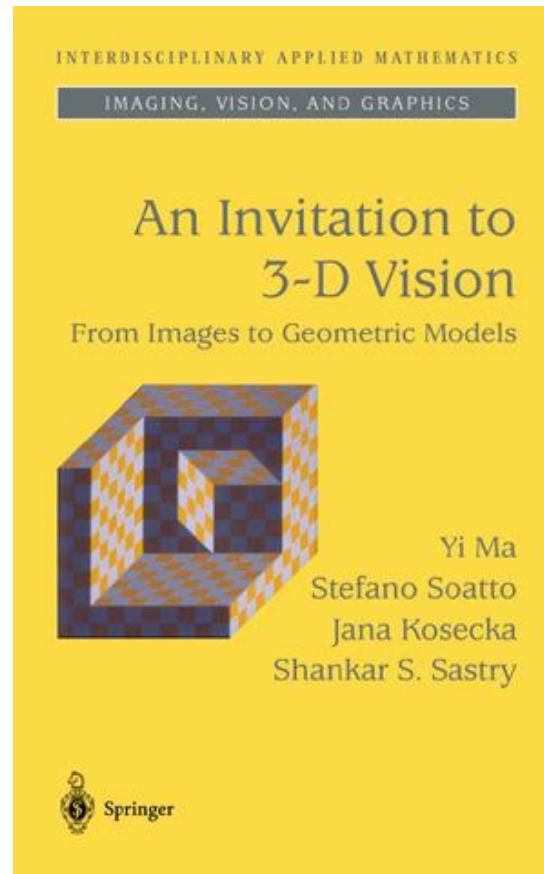
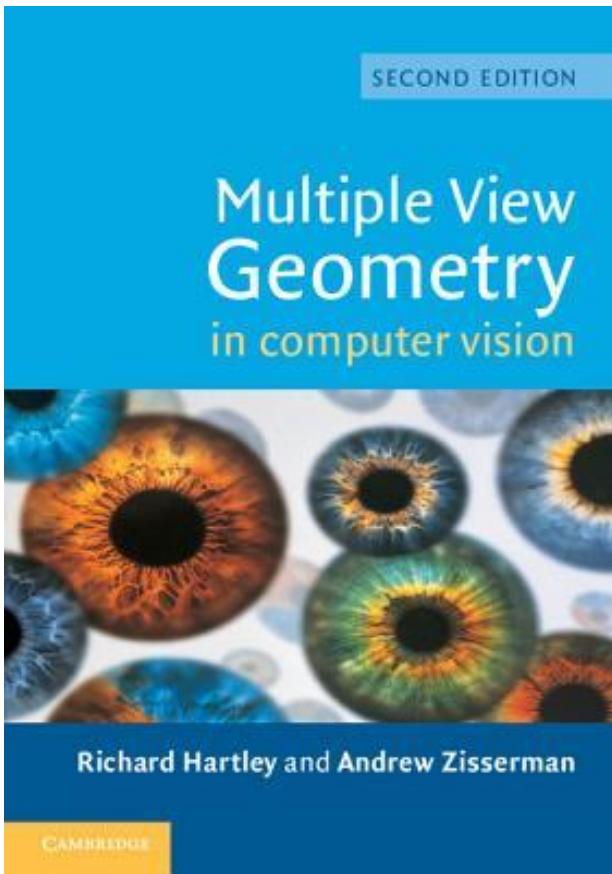
(84, 10, 18) [m]



Recognition Problems v.s. Reconstruction Problems

# What is 3D Vision?

- Reference Books

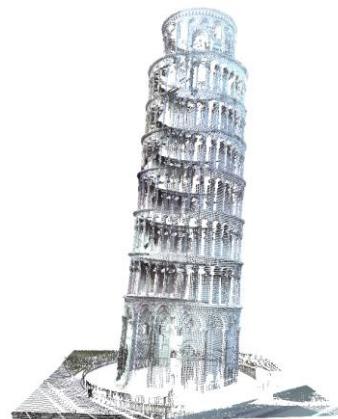




# What is 3D Vision?

- OpenCV (Open Source Computer Vision)
  - calib3d Module: Camera Calibration and 3D Reconstruction
    - OpenCV API Reference
      - Introduction
      - core. The Core Functionality
      - imgproc. Image Processing
      - imgcodecs. Image file reading and writing
      - videoio. Media I/O
      - highgui. High-level GUI and Media I/O
      - video. Video Analysis
      - calib3d. Camera Calibration and 3D Reconstruction
      - features2d. 2D Features Framework
      - objdetect. Object Detection
      - ml. Machine Learning
      - flann. Clustering and Search in Multi-Dimensional Spaces
      - photo. Computational Photography
      - stitching. Images stitching
      - cuda. CUDA-accelerated Computer Vision
      - cudaarithm. CUDA-accelerated Operations on Matrices
      - cudabgsegm. CUDA-accelerated Background Segmentation
      - cudacodec. CUDA-accelerated Video Encoding/Decoding
      - cudafeatures2d. CUDA-accelerated Feature Detection and Description
      - cudafilters. CUDA-accelerated Image Filtering
      - cudaimproc. CUDA-accelerated Image Processing
      - cudaoptflow. CUDA-accelerated Optical Flow
      - cudastereo. CUDA-accelerated Stereo Correspondence
      - cudawarping. CUDA-accelerated Image Warping
      - shape. Shape Distance and Matching
      - superres. Super Resolution
      - videostab. Video Stabilization
      - viz. 3D Visualizer

# What is 3D Vision?



point cloud, depth/range image, polygon mesh, ...



Perspective Camera

**3D Vision**

v.s.

**3D Data Processing**



RGB-D Camera  
(Stereo, Structured Light, ToF, Light Field)

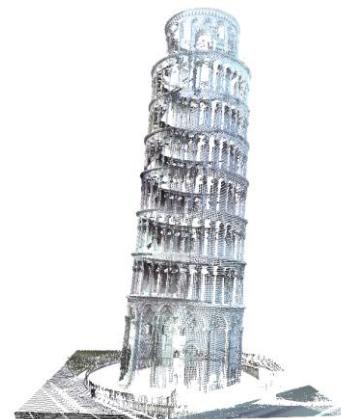


Omni-directional Camera



Range Sensor  
(LiDAR, RADAR)

# What is 3D Vision?



point cloud, depth/range image, polygon mesh, ...



Perspective Camera

**3D Vision**

v.s.

**3D Data Processing**



RGB-D Camera

(Stereo, Structured Light, ToF, Light Field)



Omni-directional Camera

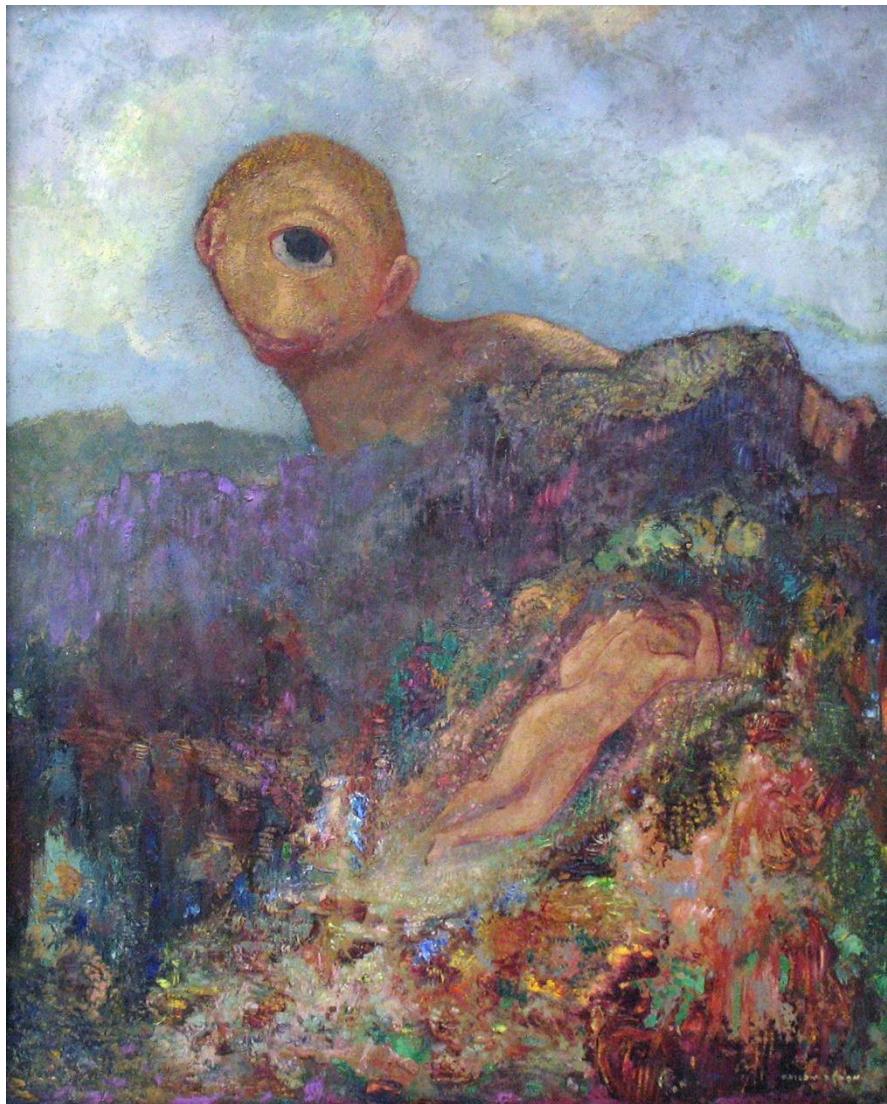


Range Sensor  
(LiDAR, RADAR)

# Table of Contents

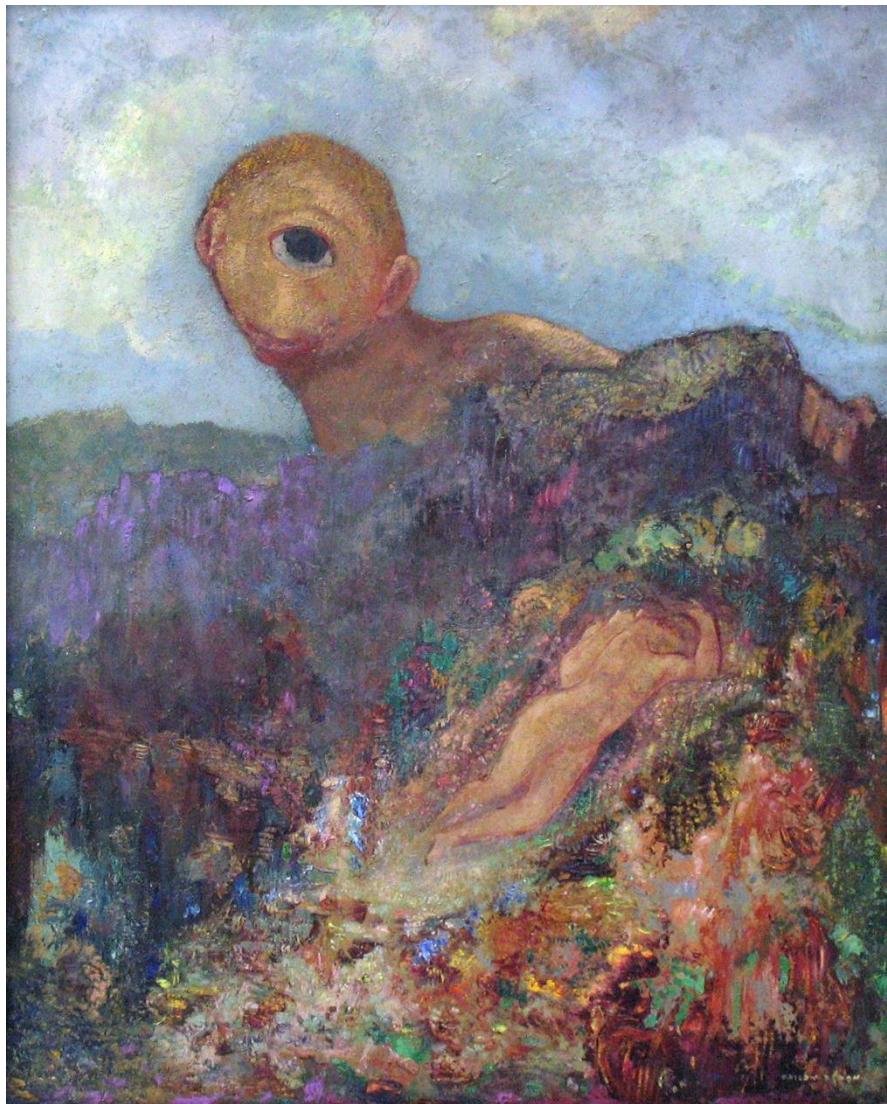
- **What is 3D Vision?**
- **Single-view Geometry**
  - Camera Projection Model
  - General 2D-3D Geometry
- **Two-view Geometry**
  - Planar 2D-2D Geometry (**Projective Geometry**)
  - General 2D-2D Geometry (**Epipolar Geometry**)
- **Multi-view Geometry**
- **Correspondence Problem**
- **Summary**

# Single-view Geometry



*The Cyclops*, gouache and oil by Odilon Redon

# Single-view Geometry



Camera Projection Model  
General 2D-3D Geometry

*The Cyclops*, gouache and oil by Odilon Redon

# Camera Projection Model

- The Pinhole Camera Model



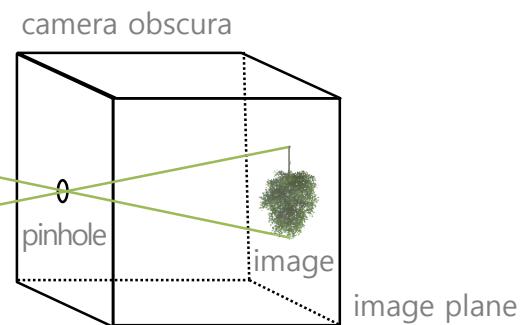
A large-scale camera obscura  
at San Francisco, California



A modern-day camera obscura

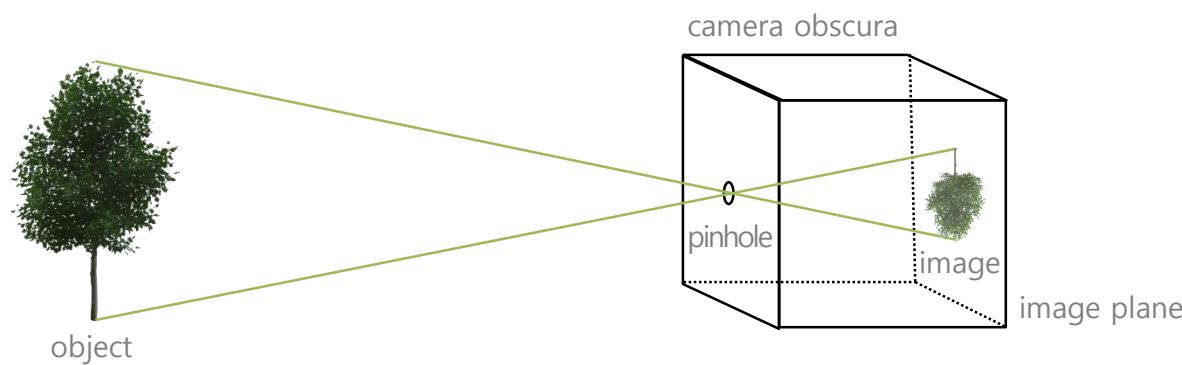
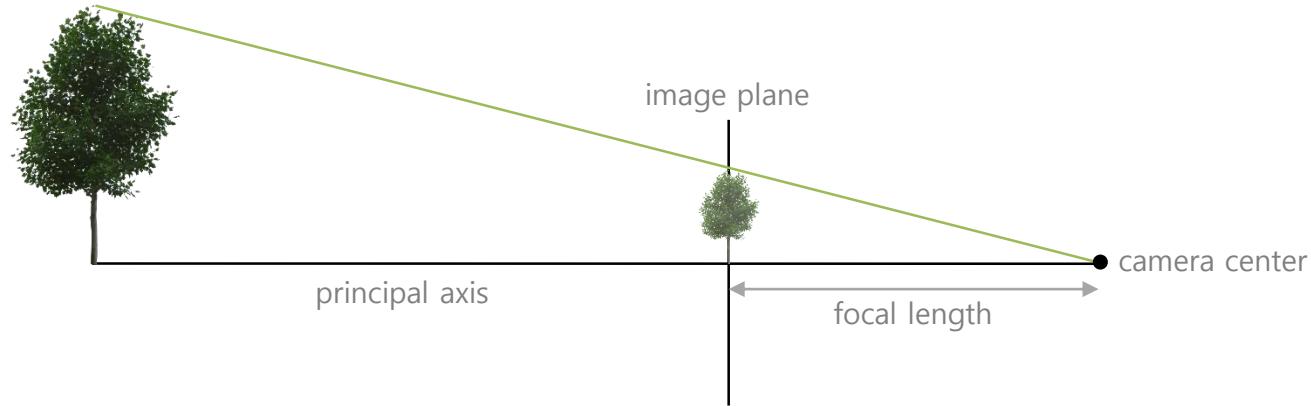


An Image in camera obscura  
at Portslade, England



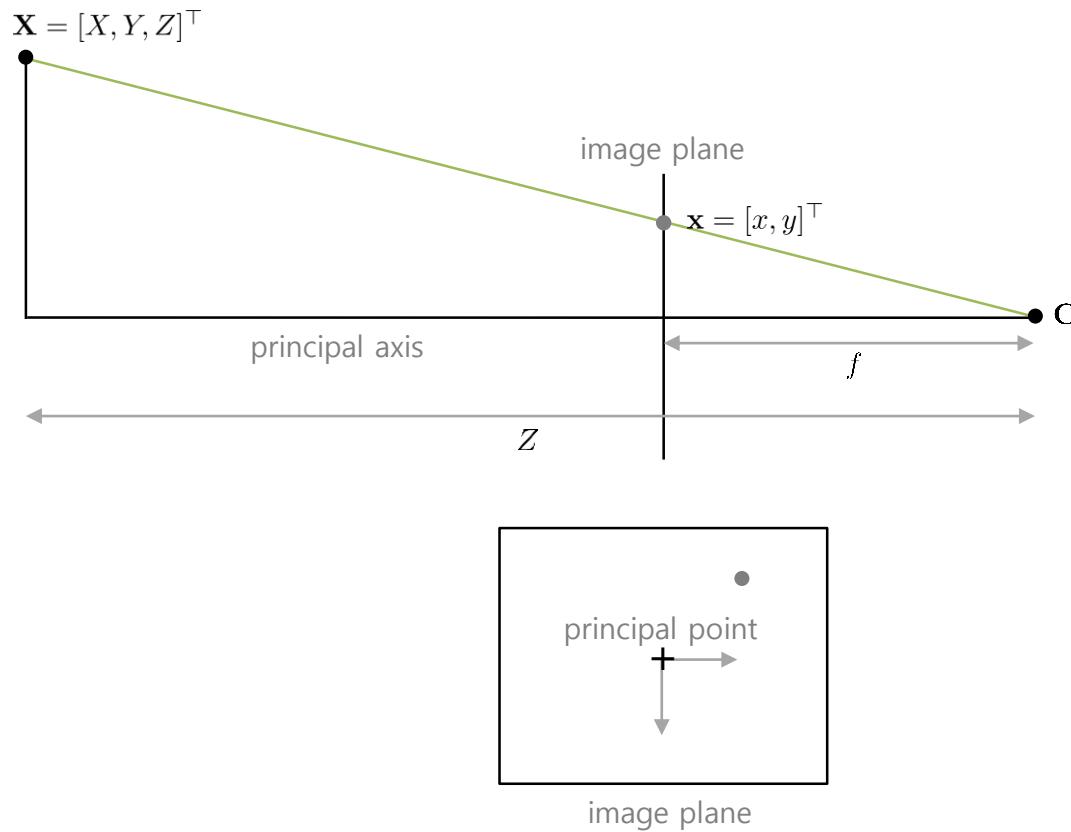
# Camera Projection Model

- The Pinhole Camera Model



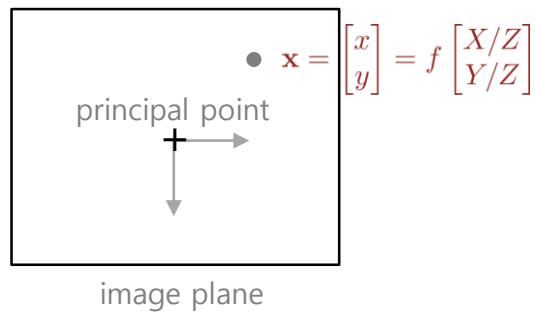
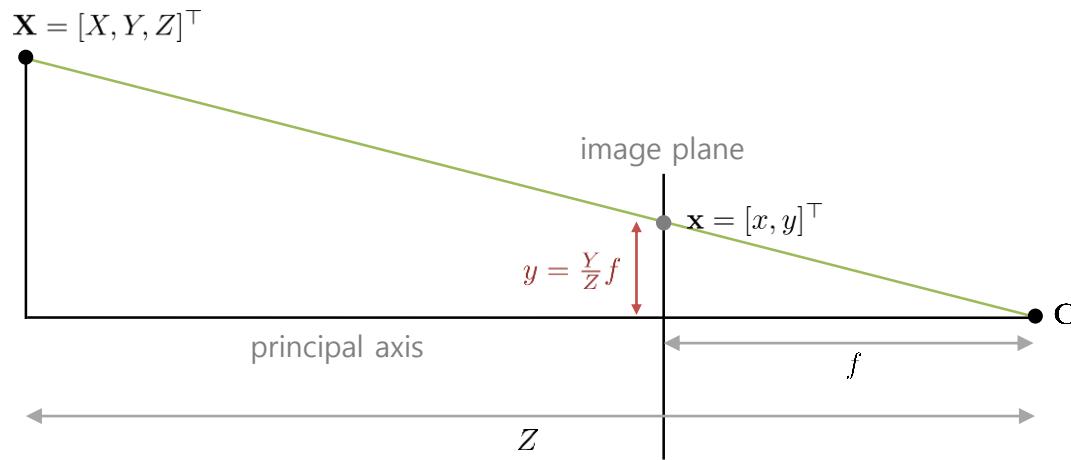
# Camera Projection Model

- The Pinhole Camera Model



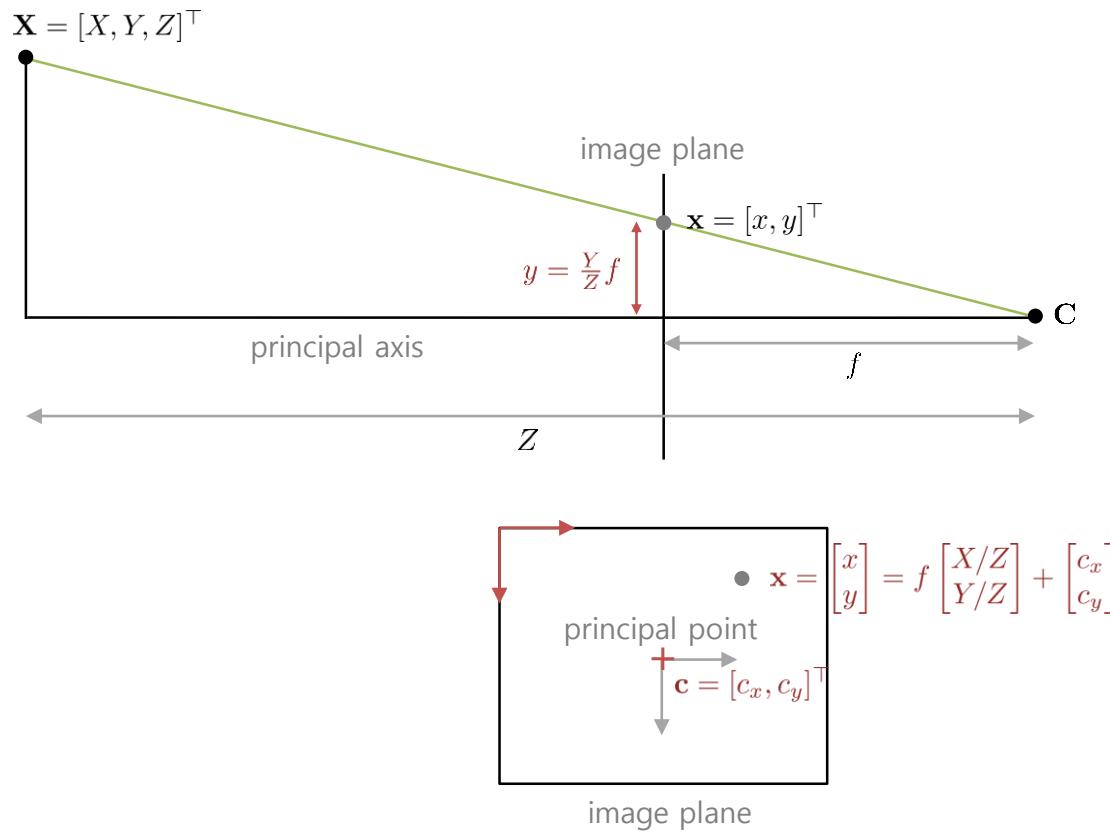
# Camera Projection Model

- The Pinhole Camera Model



# Camera Projection Model

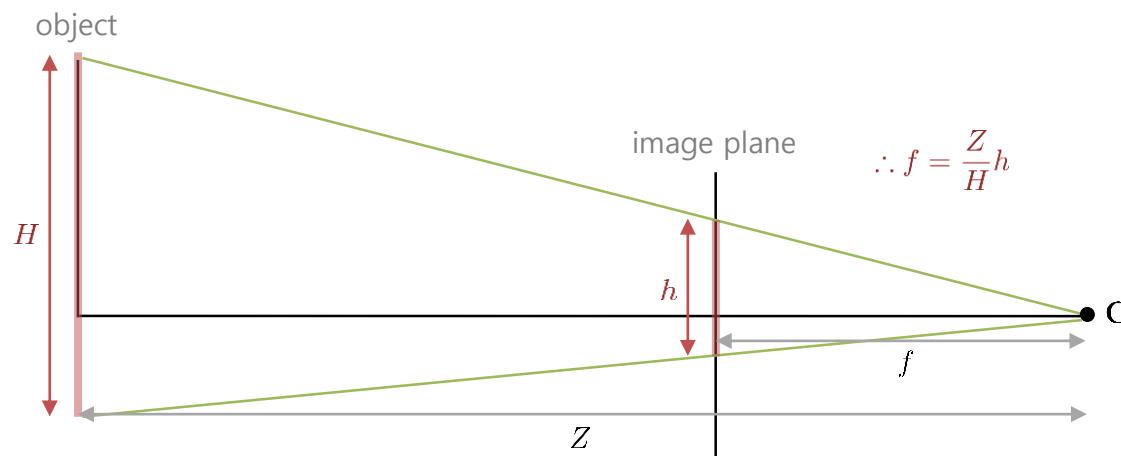
- The Pinhole Camera Model



# Camera Projection Model

- The Pinhole Camera Model

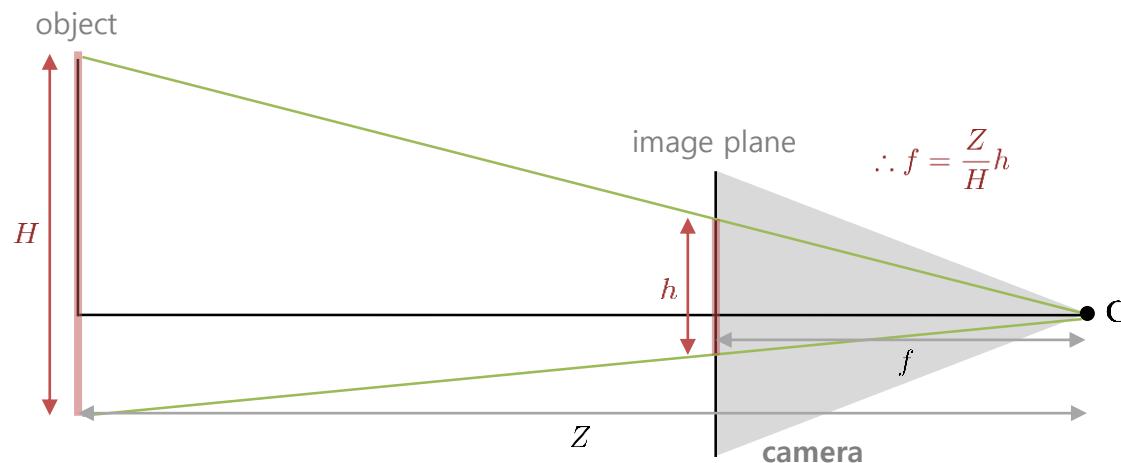
- Example: **Simple Camera Calibration #1**
  - Unknown: **Focal length** (unit: [pixel])
  - Assumptions
    - The object length and distance from the camera are known.
    - The object is aligned with the image plane.



# Camera Projection Model

- The Pinhole Camera Model

- Example: Simple Camera Calibration #1
  - Unknown: **Focal length** (unit: [pixel])
  - Assumptions
    - The object length and distance from the camera are known.
    - The object is aligned with the image plane.



# Camera Projection Model

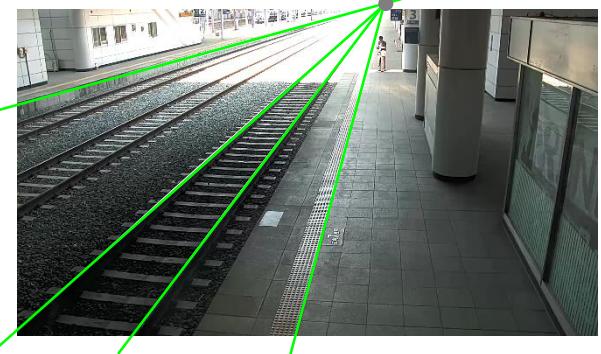
- The Pinhole Camera Model

- Example: **Simple Camera Calibration #2**

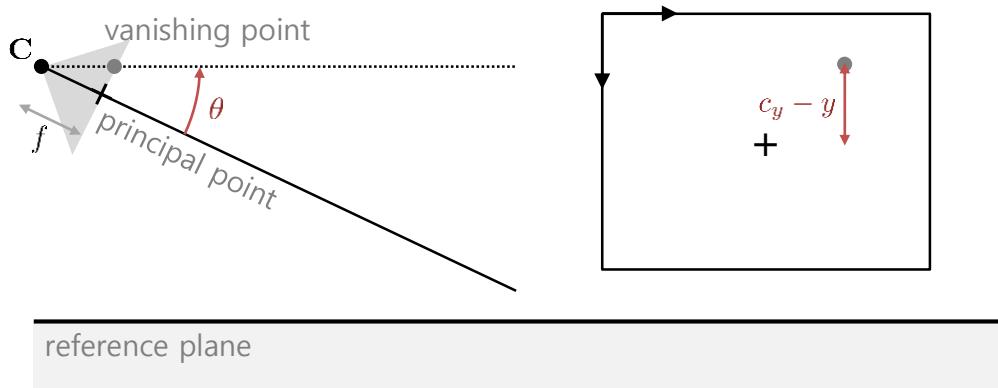
- Unknown: **Tilt angle** (unit: [rad])

- Assumptions

- The focal length is known.
      - The principal point is known or selected as the center of images.
      - A vanishing point from the reference plane is known.



$$\therefore \theta = \tan^{-1} \frac{c_y - y}{f}$$



# Camera Projection Model

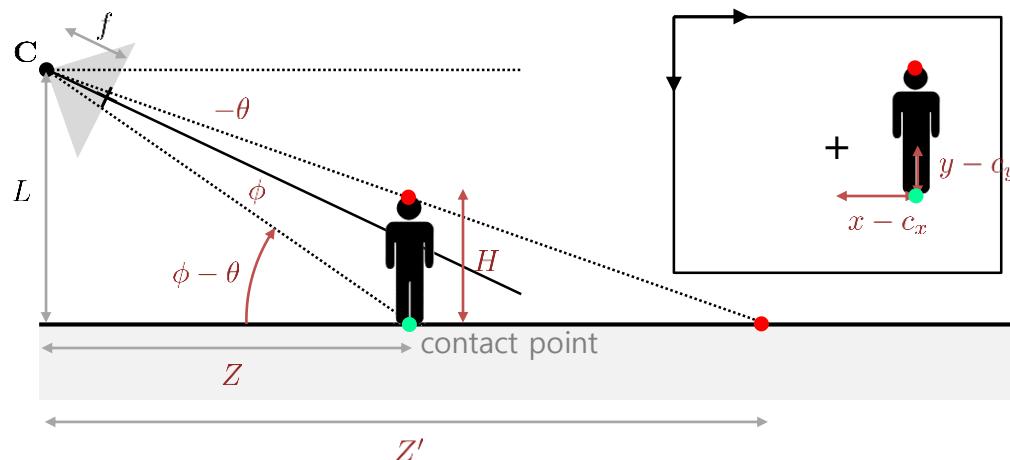
- The Pinhole Camera Model

- Example: Simple Object Localization
  - Unknown: **Position** (and also **height**)
  - Assumptions
    - The focal length, principal points, camera height, and tilt angle are known.
    - The object is on the reference plane.



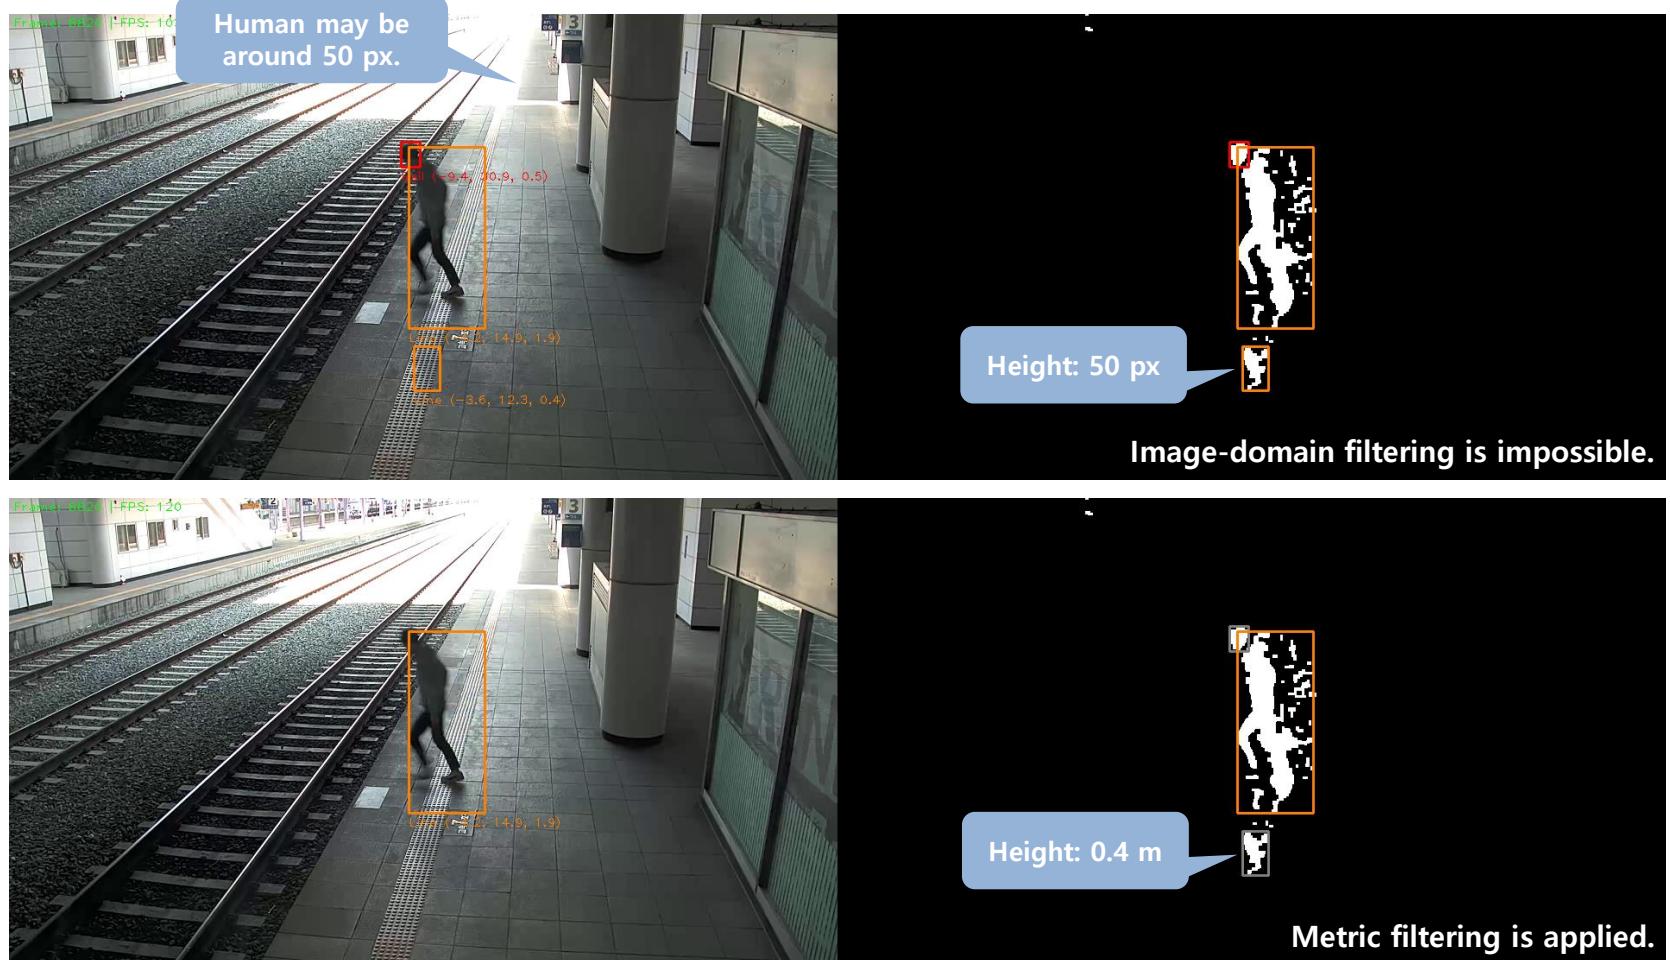
$$\therefore Z = \frac{H}{\tan(\phi - \theta)} \quad (\phi = \tan^{-1} \frac{y - c_y}{f})$$

$$X = \frac{x - c_x}{f} Z \quad H = \frac{Z' - Z}{Z'} L$$



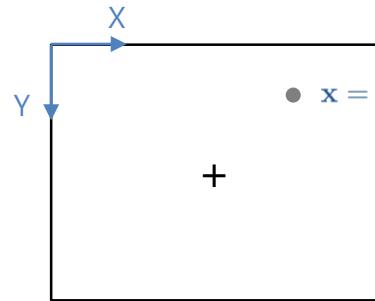
# Camera Projection Model

- The Pinhole Camera Model
  - Example: Simple Object Localization



# Camera Projection Model

- Camera Matrix (Intrinsic Parameters)



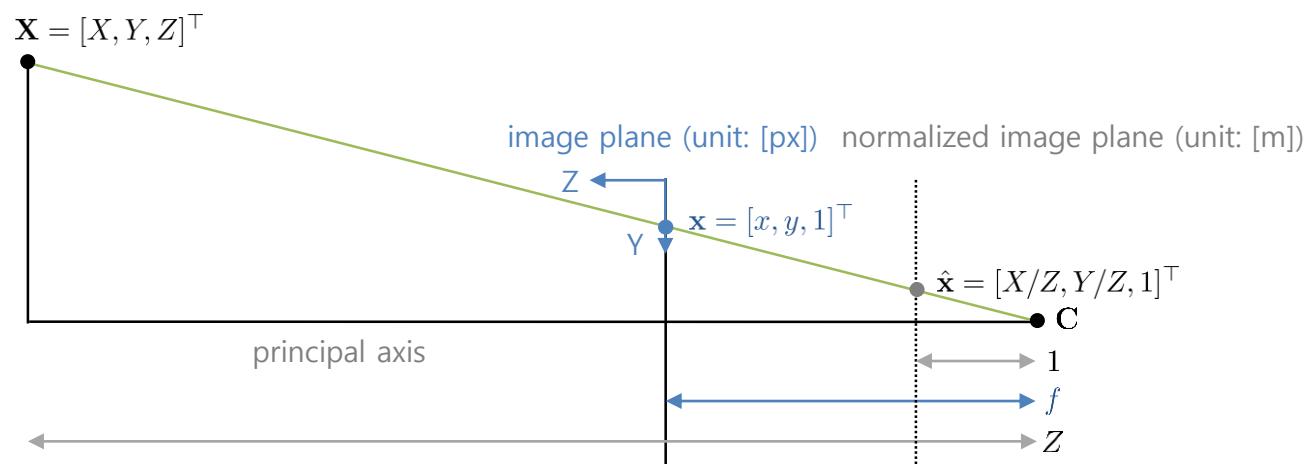
A diagram showing a 2D image plane with a coordinate system having X and Y axes. A point  $\mathbf{x}$  is shown on the plane, represented by a dot and a bracketed vector  $[x, y]$ . A plus sign (+) is placed below the image plane.

$$\bullet \quad \mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix} = f \begin{bmatrix} X/Z \\ Y/Z \end{bmatrix} + \begin{bmatrix} c_x \\ c_y \end{bmatrix} \longrightarrow \mathbf{x} = \mathbf{K}\hat{\mathbf{x}}$$

where  $\mathbf{K} = \begin{bmatrix} f & 0 & c_x \\ 0 & f & c_y \\ 0 & 0 & 1 \end{bmatrix}$ ,  $\mathbf{x} = \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$ , and  $\hat{\mathbf{x}} = \begin{bmatrix} X/Z \\ Y/Z \\ 1 \end{bmatrix}$

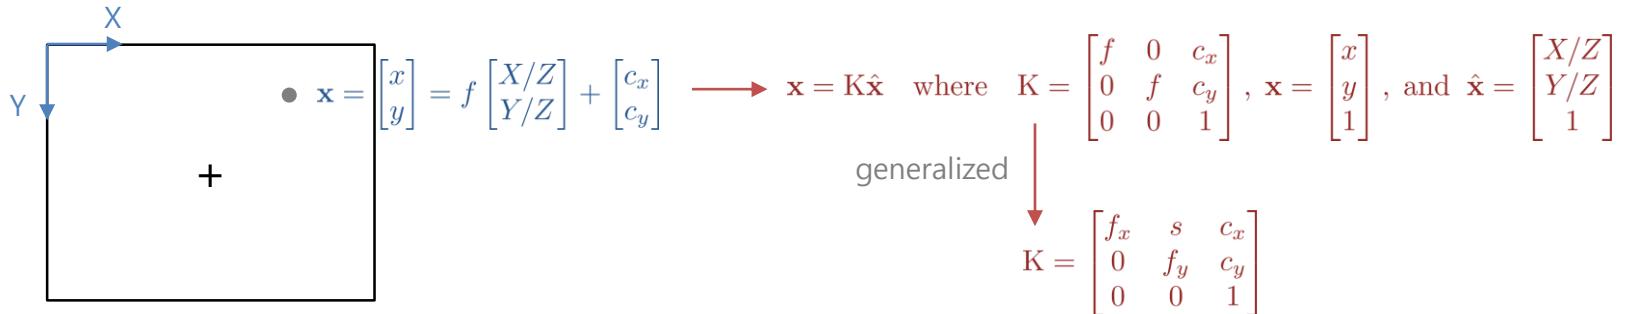
generalized

$$\mathbf{K} = \begin{bmatrix} f_x & s & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$



# Camera Projection Model

- **Camera Matrix (Intrinsic Parameters)**



- **Projection Matrix (Intrinsic/Extrinsic Parameters)**

- If a 3D point is not based on the camera coordinate, the point should be transformed.

$$\mathbf{X}' = \mathbf{R}\mathbf{X} + \mathbf{t} \longrightarrow \mathbf{X}' = [\mathbf{R} \mid \mathbf{t}] \mathbf{X} \text{ where } \mathbf{X} = \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

3x4 matrix

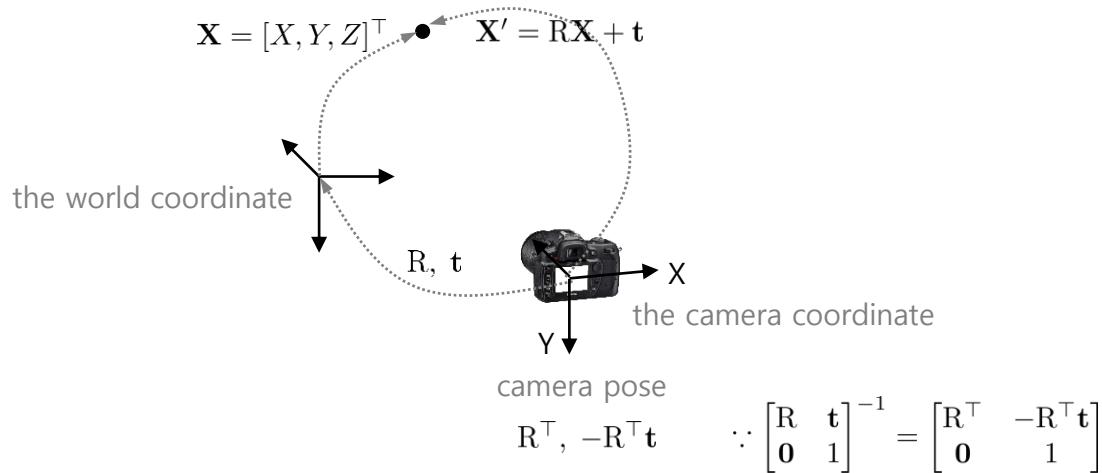
$$\mathbf{x} = \mathbf{P}\mathbf{X} \text{ where } \mathbf{P} = \mathbf{K}[\mathbf{R} \mid \mathbf{t}] \text{ and } \mathbf{x} = \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

3x4 matrix

$\because \mathbf{X}'$  is not normalized.

- No problem! To get its projected point on the image plane, we can normalize it later,  $\begin{bmatrix} x/w \\ y/w \\ 1 \end{bmatrix}$ .

# Camera Projection Model



- **Projection Matrix** (Intrinsic/Extrinsic Parameters)

- If a 3D point is not based on the camera coordinate, the point should be transformed.

$$\mathbf{X}' = \mathbf{R}\mathbf{X} + \mathbf{t} \longrightarrow \mathbf{X}' = \begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix} \mathbf{X} \quad \text{where } \mathbf{X} = \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

3x4 matrix

$$\mathbf{x} = \mathbf{P}\mathbf{X} \quad \text{where } \mathbf{P} = \mathbf{K} \begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}$$

3x4 matrix

$$\mathbf{x} = \begin{bmatrix} x \\ y \\ w \end{bmatrix} \quad \because \mathbf{X}' \text{ is not normalized.}$$

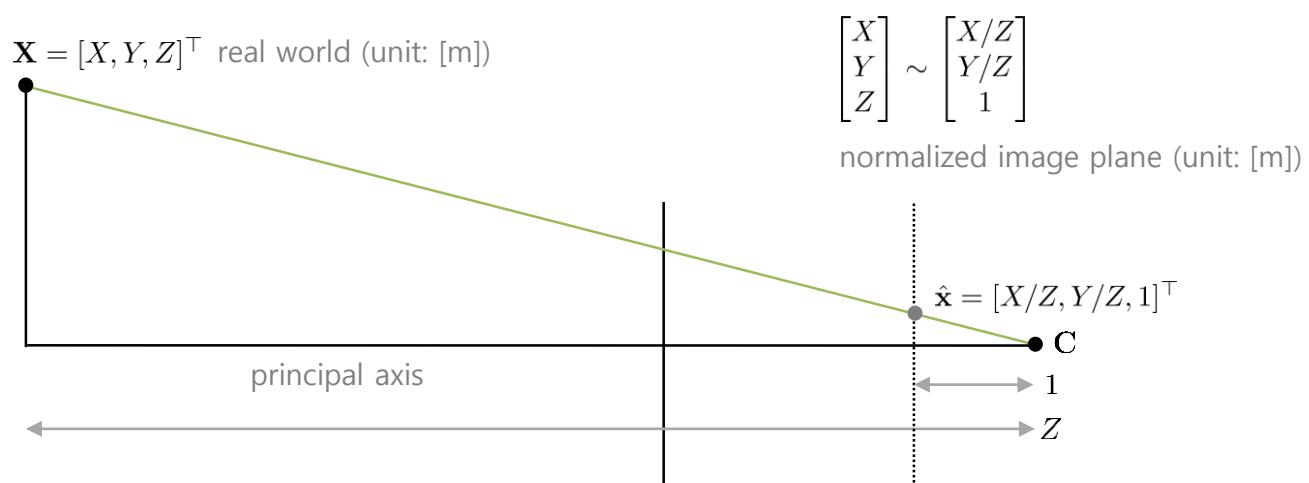
- No problem! To get its projected point on the image plane, we can normalize it later,  $\begin{bmatrix} x/w \\ y/w \\ 1 \end{bmatrix}$ .

# Camera Projection Model

$$\therefore \mathbf{x} = \mathbf{P}\mathbf{X} \quad \text{where} \quad \mathbf{P} = \mathbf{K}[\mathbf{R} | \mathbf{t}], \quad \mathbf{x} = \begin{bmatrix} x \\ y \\ w \end{bmatrix}, \quad \text{and} \quad \mathbf{X} = \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

- **Why Homogeneous Coordinate?**

- A linear transformation (rotation and translation) is applied by a single matrix multiplication.
- A light ray can be represented as a point on the image plane.
- An infinite point (a.k.a. ideal point) is numerically represented by  $w = 0$ .
- A point and line are represented beautifully as like  $\mathbf{l}^T \mathbf{x} = 0$  or  $\mathbf{x}^T \mathbf{l} = 0$ .
- ...

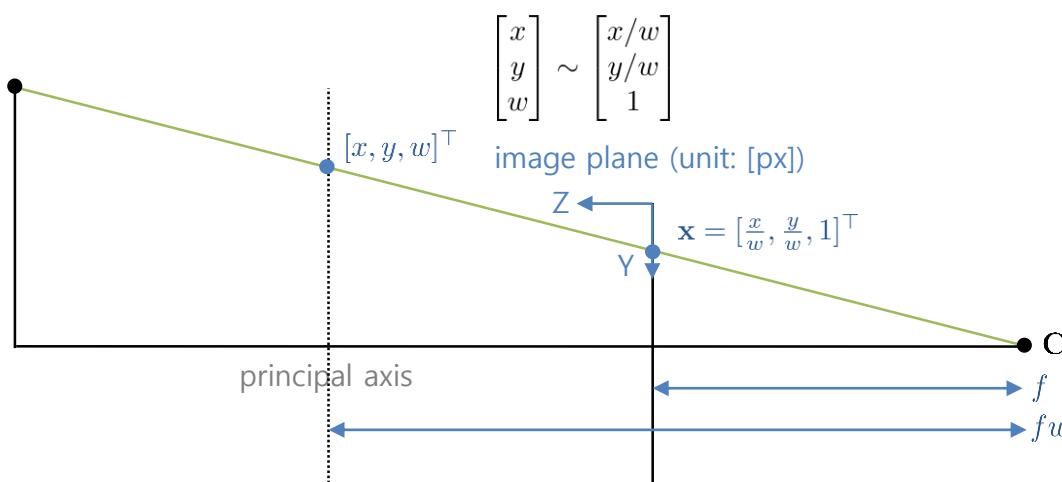


# Camera Projection Model

$$\therefore \mathbf{x} = \mathbf{P}\mathbf{X} \quad \text{where} \quad \mathbf{P} = \mathbf{K}[\mathbf{R} | \mathbf{t}], \quad \mathbf{x} = \begin{bmatrix} x \\ y \\ w \end{bmatrix}, \quad \text{and} \quad \mathbf{X} = \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

- **Why Homogeneous Coordinate?**

- A linear transformation (rotation and translation) is applied by a single matrix multiplication.
- A light ray can be represented as a point on the image plane.
- An infinite point (a.k.a. ideal point) is numerically represented by  $w = 0$ .
- A point and line are represented beautifully as like  $\mathbf{l}^\top \mathbf{x} = 0$  or  $\mathbf{x}^\top \mathbf{l} = 0$ .
- ...



# Image Formation using OpenCV

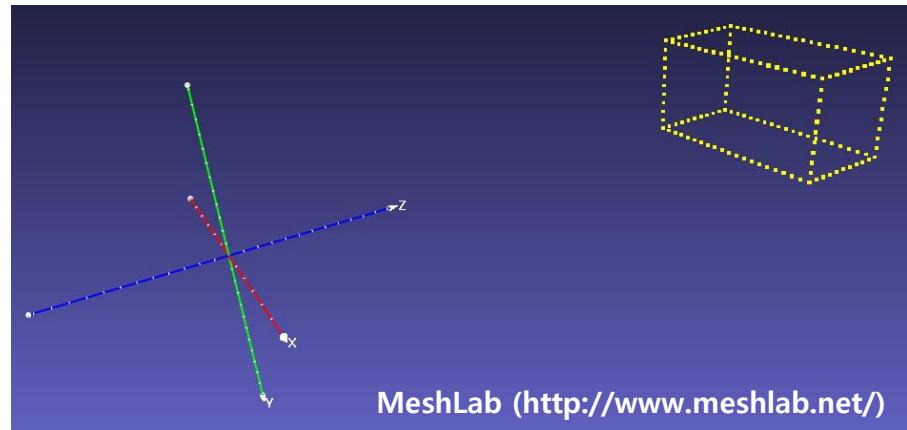


```
1. #include "opencv_all.hpp"

2. #define Rx(rx) (cv::Mat_<double>(3, 3) << 1, 0, 0, 0, cos(rx), -sin(rx), 0, sin(rx), cos(rx))
3. #define Ry(ry) (cv::Mat_<double>(3, 3) << cos(ry), 0, sin(ry), 0, 1, 0, -sin(ry), 0, cos(ry))
4. #define Rz(rz) (cv::Mat_<double>(3, 3) << cos(rz), -sin(rz), 0, sin(rz), cos(rz), 0, 0, 0, 1)

5. int main(void)
6. {
7.     // The given camera configuration: focal length, principal point, image resolution, position, and orientation
8.     double camera_focal = 1000;
9.     cv::Point2d camera_center(320, 240);
10.    cv::Size camera_res(640, 480);
11.    cv::Point3d camera_pos[] = { cv::Point3d(0, 0, 0), cv::Point3d(-2, -2, 0), ... };
12.    cv::Point3d camera_ori[] = { cv::Point3d(0, 0, 0), cv::Point3d(-CV_PI / 12, CV_PI / 12, 0), ... };
13.    double camera_noise = 1;

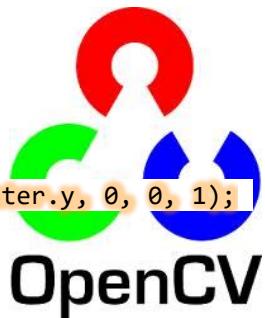
14.    // Load a point cloud in the homogeneous coordinate
15.    FILE* fin = fopen("data/box.xyz", "rt");
16.    if (fin == NULL) return -1;
17.    cv::Mat X;
18.    while (!feof(fin))
19.    {
20.        double x, y, z;
21.        if (fscanf(fin, "%lf %lf %lf", &x, &y, &z) == 3)
22.        {
23.            X.push_back<double>(x);
24.            X.push_back<double>(y);
25.            X.push_back<double>(z);
26.            X.push_back<double>(1);
27.        }
28.    }
29.    fclose(fin);
30.    X = X.reshape(1, X.rows / 4).t(); // Convert to a 4 x N matrix
```



MeshLab (<http://www.meshlab.net/>)

A point cloud: data/box.xyz

# Image Formation using OpenCV



```
31. // Generate images for each camera pose
32. cv::Mat K = (cv::Mat_<double>(3, 3) << camera_focal, 0, camera_center.x, 0, camera_focal, camera_center.y, 0, 0, 1);
33. for (int i = 0; i < sizeof(camera_pos) / sizeof(cv::Point3d); i++)
34. {
35.     // Derive a projection matrix
36.     cv::Mat Rc = Rz(camera_ori[i].z) * Ry(camera_ori[i].y) * Rx(camera_ori[i].x);
37.     cv::Mat tc = (cv::Mat_<double>(3, 1) << camera_pos[i].x, camera_pos[i].y, camera_pos[i].z);
38.     cv::Mat Rt;
39.     cv::hconcat(Rc.t(), -Rc.t() * tc, Rt);
40.     cv::Mat P = K * Rt;

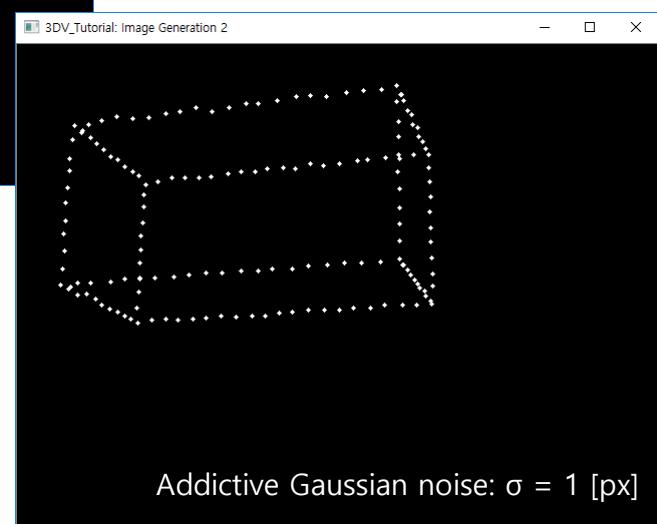
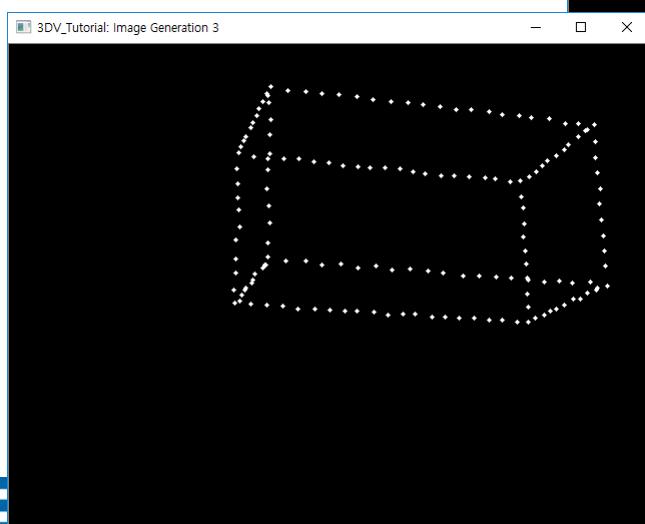
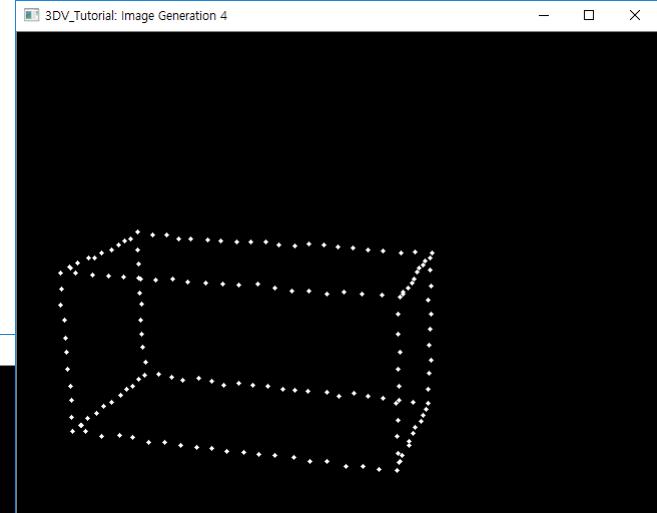
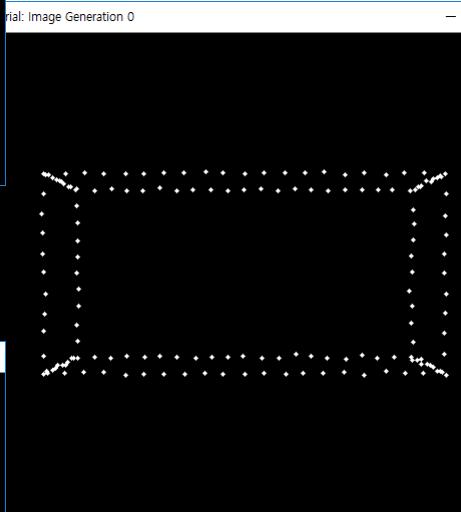
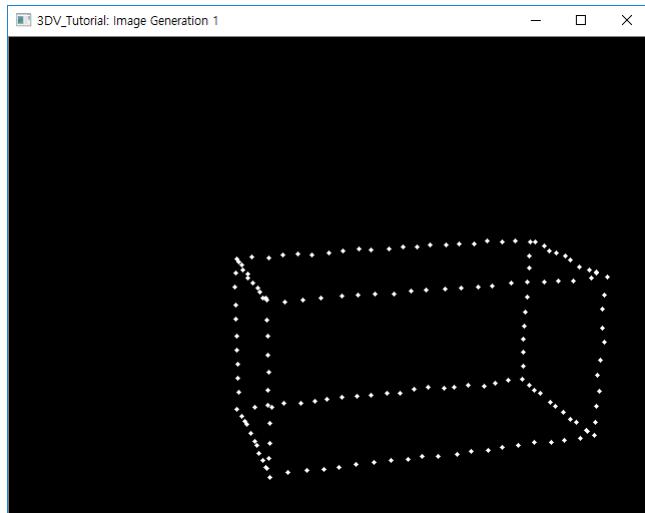
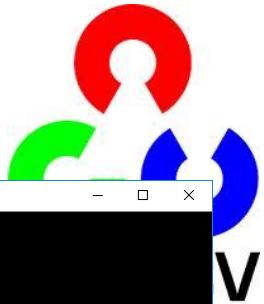
41.     // Project the points (c.f. OpenCV provide 'cv::projectPoints()' with consideration of distortion.)
42.     cv::Mat x = P * X;
43.     x.row(0) = x.row(0) / x.row(2);
44.     x.row(1) = x.row(1) / x.row(2);
45.     x.row(2) = 1;

46.     cv::Mat noise(2, x.cols, x.type());
47.     cv::randn(noise, cv::Scalar(0), cv::Scalar(camera_noise));
48.     x.rowRange(0, 2) = x.rowRange(0, 2) + noise; // Add noise

49.     // Show and store the points
50.     cv::Mat image = cv::Mat::zeros(camera_res, CV_8UC1);
51.     for (int c = 0; c < x.cols; c++)
52.     {
53.         cv::Point p(x.at<double>(0, c), x.at<double>(1, c));
54.         if (p.x >= 0 && p.x < camera_res.width && p.y >= 0 && p.y < camera_res.height)
55.             cv::circle(image, p, 2, 255, -1);
56.     }
57.     cv::imshow(cv::format("3DV_Tutorial: Image Formation %d", i), image);

58.     FILE* fout = fopen(cv::format("image_formation%d.xyz", i).c_str(), "wt");
59.     if (fout == NULL) return -1;
60.     for (int c = 0; c < x.cols; c++)
61.         fprintf(fout, "%f %f 1\n", x.at<double>(0, c), x.at<double>(1, c));
62.     fclose(fout);
63. }
```

# Image Formation using OpenCV

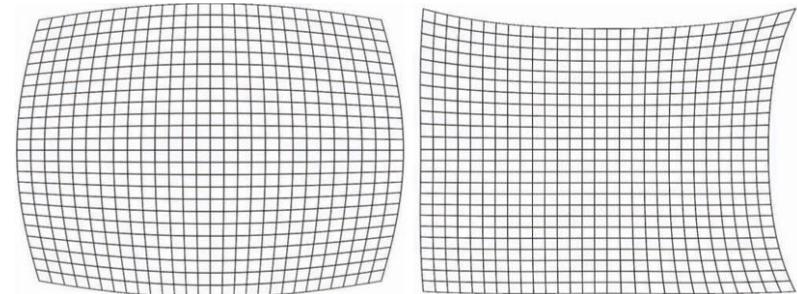


Addictive Gaussian noise:  $\sigma = 1$  [px]

# Camera Projection Model

- **Geometric Distortion**

- Radial/Tangential Distortion Model



$$\begin{bmatrix} \hat{x}_d \\ \hat{y}_d \end{bmatrix} = (1 + k_1 r^2 + k_2 r^4 + \dots) \begin{bmatrix} \hat{x} \\ \hat{y} \end{bmatrix} + \begin{bmatrix} 2p_1 \hat{x} \hat{y} + p_2 (r^2 + 2\hat{x}^2) \\ 2p_2 \hat{x} \hat{y} + p_1 (r^2 + 2\hat{y}^2) \end{bmatrix}$$

where  $r^2 = \hat{x}^2 + \hat{y}^2$

defined on the normalized image plane

- Radial distortion:  $k_1, k_2, \dots$
- Tangential distortion:  $p_1, p_2$

- Image Rectification

- OpenCV `cv::undistort()` and `cv::undistortPoints()` (c.f. included in `imgproc` module)
    - Camera Distortion Correction: Theory and Practice, <http://darkpgmr.tistory.com/31>



correction

K1: 1.105763E-01  
K2: 1.886214E-02  
K3: 1.473832E-02  
P1:-8.448460E-03  
P2:-7.356744E-03



# Distortion Correction using OpenCV



```
1. #include "opencv_all.hpp"

2. int main(void)
3. {
4.     cv::Mat K = (cv::Mat<double>(3, 3) << 432.7390364738057, 0, 476.0614994349778, ..., 0, 0, 1);
5.     cv::Mat dist_coeff = (cv::Mat<double>(5, 1) << -0.2852754904152874, 0.1016466459919075, ...);

6.     // Open an video
7.     cv::VideoCapture video;
8.     if (!video.open("data/chessboard.avi")) return -1;

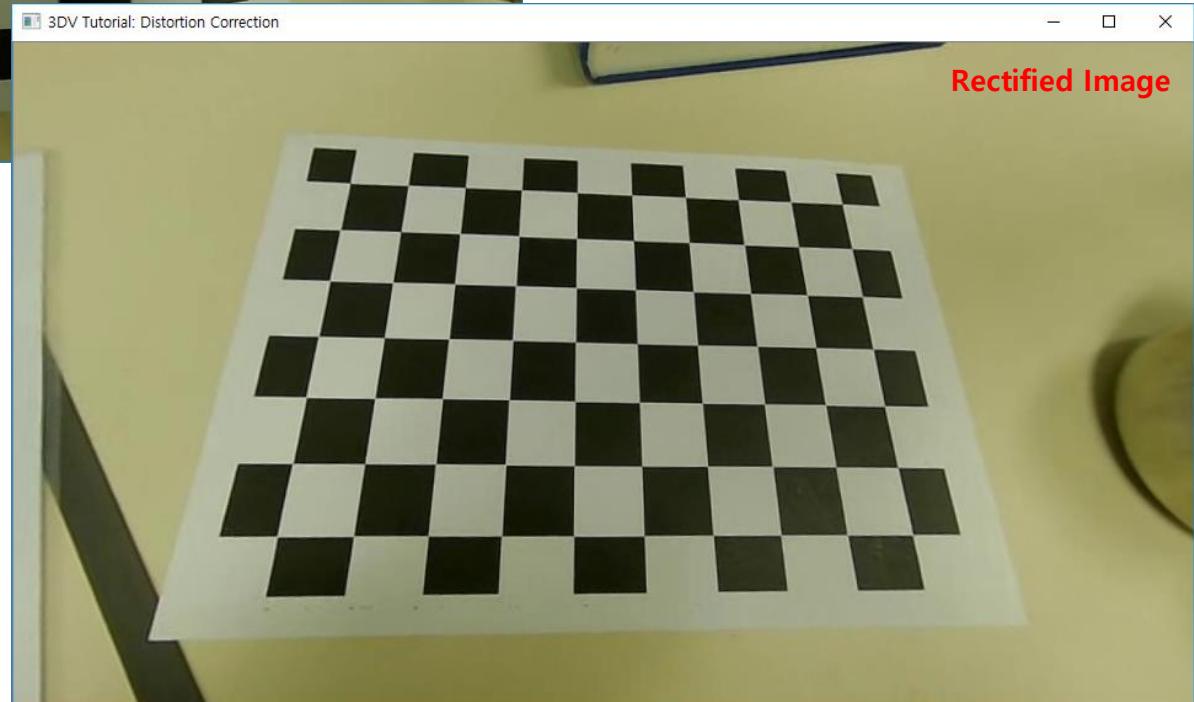
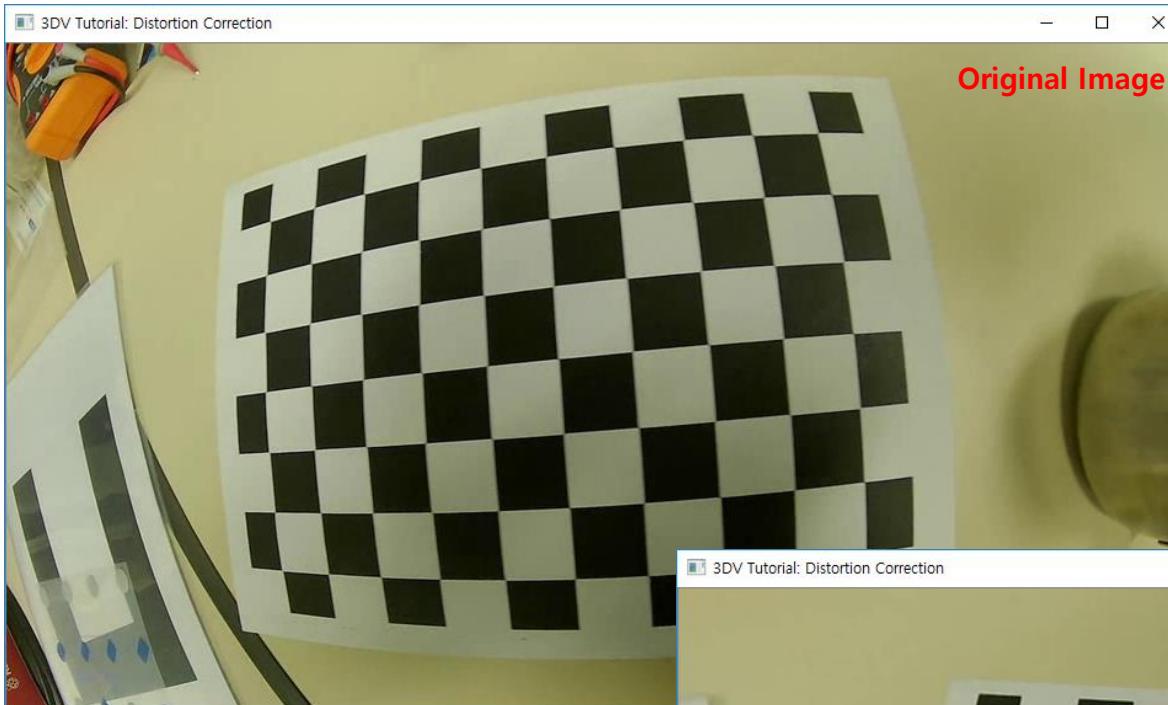
9.     // Run distortion correction
10.    bool show_rectify = true;
11.    cv::Mat map1, map2;
12.    while (true)
13.    {
14.        // Grab an image from the video
15.        cv::Mat image;
16.        video >> image;
17.        if (image.empty()) break;

18.        // Rectify geometric distortion (c.f. 'cv::undistort()' can be applied for one-time remapping.)
19.        if (show_rectify)
20.        {
21.            if (map1.empty() || map2.empty())
22.                cv::initUndistortRectifyMap(K, dist_coeff, cv::Mat(), cv::Mat(), image.size(), CV_32FC1, map1, map2);
23.            cv::remap(image, image, map1, map2, cv::InterpolationFlags::INTER_LINEAR);
24.        }

25.        // Show the image
26.        cv::imshow("3DV Tutorial: Distortion Correction", image);
27.        int key = cv::waitKey(1);
28.        if (key == 27) break;                                // 'ESC' key: Exit
29.        else if (key == 9) show_rectify = !show_rectify;    // 'Tab' key: Toggle rectification
30.        ...
31.    }

32.    video.release();
33.    return 0;
34. }
```

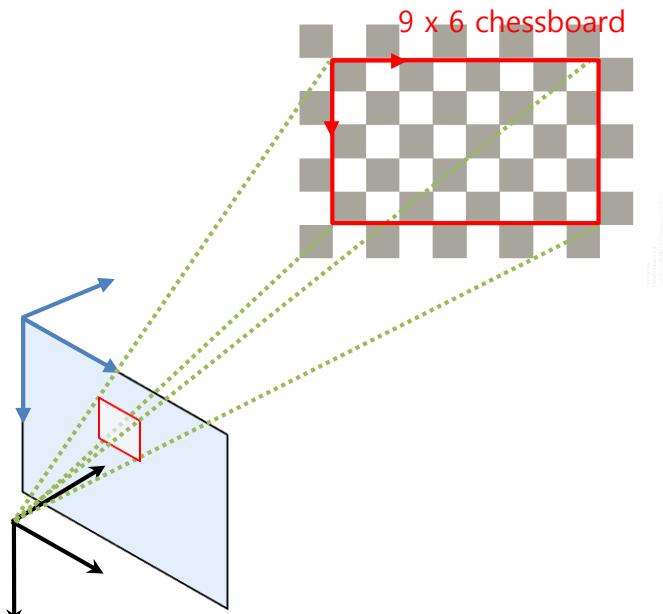
# Distortion Correction using OpenCV



# General 2D-3D Geometry

- **Camera Calibration**

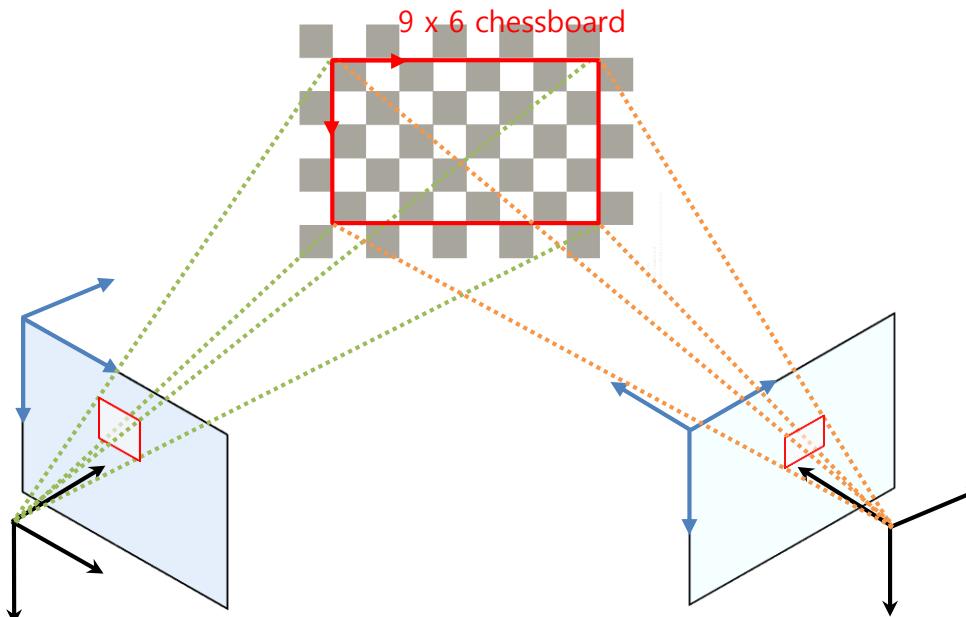
- Known: 3D points  $\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n$  and their projected points  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$
- Unknown: 5 (intrinsic w/o distortion) + 6 (extrinsic) DoF
- Constraints:  $n \times$  projection  $\mathbf{x}_i = \mathbf{P}\mathbf{X}_i$



# General 2D-3D Geometry

## ▪ Camera Calibration

- Known: 3D points  $\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n$  and their projected points from each camera  $\mathbf{x}_i^j$
- Unknown: 5 (intrinsic w/o distortion) +  $m \times 6$  (extrinsic) DoF
- Constraints:  $n \times m \times$  projection  $\mathbf{x}_i^j = \mathbf{K}[\mathbf{R}_j | \mathbf{t}_j] \mathbf{X}_i$
- Solutions
  - OpenCV `cv::calibrateCamera()` and `cv::initCameraMatrix2D()`
  - Camera Calibration Toolbox for MATLAB, [http://www.vision.caltech.edu/bouguetj/calib\\_doc/](http://www.vision.caltech.edu/bouguetj/calib_doc/)
  - GML C++ Camera Calibration Toolbox, <http://graphics.cs.msu.ru/en/node/909>
  - DarkCamCalibrator, <http://darkpgmr.tistory.com/139>



# Camera Calibration using OpenCV



```
1. #include "opencv_all.hpp"

2. int main(void)
3. {
4.     bool select_images = true;
5.     cv::Size board_pattern(10, 7);
6.     double board_cellsize = 0.025;

7.     // Select images
8.     std::vector<cv::Mat> images;
9.     ...

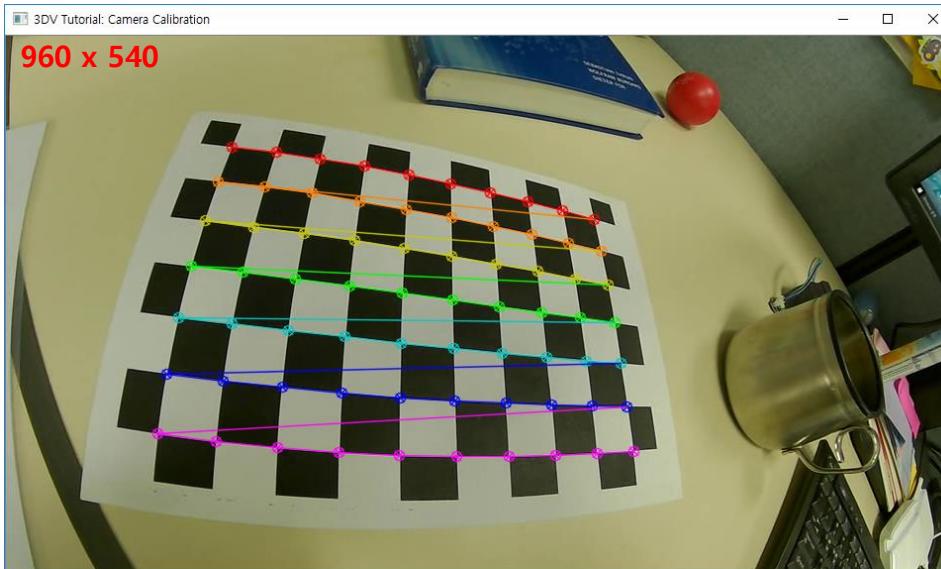
10.    // Find 2D corner points from the given images
11.    std::vector<std::vector<cv::Point2f> > image_points;
12.    for (size_t i = 0; i < images.size(); i++)
13.    {
14.        std::vector<cv::Point2f> pts;
15.        if (cv::findChessboardCorners(images[i], board_pattern, pts))
16.            image_points.push_back(pts);
17.    }
18.    if (image_points.empty()) return -1;

19.    // Prepare 3D points of the chess board
20.    std::vector<std::vector<cv::Point3f> > object_points(1);
21.    for (int r = 0; r < board_pattern.height; r++)
22.        for (int c = 0; c < board_pattern.width; c++)
23.            object_points[0].push_back(cv::Point3f(board_cellsize * c, board_cellsize * r, 0));
24.    object_points.resize(image_points.size(), object_points[0]); // Copy

25.    // Calibrate the camera
26.    cv::Mat K = cv::Mat::eye(3, 3, CV_64F);
27.    cv::Mat dist_coeff = cv::Mat::zeros(4, 1, CV_64F);
28.    std::vector<cv::Mat> rvecs, tvecs;
29.    double rms = cv::calibrateCamera(object_points, image_points, images[0].size(), K, dist_coeff, rvecs, tvecs);

30.    // Report calibration results
31.    ...
32.    return 0;
33.}
```

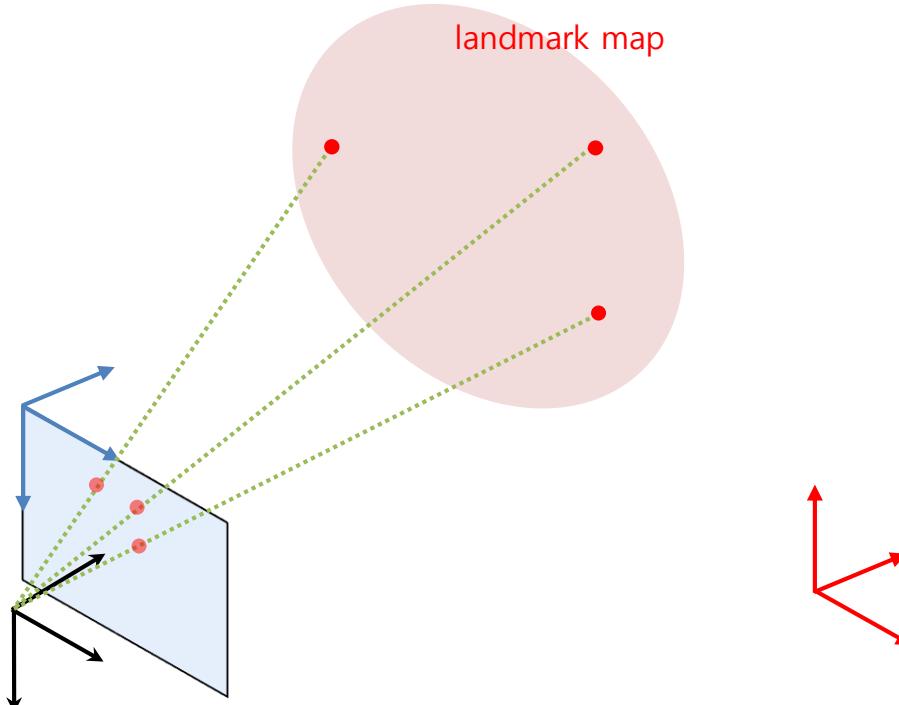
# Camera Calibration using OpenCV



```
## Camera Calibration Results
* The number of applied images = 22
* RMS error = 0.473353
* Camera matrix (K) =
[432.7390364738057, 0, 476.0614994349778]
[0, 431.2395555913084, 288.7602152621297]
[0, 0, 1]
* Distortion coefficient (k1, k2, p1, p2, k3, ...) =
[-0.2852754904152874, 0.1016466459919075,
 -0.0004420196146339175, 0.0001149909868437517,
 -0.01803978785585194]
```

# General 2D-3D Geometry

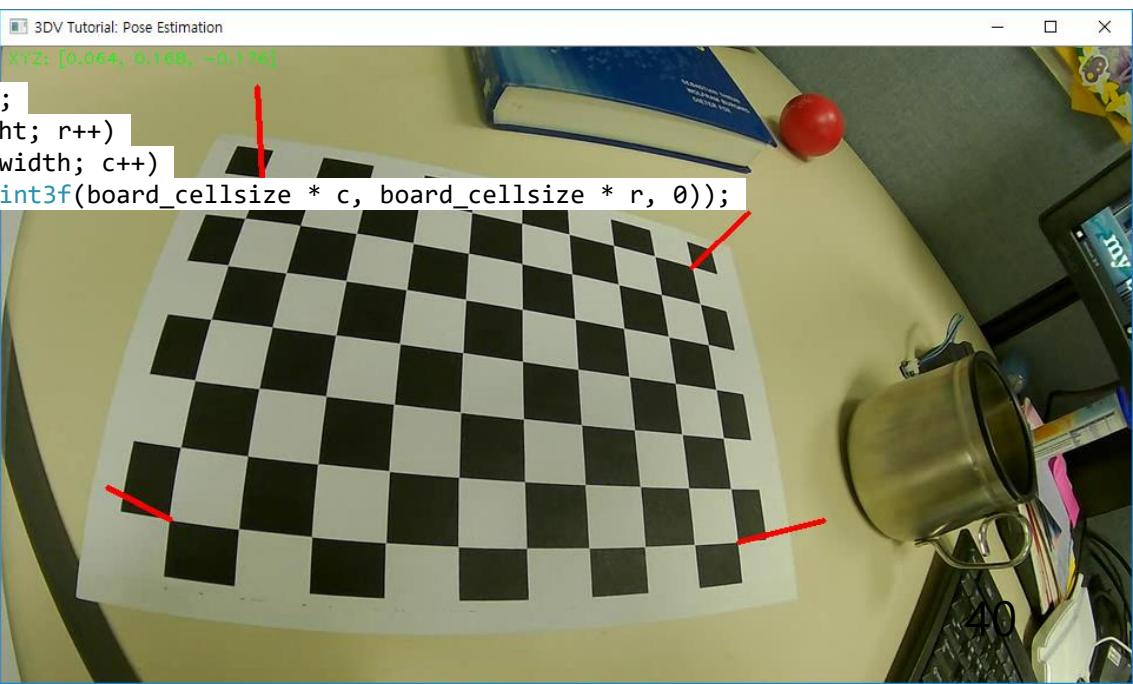
- **Absolute Camera Pose Estimation** (Perspective-n-Point; PnP)
  - Known: 3D points  $\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n$ , their projected points  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$ , and camera matrix  $\mathbf{K}$
  - Unknown: **Camera pose (6 DoF)**
  - Constraints:  $n \times$  projection  $\mathbf{x}_i = \mathbf{K}[\mathbf{R} | \mathbf{t}] \mathbf{X}_i$
  - Solutions ( $n \geq 3$ )
    - OpenCV `cv::solvePnP()` and `cv::solvePnPRansac()`
    - Efficient PnP (EPnP), <http://cvlab.epfl.ch/EPnP/>



# Pose Estimation using OpenCV



```
1. #include "opencv_all.hpp"  
2. int main(void)  
3. {  
4.     cv::Mat K = (cv::Mat<double>(3, 3) << 432.7390364738057, 0, 476.0614994349778, ...);  
5.     cv::Mat dist_coeff = (cv::Mat<double>(5, 1) << -0.2852754904152874, 0.1016466459919075, ...);  
6.     cv::Size board_pattern(10, 7);  
7.     double board_cellsize = 0.025;  
8.  
9.     // Open an video  
10.    cv::VideoCapture video;  
11.    if (!video.open("data/chessboard.avi")) return -1;  
12.  
13.    // Prepare four sticks for simple AR  
14.    std::vector<cv::Point3f> stick_points;  
15.    stick_points.push_back(cv::Point3f(0, 0, 0));  
16.    stick_points.push_back(cv::Point3f(0, 0, -board_cellsize * 2));  
17.    ...  
18.  
19.    // Run pose estimation  
20.    std::vector<cv::Point3f> object_points;  
21.    for (int r = 0; r < board_pattern.height; r++)  
22.        for (int c = 0; c < board_pattern.width; c++)  
23.            object_points.push_back(cv::Point3f(board_cellsize * c, board_cellsize * r, 0));
```



# Pose Estimation using OpenCV



```
21.     while (true)
22.     {
23.         // Grab an image from the video
24.         cv::Mat image;
25.         video >> image;
26.         if (image.empty()) break;
27.
28.         // Estimate camera pose
29.         std::vector<cv::Point2f> image_points;
30.         bool complete = cv::findChessboardCorners(image, board_pattern, image_points, ...);
31.         if (complete)
32.         {
33.             cv::Mat rvec, tvec;
34.             cv::solvePnP(object_points, image_points, K, dist_coeff, rvec, tvec);
35.
36.             // Draw four sticks
37.             std::vector<cv::Point2f> line_points;
38.             cv::projectPoints(stick_points, rvec, tvec, K, dist_coeff, line_points);
39.             for (int i = 0; i < 8; i += 2)
40.                 cv::line(image, line_points[i], line_points[i + 1], cv::Scalar(0, 0, 255), 4);
41.
42.             // Print camera position
43.             cv::Mat R;
44.             cv::Rodrigues(rvec, R);
45.             cv::Mat p = -R.t() * tvec;
46.             cv::String info = cv::format("XYZ: [% .3f, % .3f, % .3f]", p.at<double>(0), p.at<double>(1), p.at<double>(2));
47.             cv::putText(image, info, cv::Point(5, 15), cv::FONT_HERSHEY_PLAIN, 1, cv::Scalar(0, 255, 0));
48.         }
49.
50.         // Show the image
51.         ...
52.     }
53.
54.     video.release();
55.     return 0;
56. }
```

# Two-view Geometry

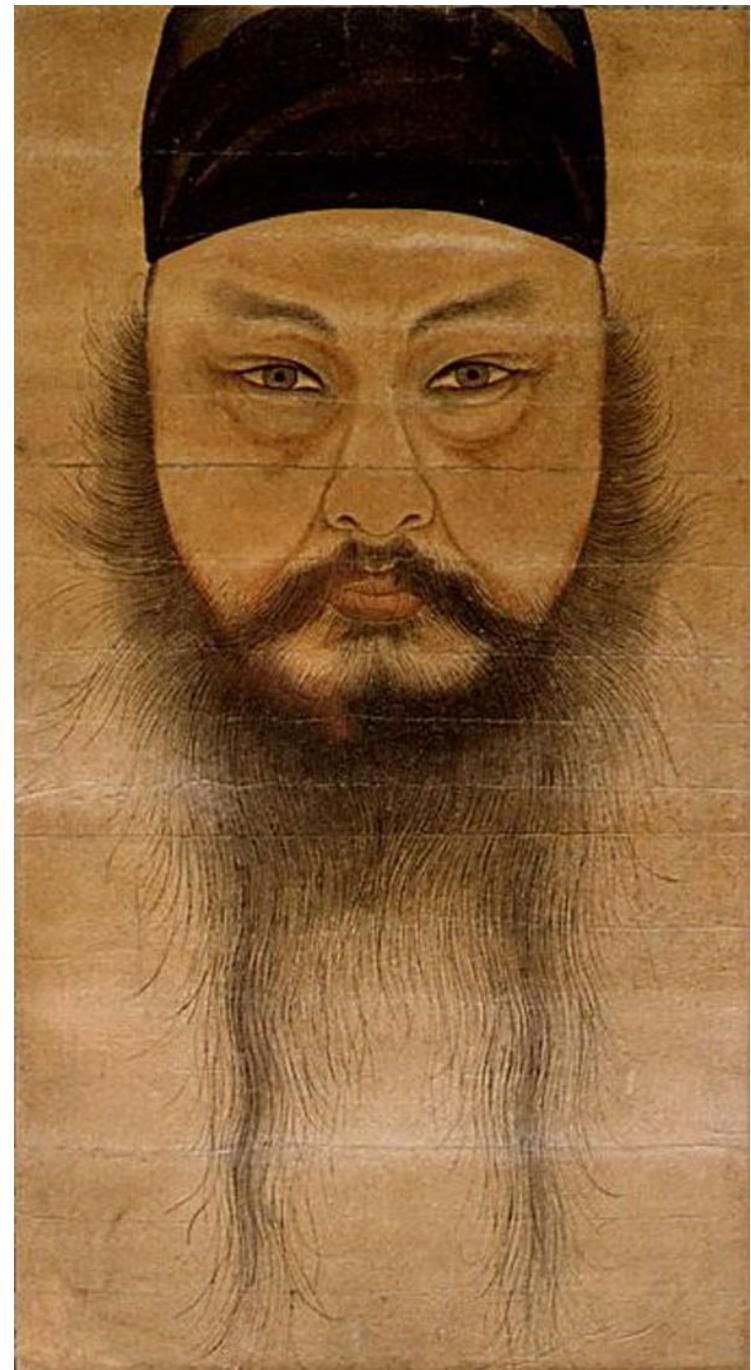


윤두서(1668-1715) 자화상, 국보 제240호  
Korean National Treasure No. 240

# Two-view Geometry

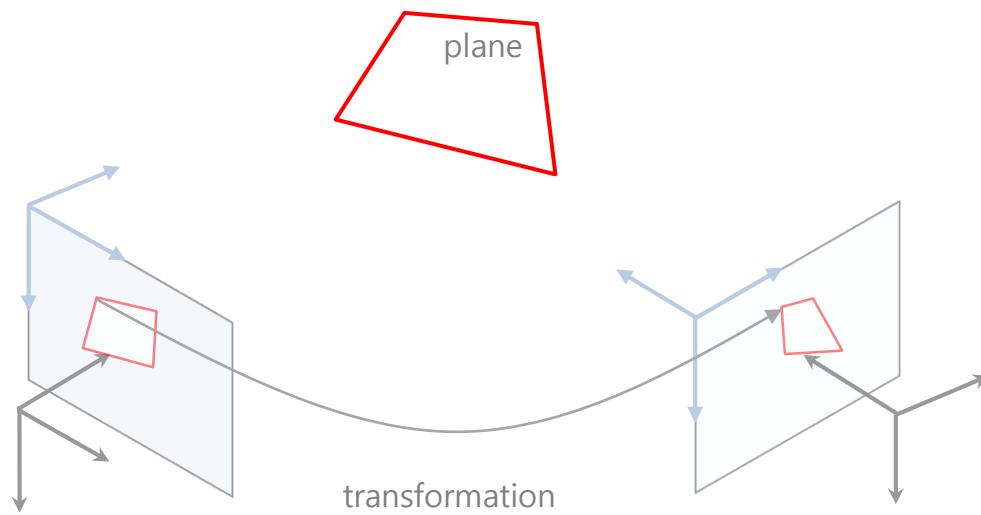
**Planar** 2D-2D Geometry (Projective Geometry)

**General** 2D-2D Geometry (Epipolar Geometry)

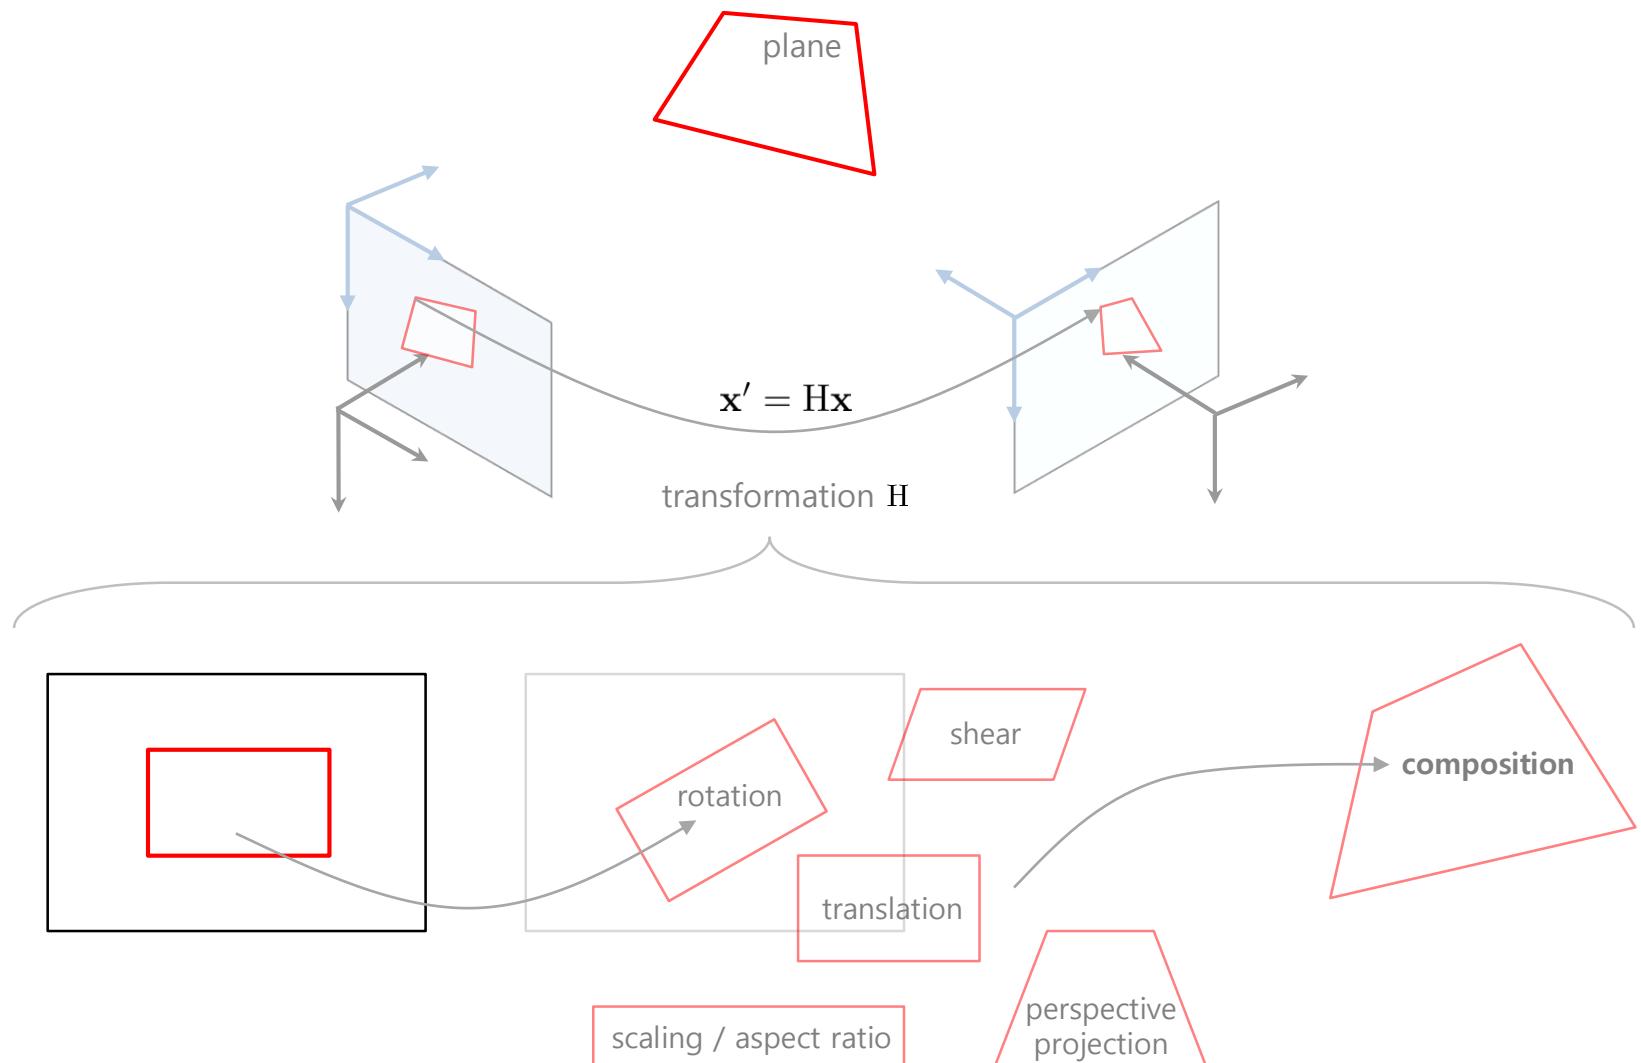


윤두서(1668-1715) 자화상, 국보 제240호  
Korean National Treasure No. 240

# Planar 2D-2D Geometry (Projective Geometry)



# Planar 2D-2D Geometry (Projective Geometry)



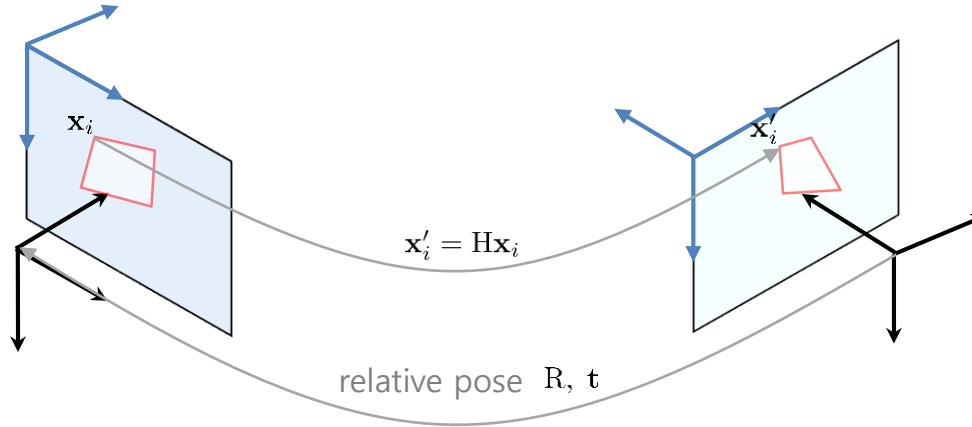
# Planar 2D-2D Geometry (Projective Geometry)

	Euclidean Transform (a.k.a. Rigid Transform)	Similarity Transform	Affine Transform	Projective Transform (a.k.a. Planar Homography)
<b>Matrix Forms H</b>	$\begin{bmatrix} \cos \theta & -\sin \theta & t_x \\ \sin \theta & \cos \theta & t_y \\ 0 & 0 & 1 \end{bmatrix}$	$\begin{bmatrix} s \cos \theta & -s \sin \theta & t_x \\ s \sin \theta & s \cos \theta & t_y \\ 0 & 0 & 1 \end{bmatrix}$	$\begin{bmatrix} a_{11} & a_{12} & t_x \\ a_{21} & a_{22} & t_y \\ 0 & 0 & 1 \end{bmatrix}$	$\begin{bmatrix} a_{11} & a_{12} & t_x \\ a_{21} & a_{22} & t_y \\ v_1 & v_2 & 1 \end{bmatrix}$
<b>DoF</b>	3	4	6	8
<b>Transformations</b> - rotation - translation - scaling - aspect ratio - shear - perspective projection	O O	O O O	O O O O O O	O O O O O O O
<b>Invariants</b> - length - angle - ratio of lengths - parallelism - incidence - cross ratio	O O O O O O	O O O O O O	O O O	O O
<b>OpenCV Functions</b>			<code>cv::getAffineTransform()</code> <code>cv::estimateRigidTransform()</code> -	<code>cv::getPerspectiveTransform()</code> - <code>cv::findHomography()</code> <code>cv::warpAffine()</code> <code>cv::warpPerspective()</code>

# Planar 2D-2D Geometry (Projective Geometry)

## ▪ Planar Homography Estimation

- Known: Point correspondence  $(\mathbf{x}_1, \mathbf{x}'_1), \dots, (\mathbf{x}_n, \mathbf{x}'_n)$
- Unknown: **Planar Homography (8 DoF)**
- Constraints:  $n \times$  projective transformation  $\mathbf{x}'_i = \mathbf{H}\mathbf{x}_i$
- Solutions ( $n \geq 4$ )
  - OpenCV `cv::getPerspectiveTransform()` and `cv::findHomography()`
  - c.f. More simplified transformations need less number of minimal correspondence.
    - Affine ( $n \geq 3$ ), similarity ( $n \geq 2$ ), Euclidean ( $n \geq 2$ )
- [Note] Planar homography can be decomposed as relative camera pose  $\mathbf{R}, \mathbf{t}$ .
  - OpenCV `cv::decomposeHomographyMat()`
  - c.f. However, the decomposition needs to know camera matrices.



# Perspective Correction using OpenCV

```
1. #include "opencv_all.hpp"

2. void MouseEventHandler(int event, int x, int y, int flags, void* param)
3. {
4.     ...
5.     std::vector<cv::Point> *points_src = (std::vector<cv::Point> *)param;
6.     points_src->push_back(cv::Point(x, y));
7.     ...
8. }

9. int main(void)
10.{ 
11.    cv::Size card_size(450, 250);

12.    // Prepare the rectified points
13.    std::vector<cv::Point> points_dst;
14.    points_dst.push_back(cv::Point(0, 0));
15.    points_dst.push_back(cv::Point(card_size.width, 0));
16.    points_dst.push_back(cv::Point(0, card_size.height));
17.    points_dst.push_back(cv::Point(card_size.width, card_size.height));

18.    // Load an image
19.    cv::Mat original = cv::imread("data/sunglok.jpg");
20.    if (original.empty()) return -1;

21.    // Get the matched points from a user's mouse
22.    std::vector<cv::Point> points_src;
23.    cv::namedWindow("3DV Tutorial: Perspective Correction");
24.    cv::setMouseCallback("3DV Tutorial: Perspective Correction", MouseEventHandler, &points_src);
25.    while (points_src.size() < 4)
26.    {
27.        ...
28.    }
29.    if (points_src.size() < 4) return -1;

30.    // Calculate planar homography and rectify perspective distortion
31.    cv::Mat H = cv::findHomography(points_src, points_dst);
32.    cv::Mat rectify;
33.    cv::warpPerspective(original, rectify, H, card_size);
```



# Perspective Correction using OpenCV



# Image Stitching using OpenCV



```
1. #include "opencv_all.hpp"  
2. int main(void)  
3. {  
4.     // Load two images (c.f. We assume that two images have the same size and type)  
5.     cv::Mat image1 = cv::imread("data/hill01.jpg");  
6.     cv::Mat image2 = cv::imread("data/hill02.jpg");  
7.     if (image1.empty() || image2.empty()) return -1;  
8.  
9.     // Retrieve matching points  
10.    cv::Mat gray1, gray2;  
11.    ...  
12.    cv::Ptr<cv::FeatureDetector> detector = cv::xfeatures2d::SURF::create(); // SURF features  
13.    std::vector<cv::KeyPoint> keypoint1, keypoint2;  
14.    detector->detect(gray1, keypoint1);  
15.    detector->detect(gray2, keypoint2);  
16.    cv::Ptr<cv::FeatureDetector> extractor = cv::xfeatures2d::SURF::create(); // SURF descriptors  
17.    cv::Mat descriptor1, descriptor2;  
18.    extractor->compute(gray1, keypoint1, descriptor1);  
19.    extractor->compute(gray2, keypoint2, descriptor2);  
20.    cv::FlannBasedMatcher matcher; // Approximate Nearest Neighbors (ANN) matcher  
21.    std::vector<cv::DMatch> match;  
22.    matcher.match(descriptor1, descriptor2, match);  
23.  
24.    // Calculate planar homography and merge them  
25.    std::vector<cv::Point2f> points1, points2;  
26.    for (size_t i = 0; i < match.size(); i++)  
27.    {  
28.        points1.push_back(keypoint1.at(match.at(i).queryIdx).pt);  
29.        points2.push_back(keypoint2.at(match.at(i).trainIdx).pt);  
30.    }  
31.    cv::Mat H = cv::findHomography(points2, points1, cv::RANSAC);  
32.    cv::Mat merge;  
33.    cv::warpPerspective(image2, merge, H, cv::Size(image1.cols * 2, image1.rows));  
34.    merge.colRange(0, image1.cols) = image1 * 1; // Copy  
35.    // Show the merged image  
36.    ...  
37.}
```



+



||



# Video Stabilization using OpenCV



```
1. #include <opencv_all.hpp>
2. int main(void)
3. {
4.     // Open an video and get the reference image and feature points
5.     cv::VideoCapture video;
6.     if (!video.open("data/traffic.avi")) return -1;
7.
8.     cv::Mat gray_ref;
9.     video >> gray_ref;
10.    if (gray_ref.empty())
11.    {
12.        video.release();
13.        return -1;
14.    }
15.    if (gray_ref.channels() > 1) cv::cvtColor(gray_ref, gray_ref, CV_RGB2GRAY);
16.
17.    std::vector<cv::Point2f> point_ref;
18.    cv::goodFeaturesToTrack(gray_ref, point_ref, 2000, 0.01, 10);
19.    if (point_ref.size() < 4)
20.    {
21.        video.release();
22.        return -1;
23.    }
```

A shaking CCTV video: data/traffic.avi





# Video Stabilization using OpenCV

```
22.     // Run and show video stabilization
23.     while (true)
24.     {
25.         // Grab an image from the video
26.         cv::Mat image, gray;
27.         video >> image;
28.         if (image.empty()) break;
29.         if (image.channels() > 1) cv::cvtColor(image, gray, CV_RGB2GRAY);
30.         else                         gray = image.clone();

31.         // Extract optical flow and calculate planar homography
32.         std::vector<cv::Point2f> point;
33.         std::vector<uchar> m_status;
34.         cv::Mat err, inlier_mask;
35.         cv::calcOpticalFlowPyrLK(gray_ref, gray, point_ref, point, m_status, err);
36.         cv::Mat H = cv::findHomography(point, point_ref, inlier_mask, cv::RANSAC);

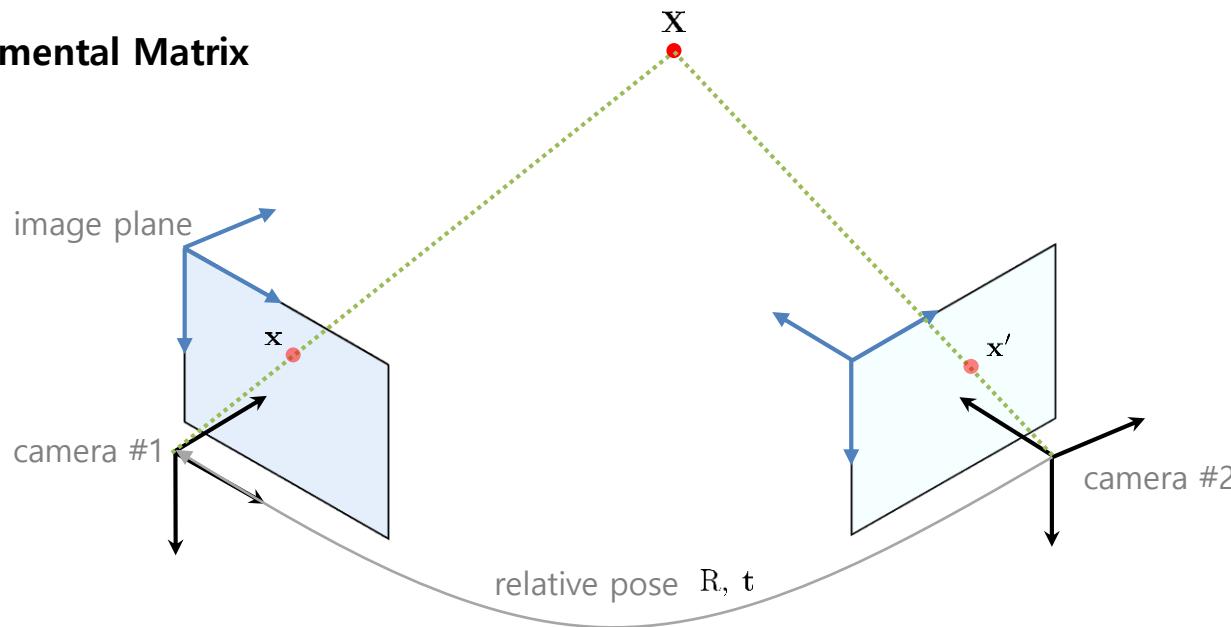
37.         // Synthesize a stabilized image
38.         cv::Mat warp;
39.         cv::warpPerspective(image, warp, H, cv::Size(image.cols, image.rows));

40.         // Show the original and rectified images together
41.         for (size_t i = 0; i < point_ref.size(); i++)
42.         {
43.             if (inlier_mask.at<uchar>(i) > 0) cv::line(image, point_ref[i], point[i], cv::Scalar(0, 0, 255));
44.             else cv::line(image, point_ref[i], point[i], cv::Scalar(0, 255, 0));
45.         }
46.         cv::hconcat(image, warp, image);
47.         cv::imshow("3DVT Tutorial: Video Stabilization", image);
48.         if (cv::waitKey(1) == 27) break; // 'ESC' key
49.     }

50.     video.release();
51.     return 0;
52. }
```

# General 2D-2D Geometry (Epipolar Geometry)

- Fundamental Matrix



**Epipolar Constraint**

$$x'^\top F x = 0$$

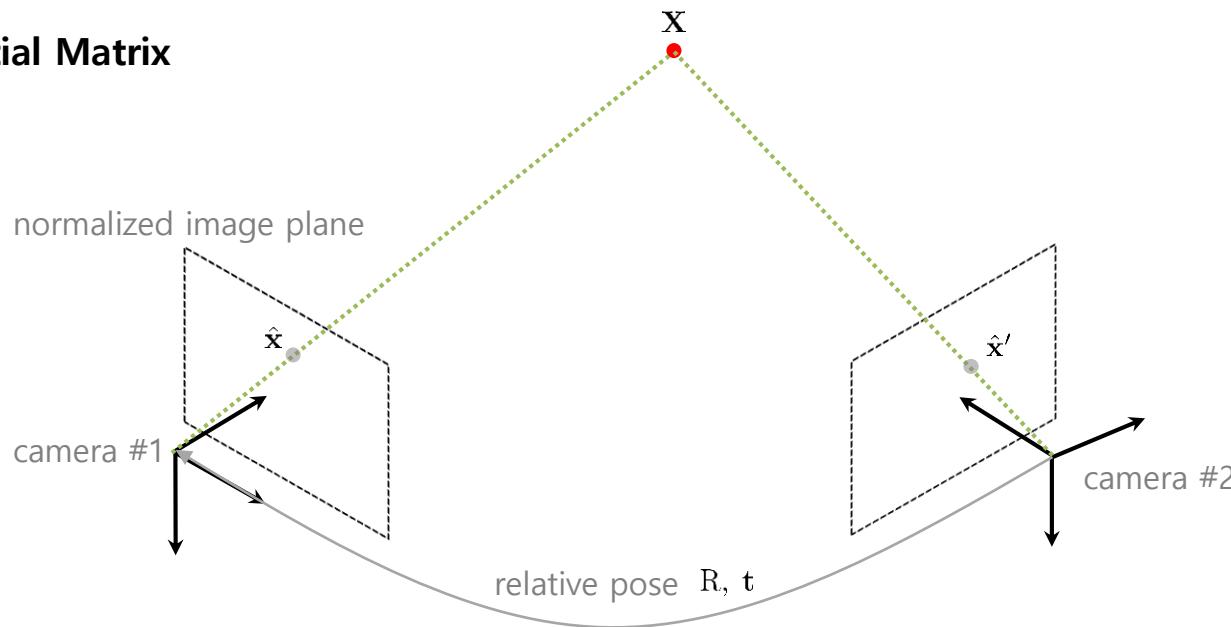
$$x = K \hat{x}$$

$$\hat{x}'^\top K'^\top F K \hat{x} = 0$$

$$\hat{x}'^\top E \hat{x} = 0 \quad (E = K'^\top F K) \quad \text{on the normalized image plane}$$

# General 2D-2D Geometry (Epipolar Geometry)

- Essential Matrix



**Epipolar Constraint**

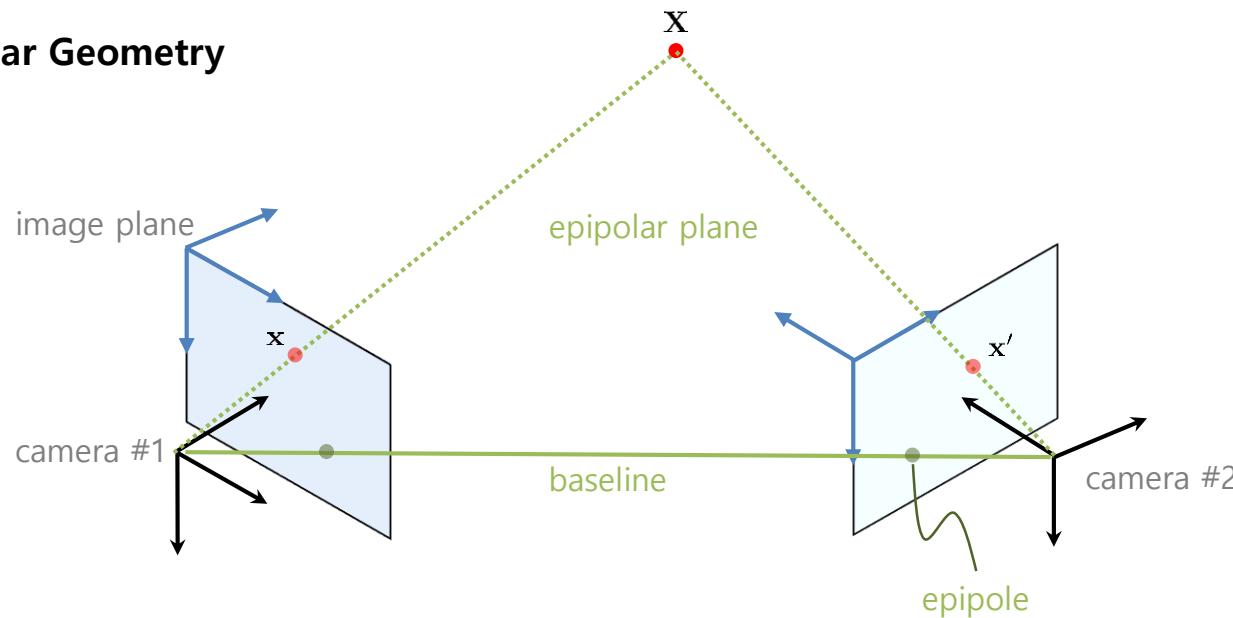
$$\hat{x}'^\top E \hat{x} = 0 \quad (E = [\mathbf{t}]_\times R)$$

c.f.  $[\mathbf{t}]_\times = \begin{bmatrix} 0 & -t_z & t_y \\ t_z & 0 & -t_x \\ -t_y & t_x & 0 \end{bmatrix}$

$$\begin{aligned} \hat{x}' \cdot (\mathbf{t} \times \hat{x}') &= 0 \\ \hat{x}' \cdot (\mathbf{t} \times \hat{x}') &= \hat{x}' \cdot \mathbf{t} \times R \hat{x} \\ \hat{x}' \cdot & \\ \mathbf{t} \times \hat{x}' &= \mathbf{t} \times R \hat{x} \\ \mathbf{t} \times & \\ \hat{x}' &= R \hat{x} + \mathbf{t} \end{aligned}$$

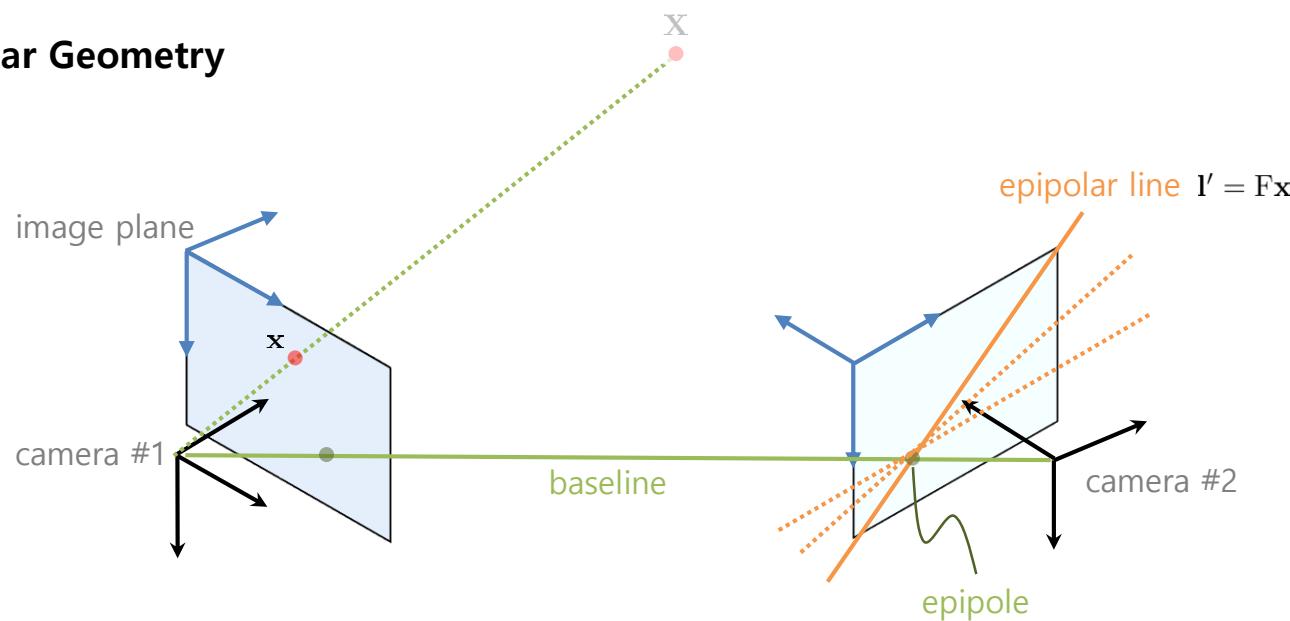
# General 2D-2D Geometry (Epipolar Geometry)

- Epipolar Geometry



# General 2D-2D Geometry (Epipolar Geometry)

- Epipolar Geometry

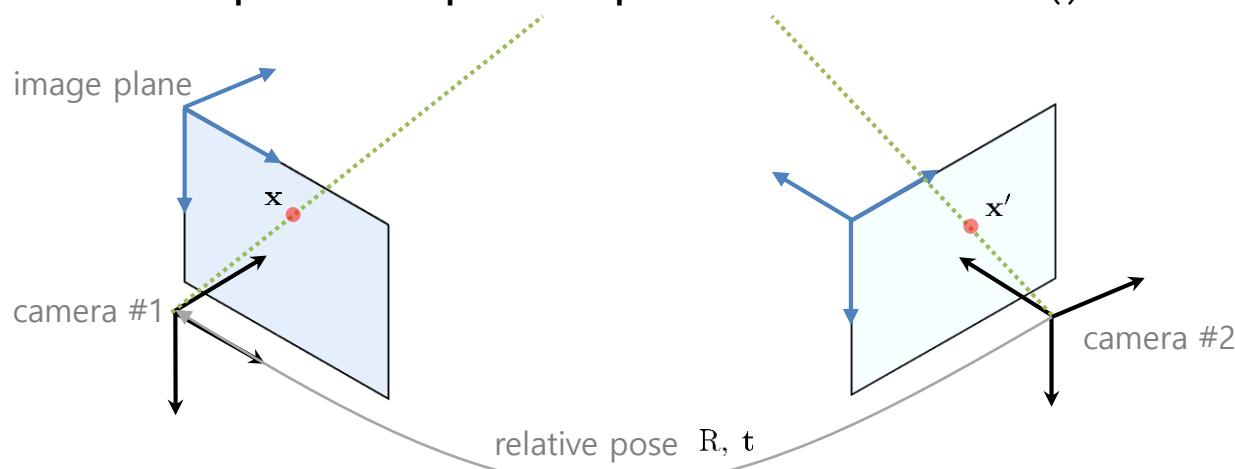
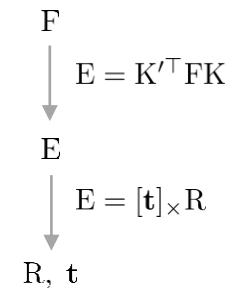


# General 2D-2D Geometry (Epipolar Geometry)

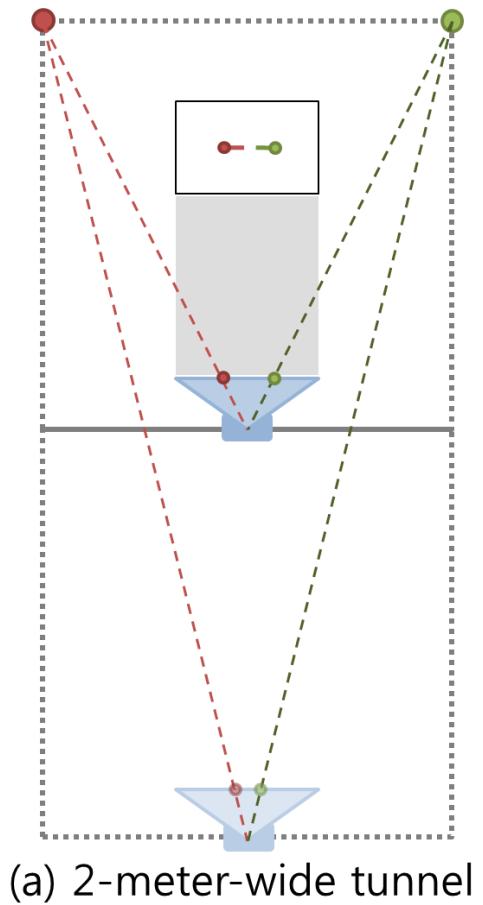
## ▪ Relative Camera Pose Estimation

- Known: Point correspondence  $(x_1, x'_1), \dots, (x_n, x'_n)$  and camera matrices  $K, K'$
- Unknown: **Rotation and translation** (5 DoF)
- Constraints:  $n \times$  epipolar constraint  $x'^\top F x = 0$
- Solutions (OpenCV)  
    • **Fundamental matrix:** 7/8-point algorithm
  - **Estimation:** `cv::findFundamentalMat()`  $\rightarrow 1$  solution
  - **Conversion to E:**  $E = K'^\top F K$
- **Solutions (OpenCV)**  
    • **Essential matrix:** 5-point algorithm
  - **Estimation:** `cv::findEssentialMat()`  $\rightarrow k$  solutions
  - **Decomposition:** `cv::decomposeEssentialMat()`  $\rightarrow 4$  solutions
  - **Decomposition with positive-depth check:** `cv::recoverPose()`  $\rightarrow 1$  solution

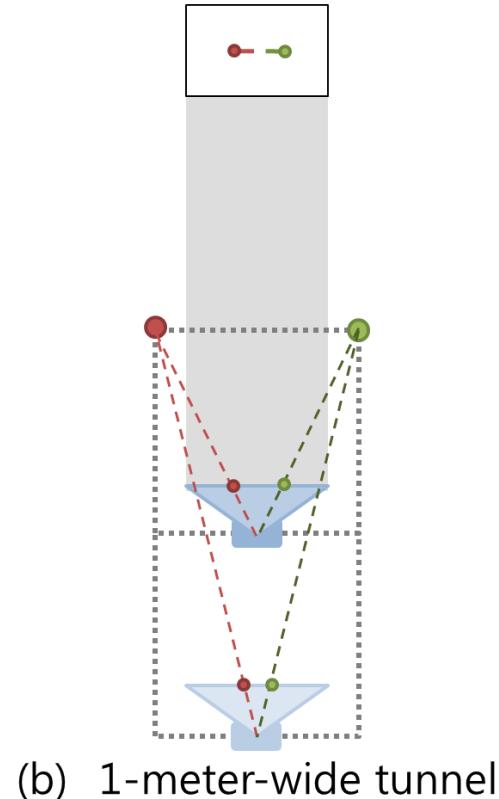
*up-to scale*  $\rightarrow$  scale ambiguity



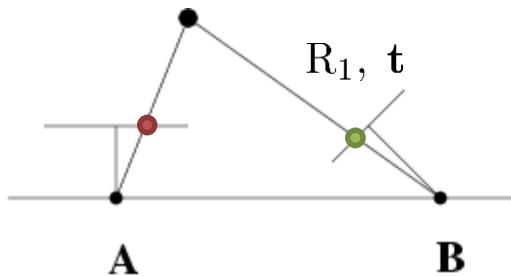
# Epipolar Geometry: Scale Ambiguity



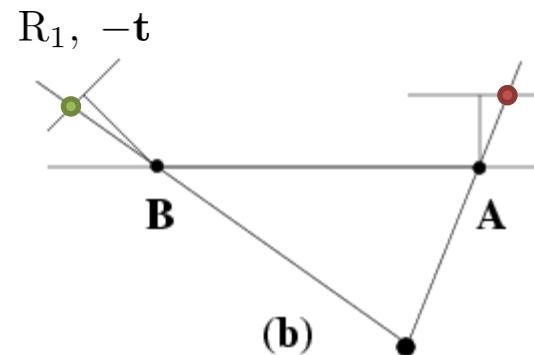
optical flow on an image



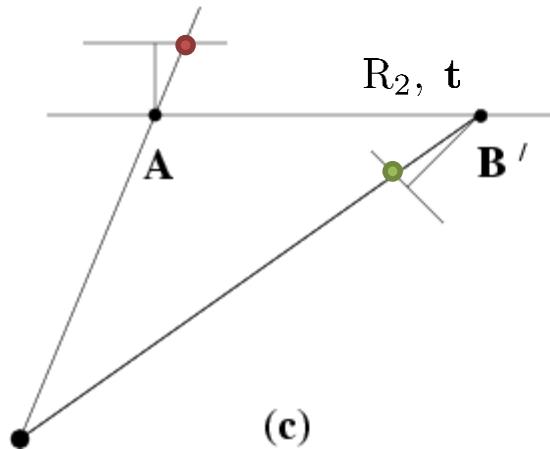
# Epipolar Geometry: Relative Pose Ambiguity



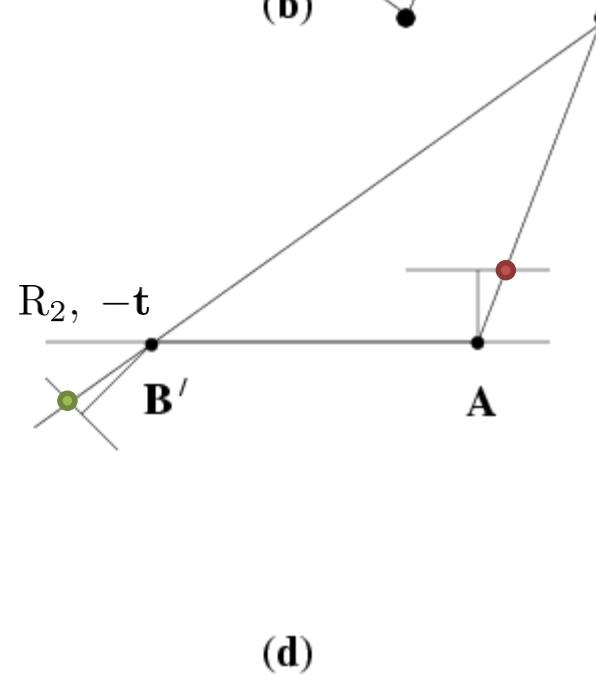
(a)



(b)



(c)



(d)

# General 2D-2D Geometry (Epipolar Geometry)

## ▪ Relative Camera Pose Estimation

- Known: Point correspondence  $(\mathbf{x}_1, \mathbf{x}'_1), \dots, (\mathbf{x}_n, \mathbf{x}'_n)$  and camera matrices  $\mathbf{K}, \mathbf{K}'$
- Unknown: Rotation and translation (5 DoF)
- Constraints:  $n \times$  epipolar constraint  $\mathbf{x}'^\top \mathbf{F} \mathbf{x} = 0$
- Solutions (OpenCV)

intrinsic & extrinsic  
camera parameters

- **Fundamental matrix:** 7/8-point algorithm (7 DoF)

$$\mathbf{F} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix} \quad \& \quad \text{rank}(\mathbf{F}) = 2 \quad \text{det}(\mathbf{F}) = 0$$

- **Estimation:** `cv::findFundamentalMat()`  $\rightarrow$  1 solution
- **Conversion to E:**  $\mathbf{E} = \mathbf{K}'^\top \mathbf{F} \mathbf{K}$
- **Degenerate cases:** No translation, correspondence from a single plane

extrinsic  
camera parameters

- **Essential matrix:** 5-point algorithm (5 DoF)

$$2\mathbf{E}\mathbf{E}^\top - \text{tr}(\mathbf{E}\mathbf{E}^\top)\mathbf{E} = 0$$

- **Estimation:** `cv::findEssentialMat()`  $\rightarrow k$  solutions
- **Decomposition:** `cv::decomposeEssentialMat()`  $\rightarrow$  4 solutions
- **Decomposition with positive-depth check:** `cv::recoverPose()`  $\rightarrow$  1 solution
- **Degenerate cases:** No translation ( $\because \mathbf{E} = [\mathbf{t}]_\times \mathbf{R}$ )

# General 2D-2D Geometry (Epipolar Geometry)

- Relative Camera Pose Estimation

- Known: Point correspondence  $(\mathbf{x}_1, \mathbf{x}'_1), \dots, (\mathbf{x}_n, \mathbf{x}'_n)$  and camera matrices  $\mathbf{K}, \mathbf{K}'$
- Unknown: Rotation and translation (5 DoF)
- Constraints:  $n \times$  epipolar constraint  $\mathbf{x}'^\top \mathbf{F} \mathbf{x} = 0$
- Solutions (OpenCV)

$$\mathbf{F} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix} \quad \text{rank}(\mathbf{F}) = 2 \quad \text{and} \quad \det(\mathbf{F}) = 0$$

intrinsic & extrinsic  
camera parameters

- Fundamental matrix:** 7/8-point algorithm (7 DoF)

- Estimation:** `cv::findFundamentalMat()` → 1 solution
- Conversion to E:**  $\mathbf{E} = \mathbf{K}'^\top \mathbf{F} \mathbf{K}$
- Degenerate cases:** No translation, correspondence from a single plane

extrinsic  
camera parameters

- Essential matrix:** 5-point algorithm (5 DoF)

$$2\mathbf{E}\mathbf{E}^\top - \text{tr}(\mathbf{E}\mathbf{E}^\top)\mathbf{E} = 0$$

- Estimation:** `cv::findEssentialMat()` →  $k$  solutions
- Decomposition:** `cv::decomposeEssentialMat()` → 4 solutions
- Decomposition with positive-depth check:** `cv::recoverPose()` → 1 solution
- Degenerate cases:** No translation ( $\because \mathbf{E} = [\mathbf{t}]_\times \mathbf{R}$ )

intrinsic & extrinsic  
camera parameters  
+ plane normal

- Planar homography:** 4-point algorithm (8 DoF)

- Estimation:** `cv::findHomography()` → 1 solutions
- Conversion to the normalized coordinate:**  $\hat{\mathbf{H}} = \mathbf{K}'^{-1} \mathbf{H} \mathbf{K}$
- Decomposition:** `cv::decomposeHomographyMat()` → 4 solutions
- Degenerate cases:** Correspondence not from a single plane

$$\mathbf{H} = \mathbf{K}' \hat{\mathbf{H}} \mathbf{K}^{-1}$$

$$\hat{\mathbf{H}} = \lambda \left( \mathbf{R} + \frac{1}{d} \mathbf{t} \mathbf{n}^\top \right) \quad \leftarrow \quad \hat{\mathbf{x}}' = \mathbf{R} \hat{\mathbf{x}} + \mathbf{t} \quad \text{and} \quad \frac{1}{d} \mathbf{t} \mathbf{n}^\top = 1 \quad (\because n_x \hat{x} + n_y \hat{y} + n_z \hat{z} + d = 0)$$



# Visual Odometry (Epipolar) using OpenCV

```
1. #include "opencv_all.hpp"  
2. int main(void)  
3. {  
4.     bool use_5pt = false;  
5.     double camera_focal = 718.8560;  
6.     cv::Point2d camera_center(607.1928, 185.2157);  
7.     // Open a file to write camera trajectory  
8.     FILE* camera_traj = fopen("visual_odometry_epipolar.xyz", "wt");  
9.     if (camera_traj == NULL) return -1;  
10.    // Open an video and get the initial image  
11.    cv::VideoCapture video;  
12.    if (!video.open("data/KITTI_00_L/%06d.png")) return -1;  
13.    cv::Mat gray_prev;  
14.    video >> gray_prev;  
15.    ...  
16.    // Run and store monocular visual odometry  
17.    cv::Mat camera_pose = cv::Mat::eye(4, 4, CV_64F);  
18.    while (true)  
19.    {  
20.        // Grab an image from the video  
21.        cv::Mat image, gray;  
22.        video >> image;  
23.        ...  
24.        // Extract optical flow  
25.        std::vector<cv::Point2f> point_prev, point;  
26.        cv::goodFeaturesToTrack(gray_prev, point_prev, 2000, 0.01, 10);  
27.        std::vector<uchar> m_status;  
28.        cv::Mat err;  
29.        cv::calcOpticalFlowPyrLK(gray_prev, gray, point_prev, point, m_status, err);  
30.        gray_prev = gray;
```



KITTI Odometry Dataset (#: 00, Left): data/KITTI\_00\_L/%06d.png

# Visual Odometry (Epipolar) using OpenCV



```
31.     // Calculate relative pose
32.     cv::Mat E, inlier_mask;
33.     if (use_5pt)
34.     {
35.         E = cv::findEssentialMat(point_prev, point, camera_focal, camera_center, cv::RANSAC, 0.99, 1,
36.                                 inlier_mask);
37.     }
38.     else
39.     {
40.         cv::Mat F = cv::findFundamentalMat(point_prev, point, cv::FM_RANSAC, 1, 0.99, inlier_mask);
41.         cv::Mat K = (cv::Mat<double>(3, 3) << camera_focal, 0, camera_center.x, 0, camera_focal,
42.                               camera_center.y, 0, 0, 1);
43.         E = K.t() * F * K;
44.     }
45.     cv::Mat R, t;
46.     int inlier_num = cv::recoverPose(E, point_prev, point, R, t, camera_focal, camera_center, inlier_mask);

47.     // Accumulate pose
48.     cv::Mat T = cv::Mat::eye(4, 4, R.type());
49.     T(cv::Range(0, 3), cv::Range(0, 3)) = R * 1.0;
50.     T(cv::Range(0, 3), cv::Range(3, 4)) = t * 1.0;
51.     camera_pose = camera_pose * T.inv();

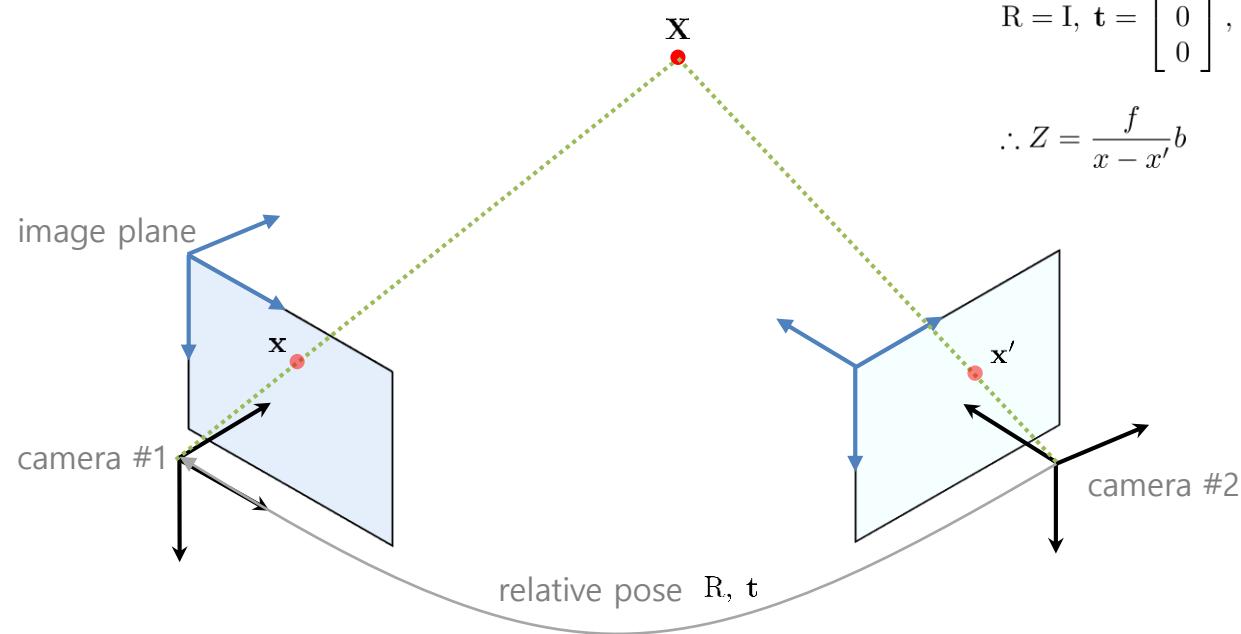
52.     // Show the image and write camera pose
53.     if (image.channels() < 3) cv::cvtColor(image, image, CV_GRAY2RGB);
54.     for (size_t i = 0; i < point_prev.size(); i++)
55.     {
56.         if (inlier_mask.at<uchar>(i) > 0) cv::line(image, point_prev[i], point[i], cv::Scalar(0, 0, 255));
57.         else cv::line(image, point_prev[i], point[i], cv::Scalar(0, 255, 0));
58.     }
59.     cv::imshow("3DVT Tutorial: Monocular Visual Odometry", image);
60.     fprintf(camera_traj, "%f %f %f\n", camera_pose.at<double>(0, 3), camera_pose.at<double>(1, 3),
61.             camera_pose.at<double>(2, 3));
62.     if (cv::waitKey(1) == 27) break; // 'ESC' key
63. }
64. video.release();
65. fclose(camera_traj);
66. return 0;
```

# General 2D-2D Geometry (Epipolar Geometry)

- **Relative Point Localization (Triangulation)**

- Known: Point correspondence, camera matrices, and relative pose
- Unknown: **Position of a 3D point** (3 DoF)
- Constraints:  $2 \times$  projection  $\mathbf{x} = \mathbf{K}[\mathbf{I} | \mathbf{0}] \mathbf{X}$ ,  $\mathbf{x}' = \mathbf{K}'[\mathbf{R} | \mathbf{t}] \mathbf{X}$
- Solution (OpenCV): `cv::triangulatePoints()`

**Special case) Stereo cameras**



$$\mathbf{R} = \mathbf{I}, \quad \mathbf{t} = \begin{bmatrix} -b \\ 0 \\ 0 \end{bmatrix}, \quad \mathbf{K} = \mathbf{K}' = \begin{bmatrix} f & 0 & c_x \\ 0 & f & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

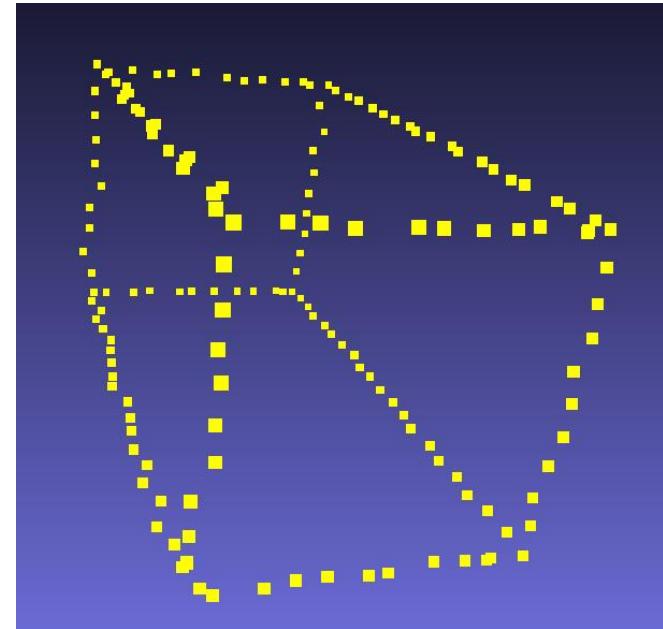
$$\therefore Z = \frac{f}{x - x'} b$$

# Triangulation using OpenCV

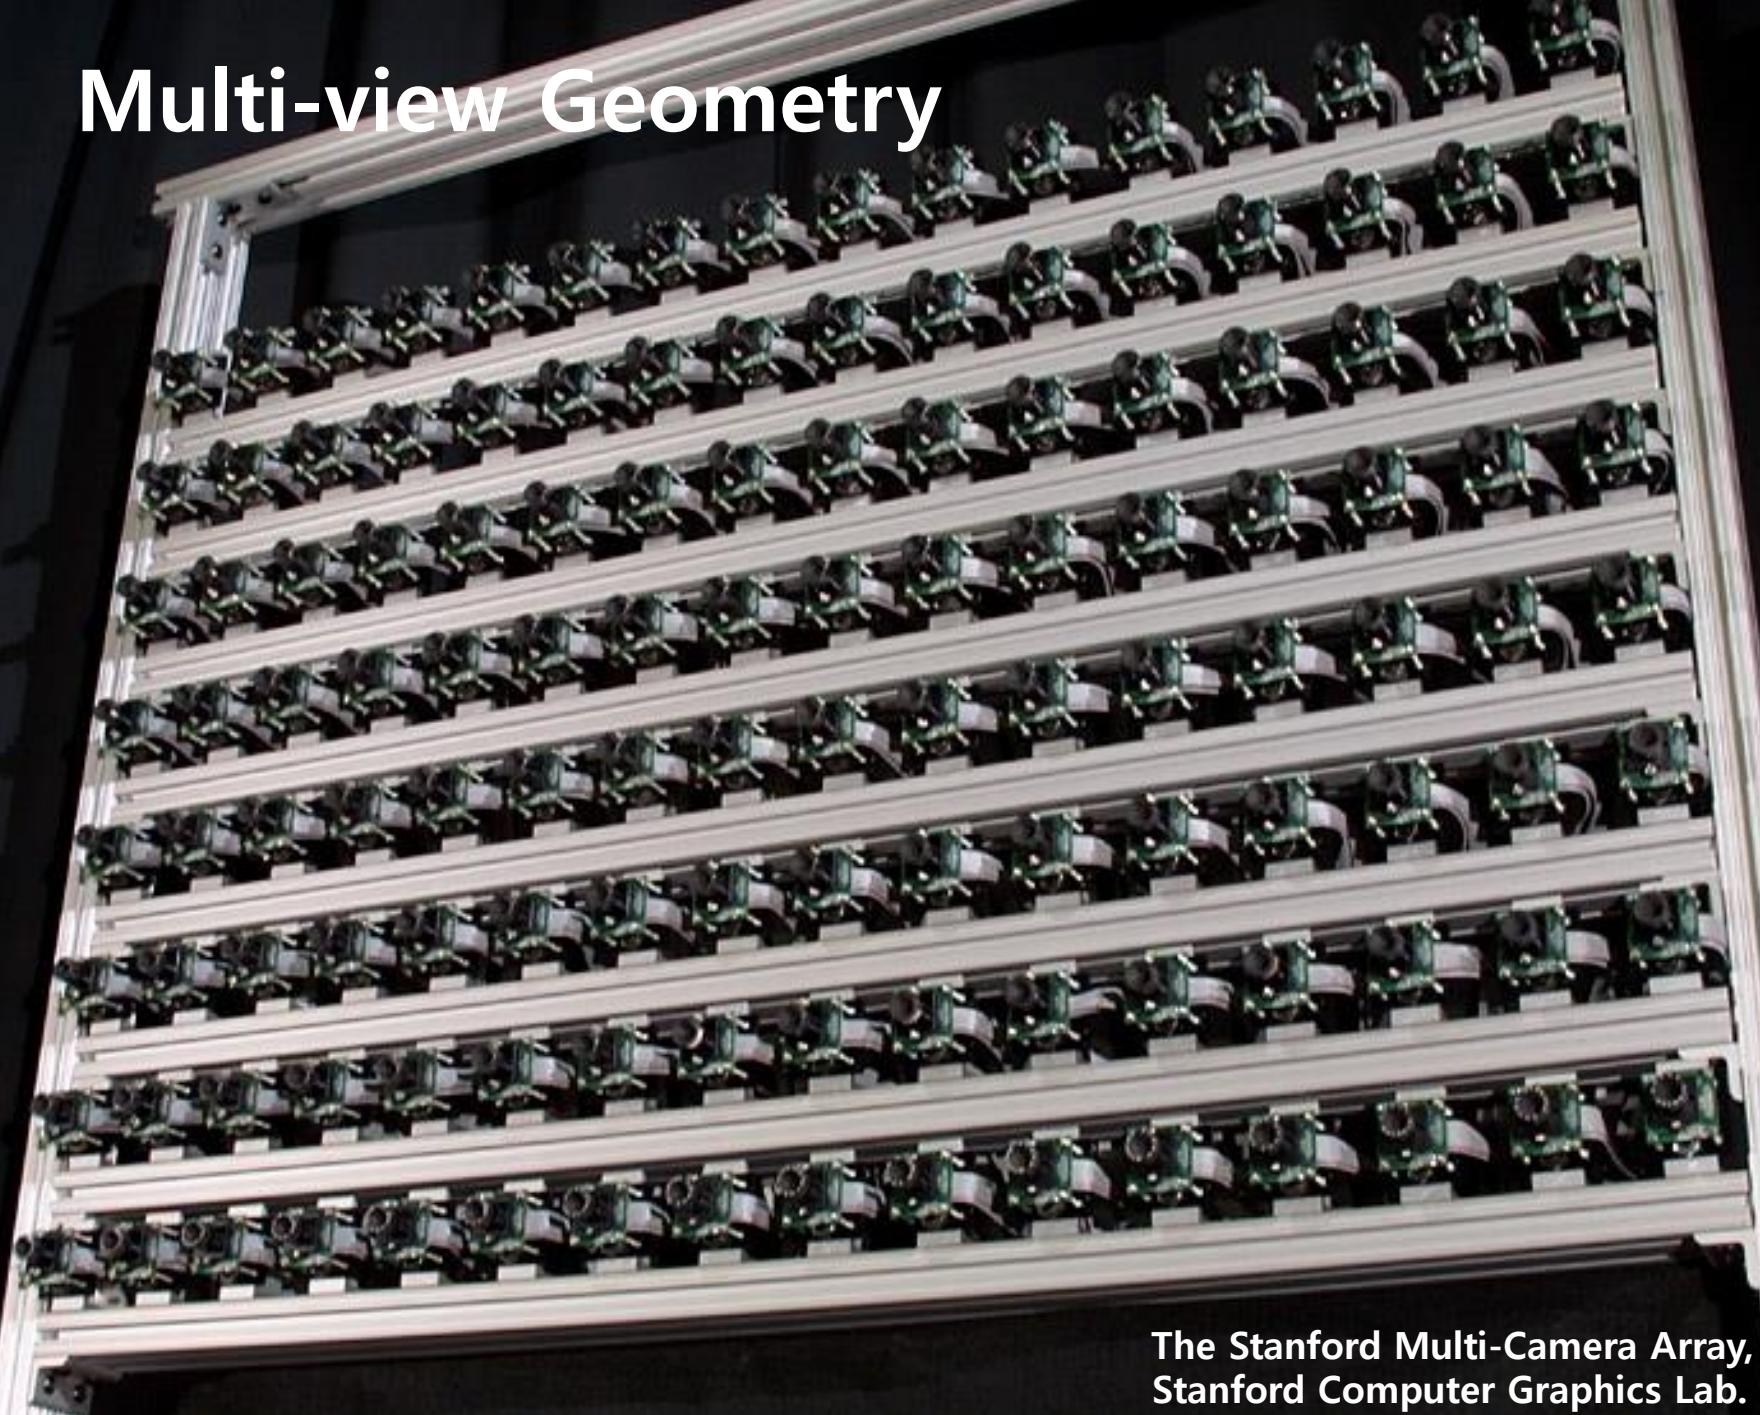
```
1. #include "opencv_all.hpp"
2.
3. int main(void)
4. {
5.     double camera_focal = 1000;
6.     cv::Point2d camera_center(320, 240);
7.
8.     // Load two views of 'box.xyz'
9.     std::vector<cv::Point2d> points0, points1;
10.    FILE* fin0 = fopen("image_formation0.xyz", "rt");
11.    FILE* fin1 = fopen("image_formation1.xyz", "rt");
12.    ...
13.
14.    // Estimate relative pose of two views
15.    cv::Mat F = cv::findFundamentalMat(points0, points1, cv::FM_8POINT);
16.    cv::Mat K = (cv::Mat<double>(3, 3) << camera_focal, ...);
17.    cv::Mat E = K.t() * F * K;
18.    cv::Mat R, t;
19.    cv::recoverPose(E, points0, points1, K, R, t);
20.
21.    // Reconstruct 3D points of 'box.xyz' (triangulation)
22.    cv::Mat P0 = K * cv::Mat::eye(3, 4, CV_64F);
23.    cv::Mat Rt, X;
24.    cv::hconcat(R, t, Rt);
25.    cv::Mat P1 = K * Rt;
26.    cv::triangulatePoints(P0, P1, points0, points1, X);
27.    X.row(0) = X.row(0) / X.row(3);
28.    X.row(1) = X.row(1) / X.row(3);
29.    X.row(2) = X.row(2) / X.row(3);
30.    X.row(3) = 1;
31.
32.    // Store the 3D points
33.    FILE* fout = fopen("triangulation.xyz", "wt");
34.    if (fout == NULL) return -1;
35.    for (int c = 0; c < X.cols; c++)
36.        fprintf(fout, "%f %f %f\n", X.at<double>(0, c), X.at<double>(1, c), X.at<double>(2, c));
37.    fclose(fout);
38.    return 0;
39.}
```



Result: triangulation.xyz



# Multi-view Geometry



The Stanford Multi-Camera Array,  
Stanford Computer Graphics Lab.

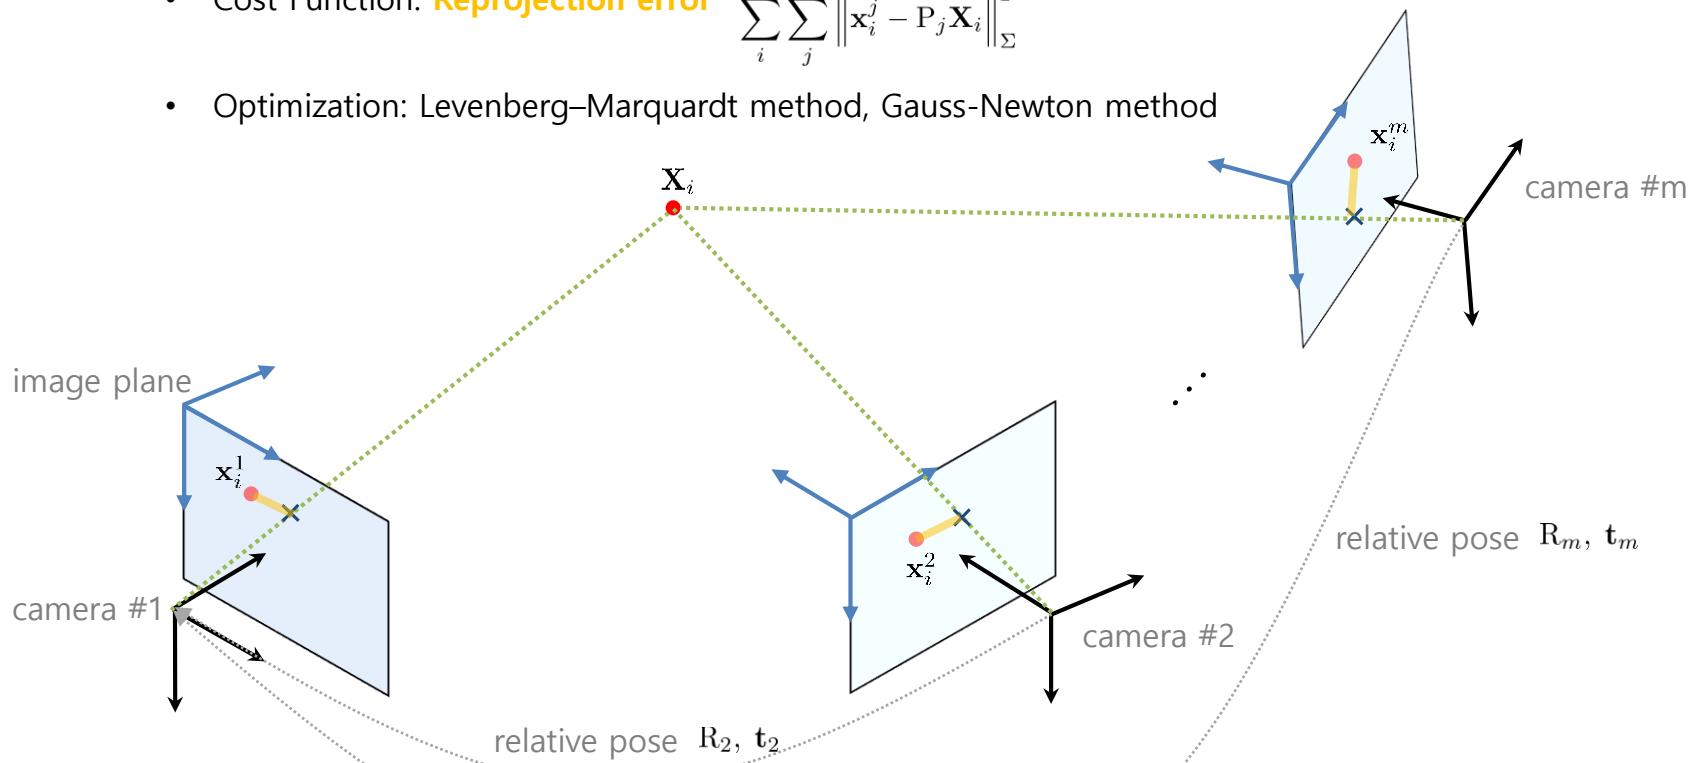
# Multi-view Geometry



The Stanford Multi-Camera Array,  
Stanford Computer Graphics Lab.

# Bundle Adjustment

- **Bundle Adjustment** c.f. initial values
  - Known: Point correspondence, camera matrices, position of 3D points, and each camera's relative pose
  - Unknown: **Position of 3D points and each camera's relative pose (6n + 3m DoF)**
  - Constraints:  $n \times m \times$  projection  $\mathbf{x}_i^j = \mathbf{P}_j \mathbf{X}_i = \mathbf{K}_j [\mathbf{R}_j \mid \mathbf{t}_j] \mathbf{X}_i$
  - Solution: **Non-linear least-square optimization**
    - Cost Function: **Reprojection error**  $\sum_i^n \sum_j^m \left\| \mathbf{x}_i^j - \mathbf{P}_j \mathbf{X}_i \right\|_{\Sigma}^2$
    - Optimization: Levenberg–Marquardt method, Gauss–Newton method



# Bundle Adjustment

- **Bundle Adjustment** c.f. initial values
  - Known: Point correspondence, camera matrices, position of 3D points, and each camera's relative pose
  - Unknown: **Position of 3D points and each camera's relative pose** ( $6n + 3m$  DoF)
  - Constraints:  $n \times m \times$  projection  $\mathbf{x}_i^j = \mathbf{P}_j \mathbf{X}_i = \mathbf{K}_j [\mathbf{R}_j \mid \mathbf{t}_j] \mathbf{X}_i$
  - Solution: **Non-linear least-square optimization**
    - Cost Function: **Reprojection error** 
$$\sum_i^n \sum_j^m \left\| \mathbf{x}_i^j - \mathbf{P}_j \mathbf{X}_i \right\|_{\Sigma}^2$$
    - Optimization: Levenberg–Marquardt method, Gauss–Newton method
- **Bundle Adjustment and Optimization Tools**
  - OpenCV: No function (but [cvSba](#) is available as a [sba](#) wrapper.)
  - Bundle adjustment: [sba](#) (Sparse Bundle Adjustment), [SSBA](#) (Simple Sparse Bundle Adjustment), [pba](#) (Multicore Bundle Adjustment)
    - Incremental structure-from-motion: [Bundler](#), [CM3PMVS](#), [VisualSfM](#)
  - Optimization: [g2o](#) (General Graph Optimization), [iSAM](#) (Incremental Smoothing and Mapping), [GTSAM](#), [Ceres Solver](#) (Google)

# Bundle Adjustment using cvSBA



```
1. #include "opencv_all.hpp"
2. #include "cvSBA.h"

3. int main(void)
4. {
5.     double camera_focal = 1000;
6.     cv::Point2d camera_center(320, 240);
7.     int n_views = 5;

8.     // Load multiple views of 'box.xyz'
9.     std::vector<std::vector<cv::Point2d>> xs;
10.    for (int i = 0; i < n_views; i++)
11.    {
12.        FILE* fin = fopen(cv::format("image_formation%d.xyz", i).c_str(), "rt");
13.        ...
14.    }

15.    // Assume that all feature points are visible
16.    std::vector<int> visible_all(xs.front().size(), 1);
17.    std::vector<std::vector<int>> visibility(n_views, visible_all);

18.    // Prepare each camera projection matrix
19.    std::vector<cv::Mat> Ks, dist_coeffs, Rs, ts;
20.    cv::Mat K = (cv::Mat_<double>(3, 3) << camera_focal, 0, camera_center.x, ...);
21.    Ks.resize(n_views, K);                                // K for all cameras
22.    dist_coeffs.resize(n_views, cv::Mat::zeros(5, 1, CV_64F)); // dist_coeff for all cameras
23.    Rs.push_back(cv::Mat::eye(3, 3, CV_64F));           // R for the first camera (index: 0)
24.    ts.push_back(cv::Mat::zeros(3, 1, CV_64F));          // t for the first camera (index: 0)

25.    // Estimate relative pose of the initial two views (epipolar geometry)
26.    cv::Mat F = cv::findFundamentalMat(xs[0], xs[1], cv::FM_8POINT);
27.    cv::Mat E = K.t() * F * K;
28.    cv::Mat R, t;
29.    cv::recoverPose(E, xs[0], xs[1], K, R, t);
30.    Rs.push_back(R);                                    // R for the second camera
31.    ts.push_back(t);                                    // t for the second camera
```



# Bundle Adjustment using cvSba

```
32. // Reconstruct the initial 3D points of 'box.xyz' (triangulation)
33. cv::Mat P0 = K * cv::Mat::eye(3, 4, CV_64F);
34. cv::Mat Rt, X;
35. cv::hconcat(R, t, Rt);
36. cv::Mat P1 = K * Rt;
37. cv::triangulatePoints(P0, P1, xs[0], xs[1], X);
38. std::vector<cv::Point3d> Xs;
39. X.row(0) = X.row(0) / X.row(3);
40. X.row(1) = X.row(1) / X.row(3);
41. X.row(2) = X.row(2) / X.row(3);
42. X.row(3) = 1;
43. for (int c = 0; c < X.cols; c++)
44.     Xs.push_back(cv::Point3d(X.at<double>(0, c), X.at<double>(1, c), X.at<double>(2, c)));

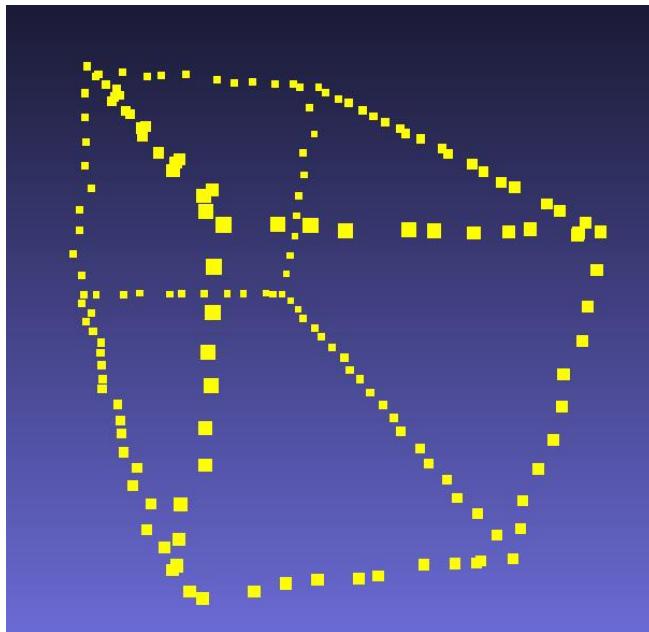
45. // Estimate the initial relative pose of other views (PnP)
46. for (int i = 2; i < n_views; i++)
47. {
48.     cv::Mat rvec;
49.     cv::solvePnP(Xs, xs[i], Ks[i], dist_coeffs[i], rvec, t);
50.     cv::Rodrigues(rvec, R);
51.     Rs.push_back(R);                                // R for the third and other cameras
52.     ts.push_back(t);                                // t for the third and other cameras
53. }

54. // Optimize camera pose and 3D points (bundle adjustment)
55. try
56. {
57.     cvSba::Sba sba;
58.     cvSba::Sba::Params param;
59.     param.type = cvSba::Sba::MOTIONSTRUCTURE;
60.     param.fixedIntrinsics = 5;
61.     param.fixedDistortion = 5;
62.     param.verbose = true;
63.     sba.setParams(param);
64.     double error = sba.run(Xs, xs, visibility, Ks, Rs, ts, dist_coeffs);
65. }
66. catch (cv::Exception) { }
67. // Store the 3D points
```

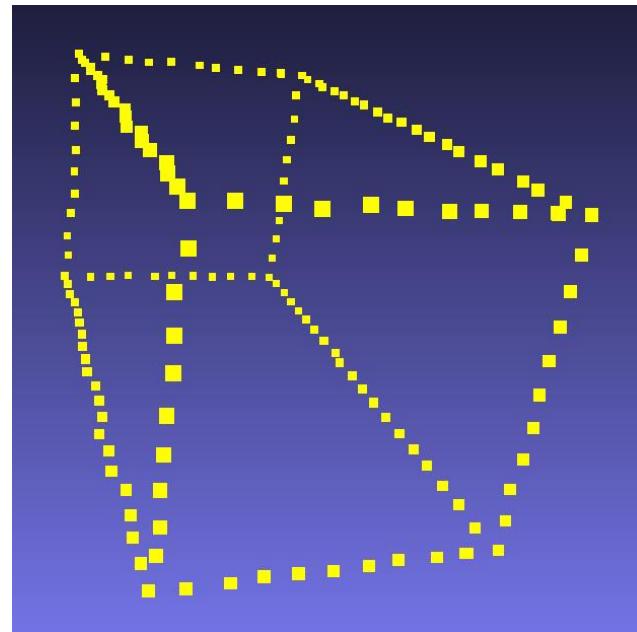
# Bundle Adjustment using cvSBA



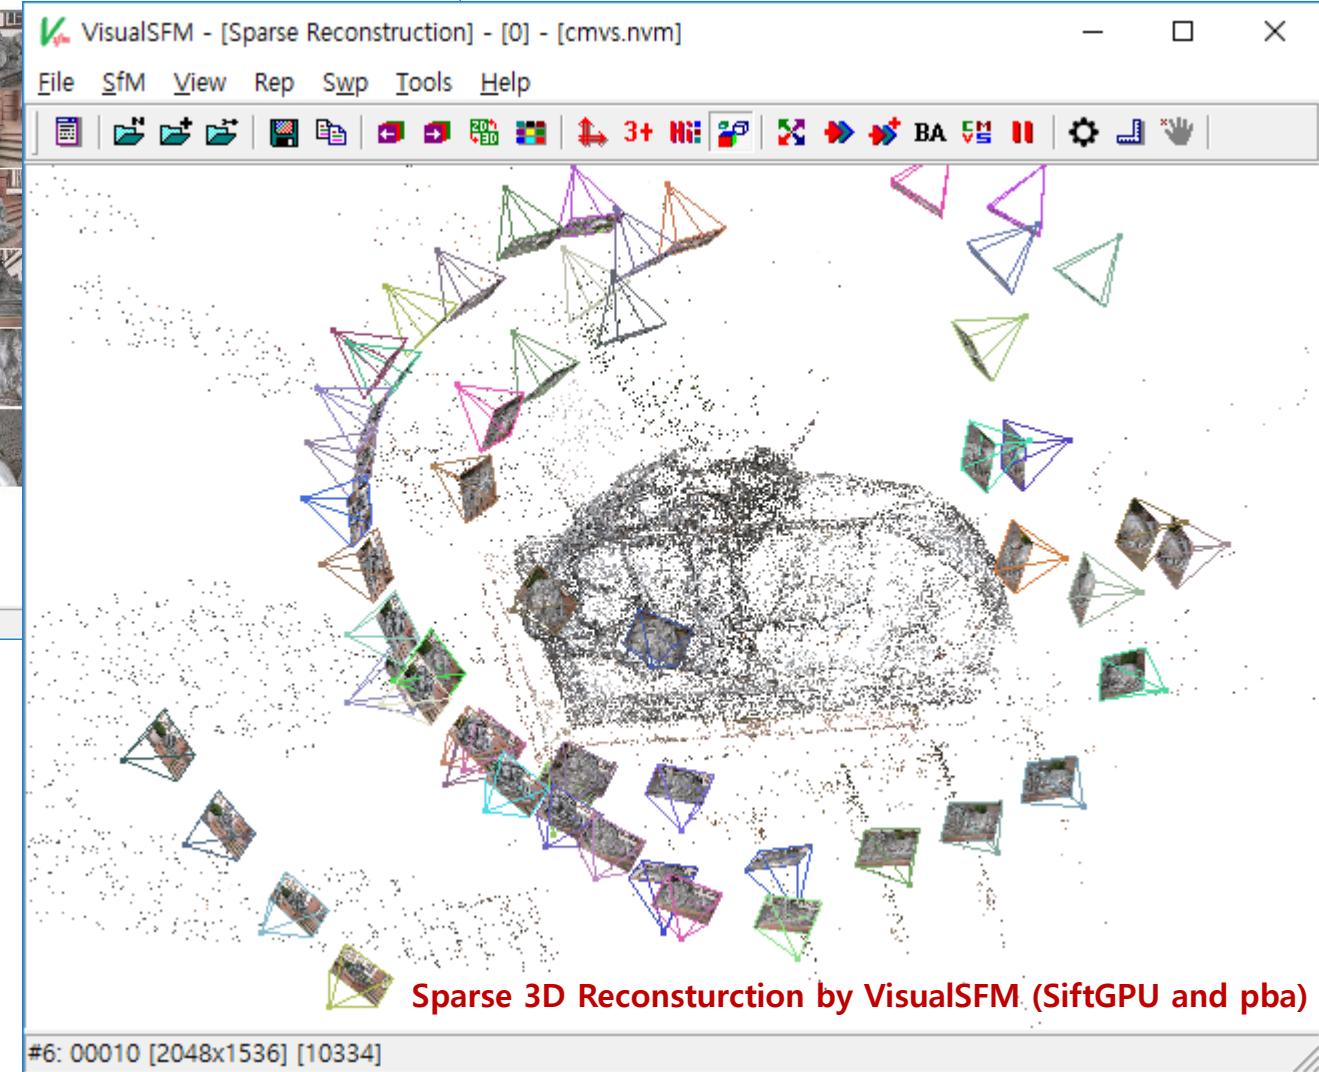
Result: triangulation.xyz



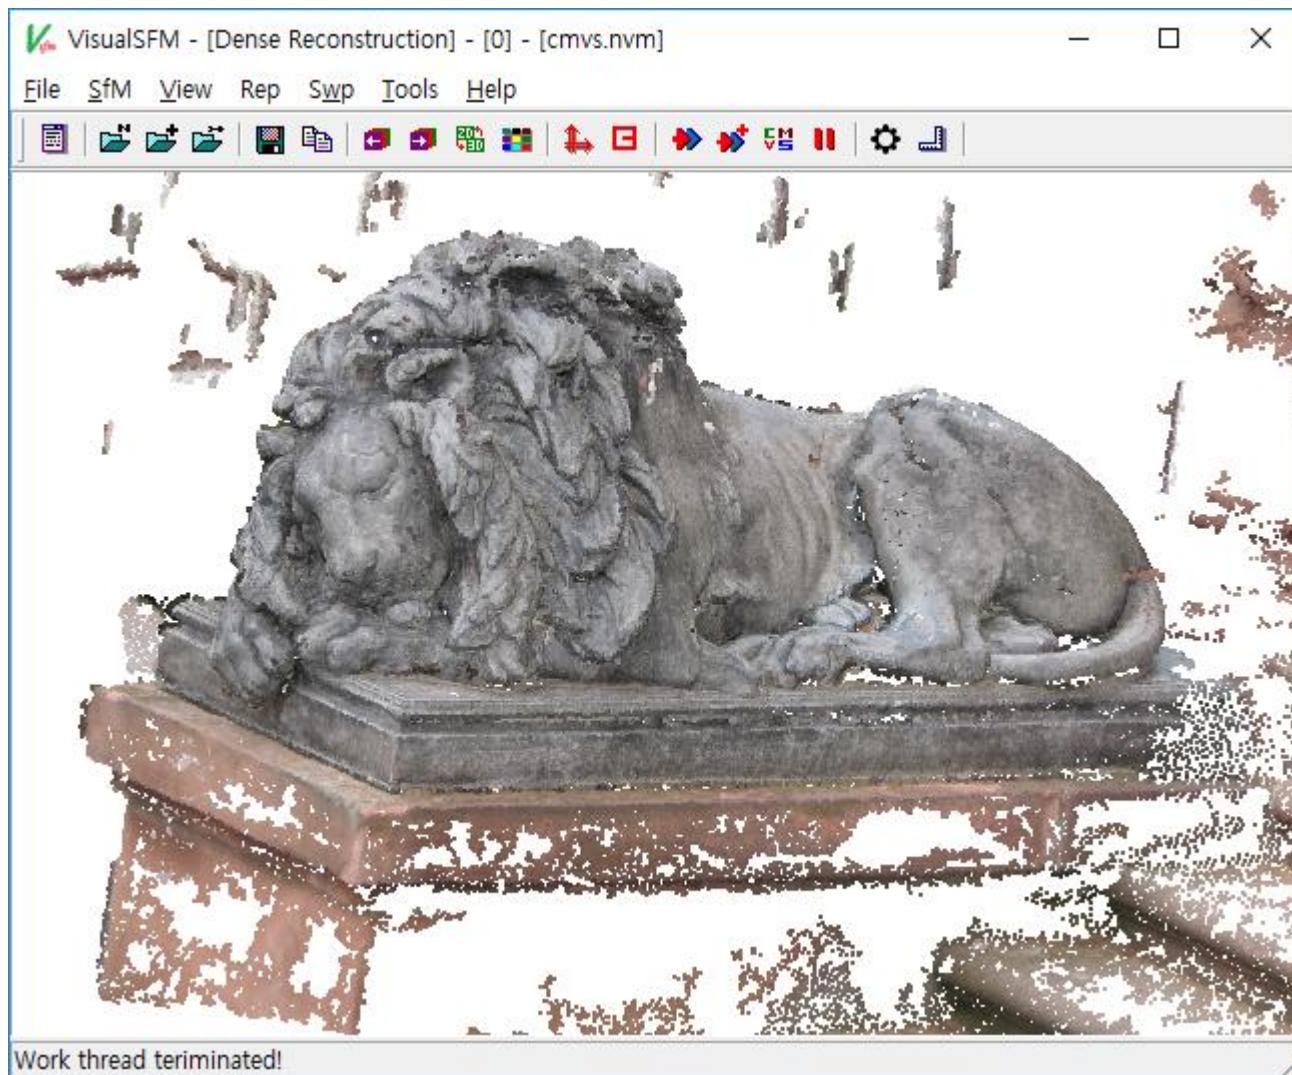
Result: bundle\_adjustment.xyz



# 3D Reconstruction using VisualSfM



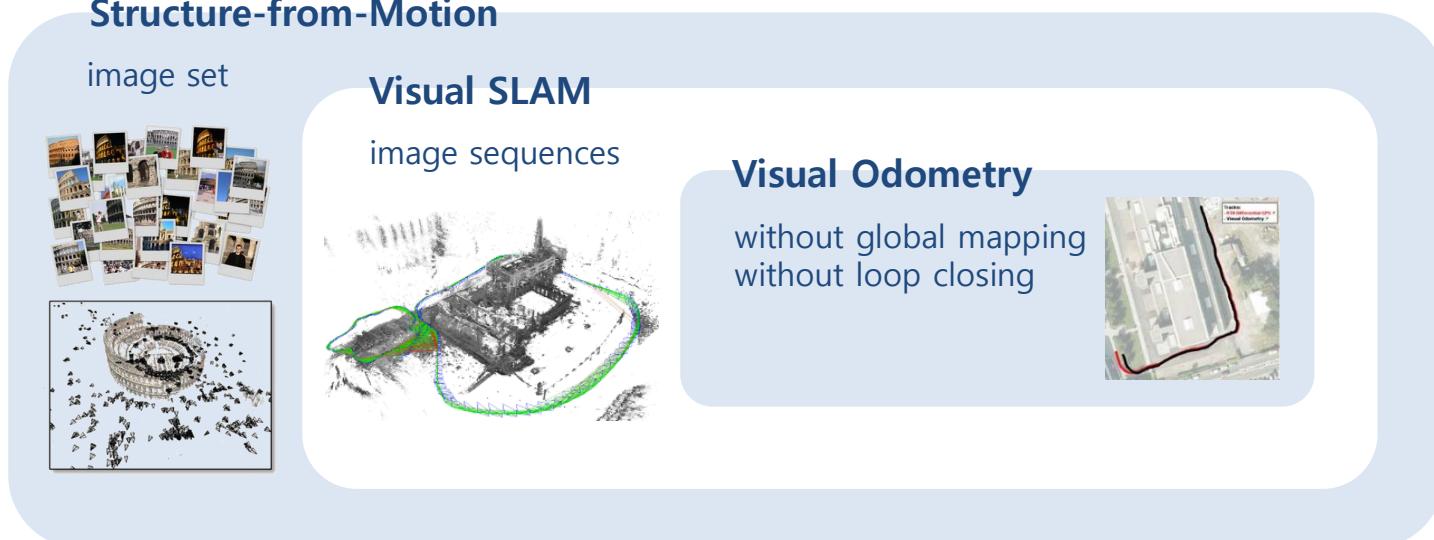
# 3D Reconstruction using VisualSfM



Dense 3D Reconstruction by VisualSfM (CMVS and PMVS2)

# Multiple-view Geometry: Applications

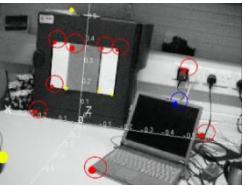
- **Structure-from-Motion (SfM)** → 3D Reconstruction, Photo Browsing
  - [Photo Toursim](#), [PhotoSynth](#), [PhotoCity](#)
  - [Building Rome in a Day](#), [Building Rome on a Cloudless Day](#)
- **Visual SLAM** → Augmented Reality, Navigation (Mapping and Localization)
  - [PTAM](#) (Parallel Tracking and Mapping), [ORB-SLAM](#), [LSD-SLAM](#), [DTAM](#) (Dense Tracking and Mapping)
    - [ORB-SLAM for Windows](#)
- **Visual Odometry** → Navigation (Localization)
  - [LIBVISO2](#) (C++ Library for Visual Odometry 2), [SVO](#) (Semi-direct Monocular Visual Odometry)



# Multiple-view Geometry: Visual SLAM

- **15 Years of Visual SLAM**

Monocular Camera



**MonoSLAM (2003)**: sparse features, filtering

SLAM in Robotics

V.S.

SfM in Computer Vision



**PTAM (2007)**: less sparse features, optimization

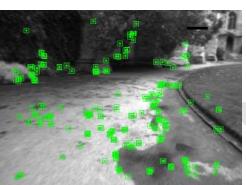
RGB-D Camera



**KinectFusion (2011)**: dense volumetric



**DTAM (2011)**: dense volumetric



**ORB-SLAM (2014)** v.s.  
less sparse features



**LSD-SLAM (2014)**  
direct, semi-dense

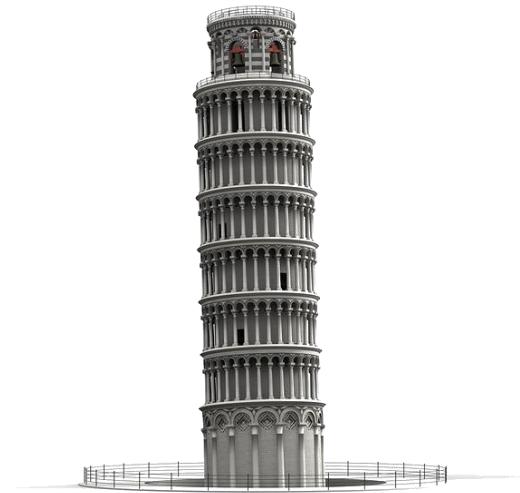


**ElasticFusion (2015)**  
dense surfels



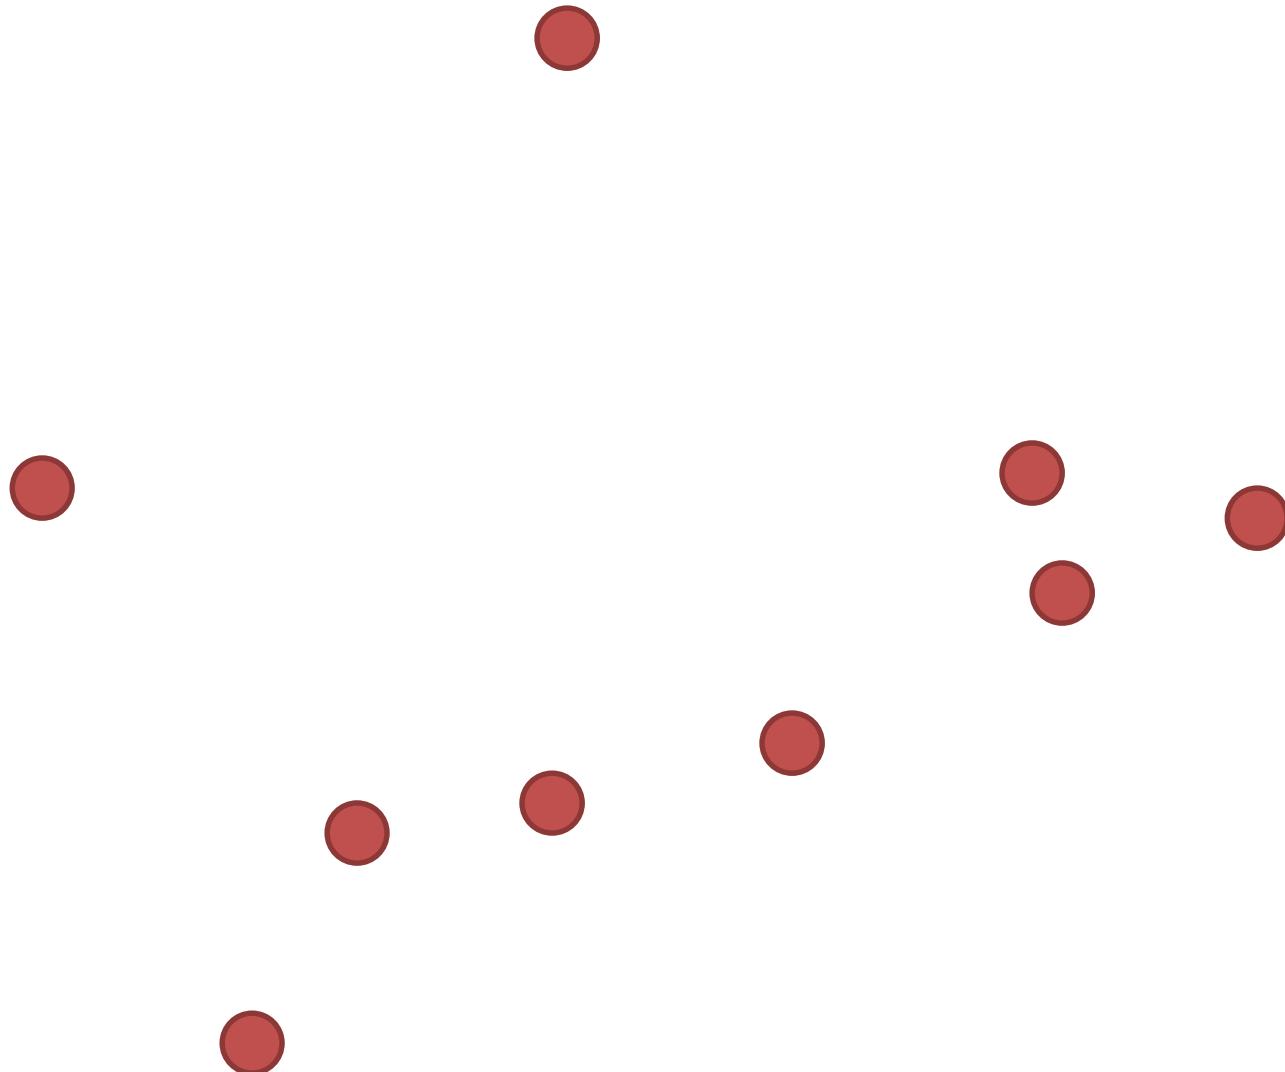
**DynamicFusion (2015)**  
dynamic dense

# Correspondence Problem

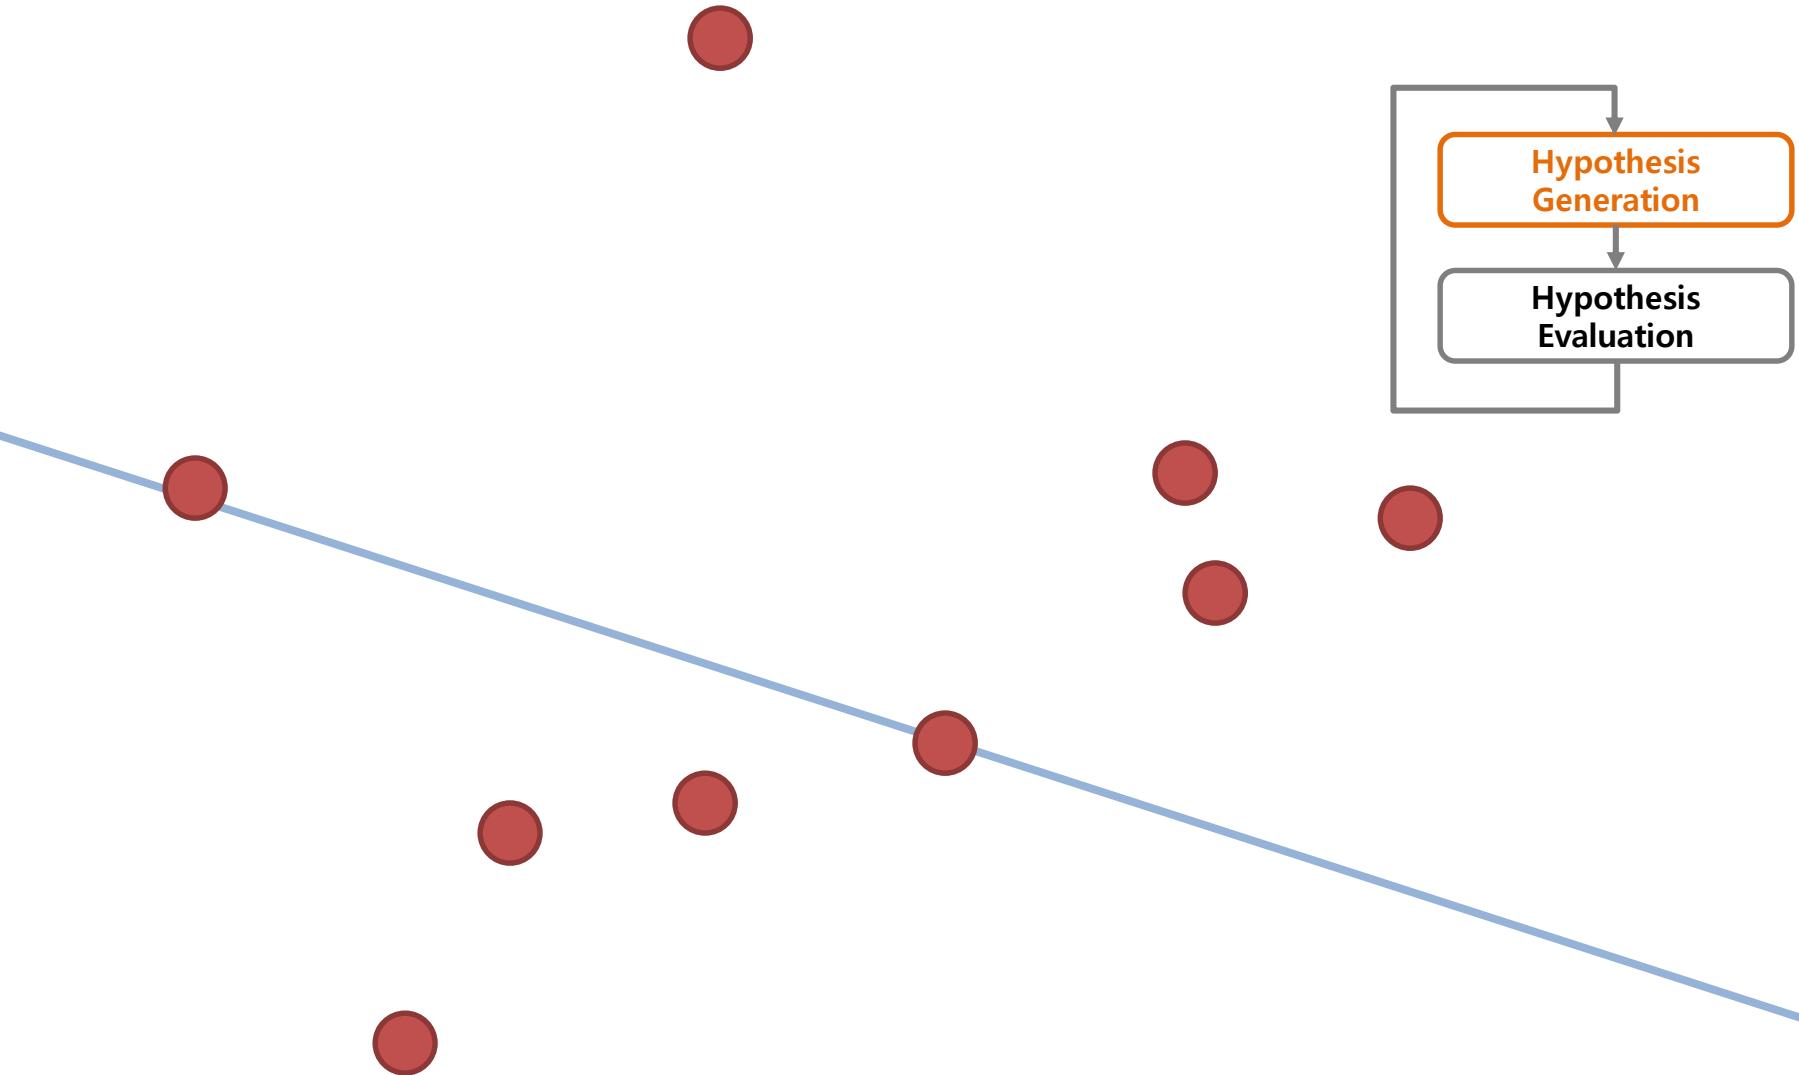


Images are from [pixabay](#).

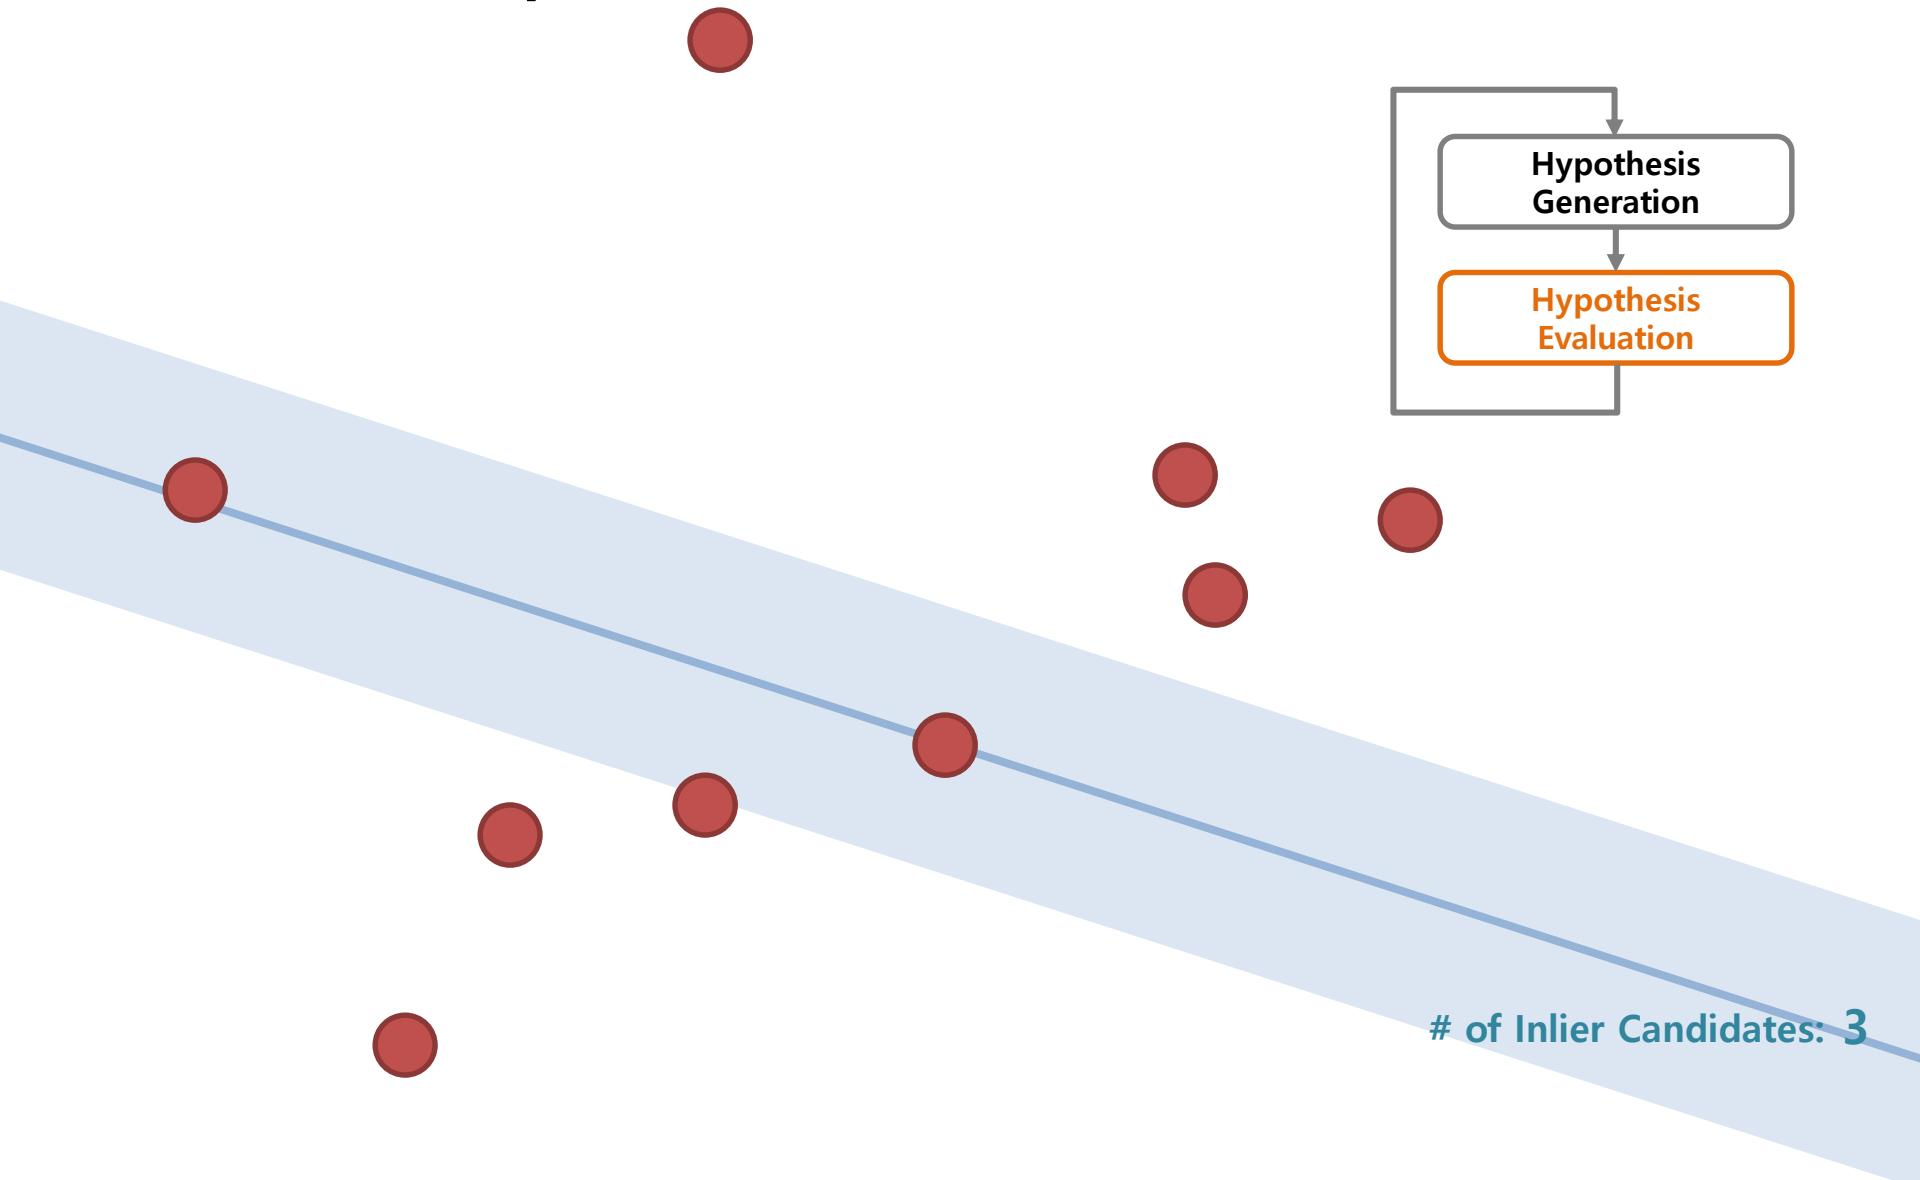
# Random Sample Consensus (RANSAC)



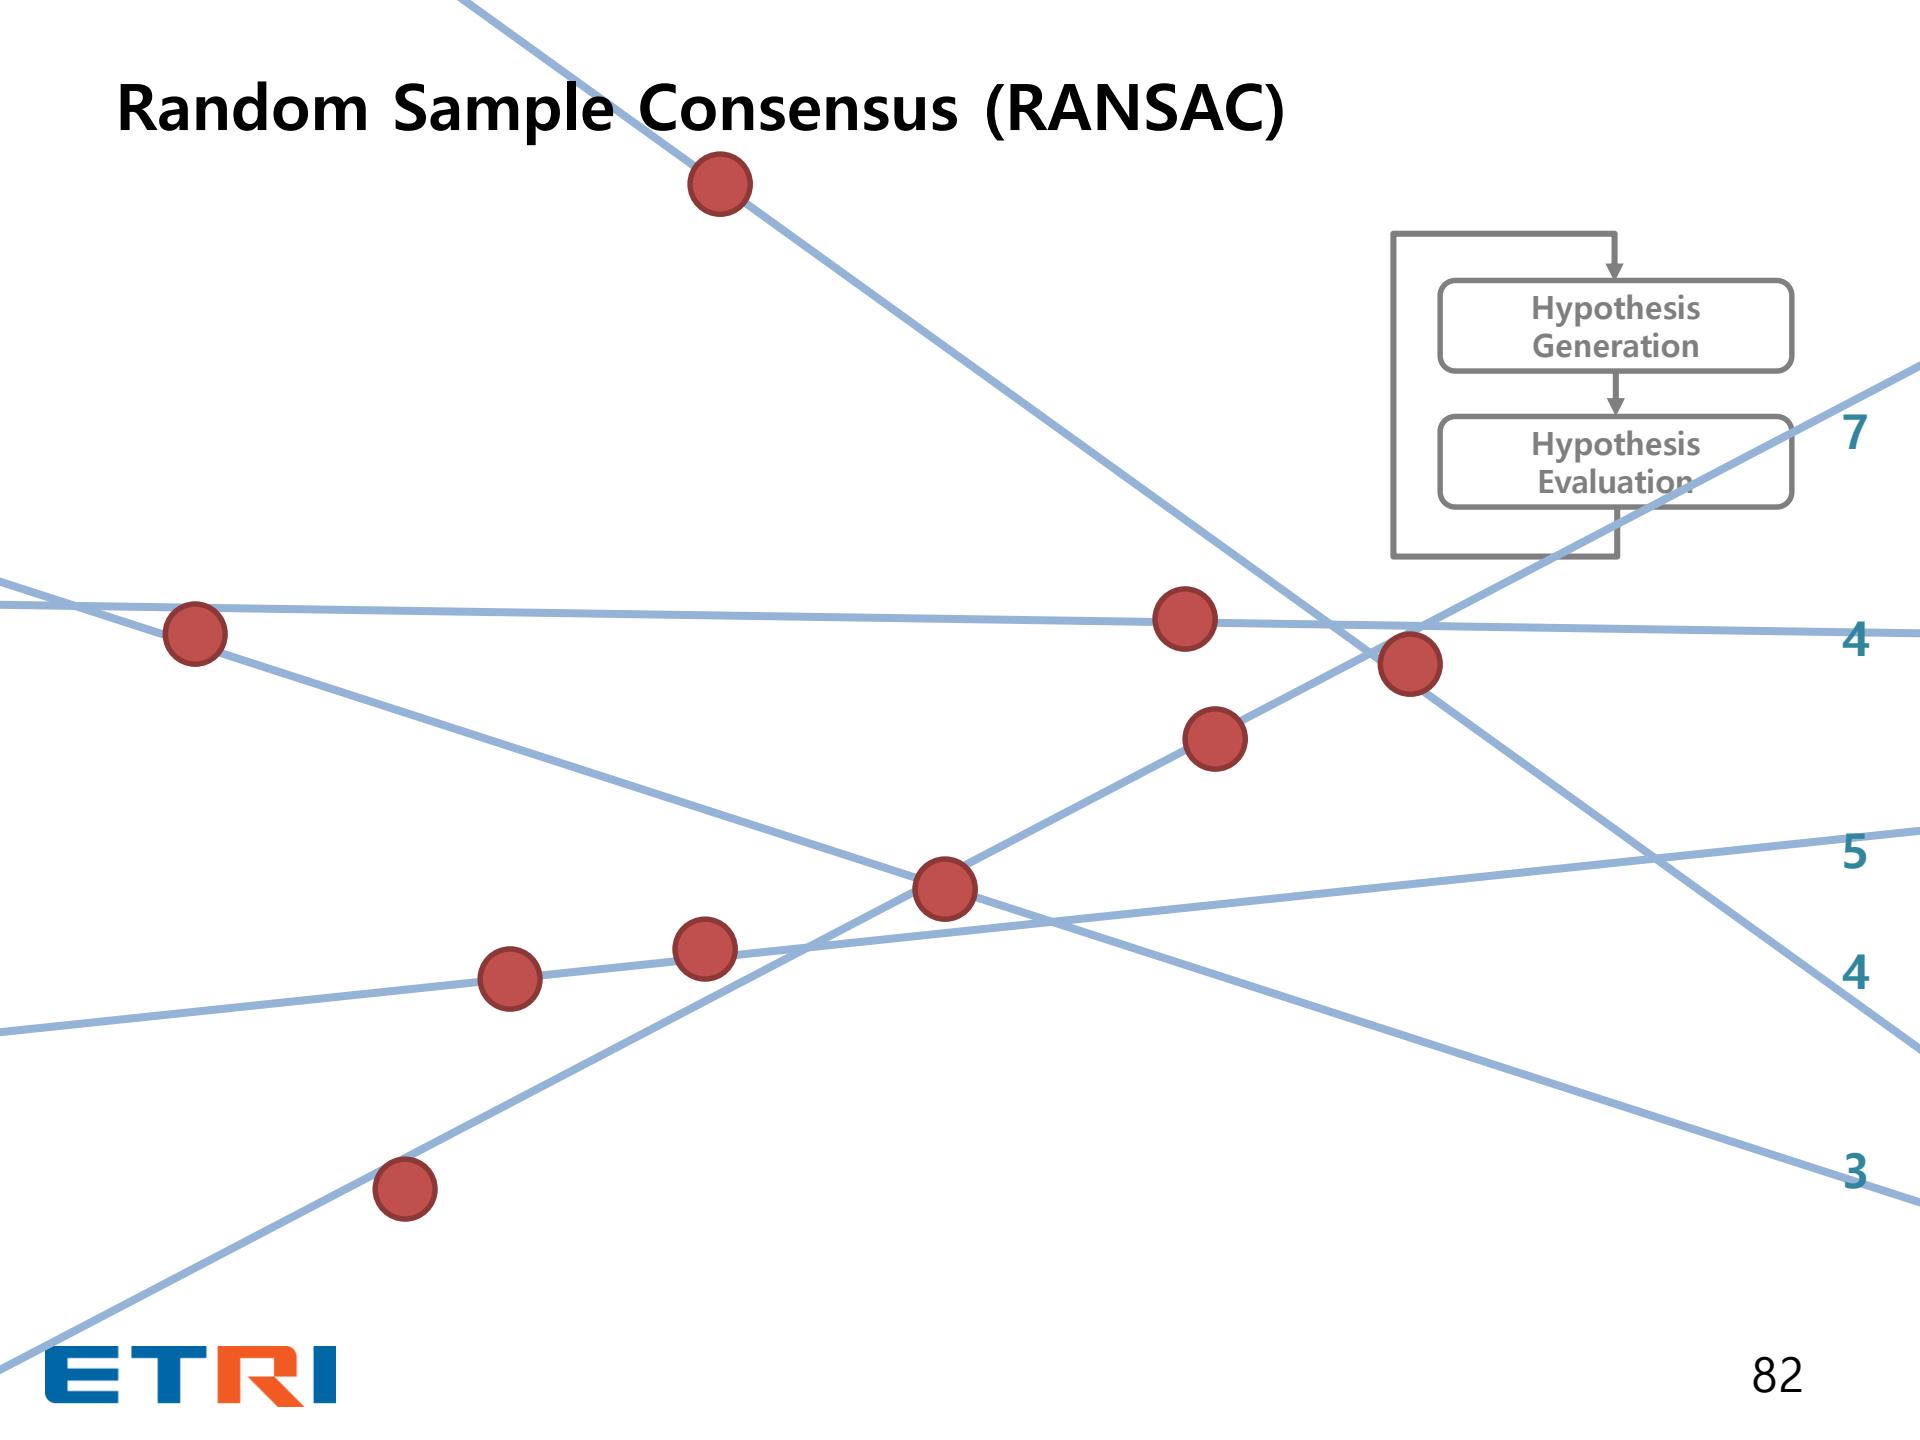
# Random Sample Consensus (RANSAC)



# Random Sample Consensus (RANSAC)



# Random Sample Consensus (RANSAC)

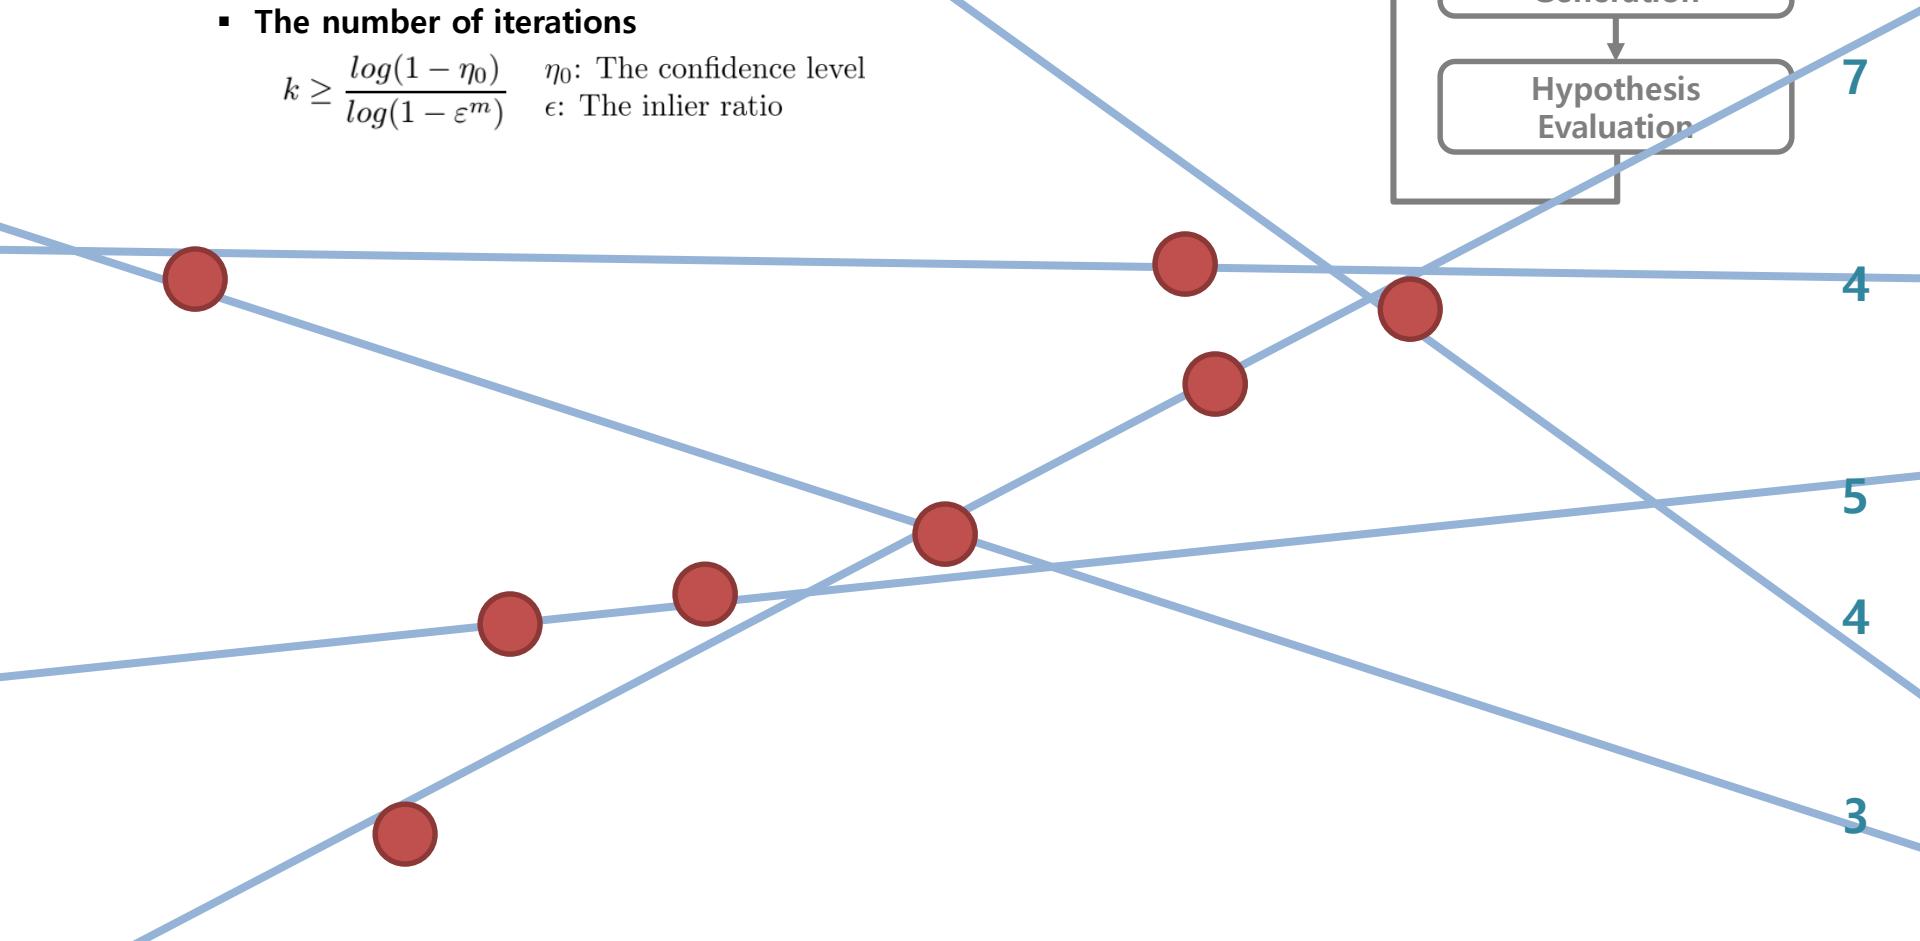
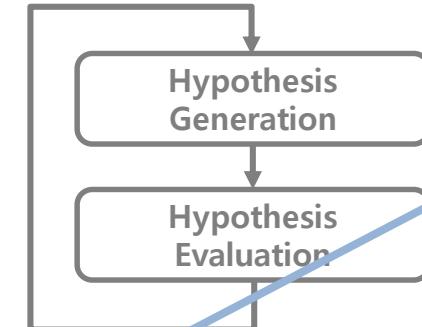


# Random Sample Consensus (RANSAC)

Parameters:

- The inlier threshold
- The number of iterations

$$k \geq \frac{\log(1 - \eta_0)}{\log(1 - \epsilon^m)} \quad \eta_0: \text{The confidence level}$$
$$\epsilon: \text{The inlier ratio}$$



# Line Fitting with RANSAC



```
1. #include "opencv_all.hpp"

2. #define CONVERT_LINE(line) (cv::Vec3d(line(0), -line(1), -line(0) * line(2) + line(1) * line(3)))

3. int main(void)
4. {
5.     int ransac_trial = 50;
6.     double ransac_thresh = 3.0;
7.     int ransac_n_sample = 2;
8.     int sim_n_data = 1000;
9.     double sim_inlier_ratio = 0.5, sim_inlier_noise = 1.0;
10.    cv::Vec3d truth(1.0 / sqrt(2.0), 1.0 / sqrt(2.0), -240.0); // The line model: a*x + b*y + c = 0

11.    // Generate data
12.    std::vector<cv::Point2d> data;
13.    cv::RNG rng;
14.    for (int i = 0; i < sim_n_data; i++)
15.    {
16.        if (rng.uniform(0.0, 1.0) < sim_inlier_ratio)
17.        {
18.            double x = rng.uniform(0.0, 480.0);
19.            double y = (truth(0) * x + truth(2)) / -truth(1);
20.            x += rng.gaussian(sim_inlier_noise);
21.            y += rng.gaussian(sim_inlier_noise);
22.            data.push_back(cv::Point2d(x, y)); // Inlier
23.        }
24.        else data.push_back(cv::Point2d(rng.uniform(0.0, 640.0), rng.uniform(0.0, 480.0))); // Outlier
25.    }

26.    // Estimate a line using RANSAC ...
27.    // Estimate a line using least-squares method (for reference) ...

28.    // Display estimates
29.    printf("* The Truth: %.3f, %.3f, %.3f\n", truth(0), truth(1), truth(2));
30.    printf("* Estimate (RANSAC): %.3f, %.3f, %.3f (Score: %d)\n", best_line(0), best_line(1), ..., best_score);
31.    printf("* Estimate (LSM): %.3f, %.3f, %.3f\n", lsm_line(0), lsm_line(1), lsm_line(2));

32.    return 0;
33.}
```

# Line Fitting with RANSAC



```
1. // Estimate a line using RANSAC
2. int best_score = -1;
3. cv::Vec3d best_line;
4. for (int i = 0; i < ransac_trial; i++)
5. {
6.     // Step 1: Hypothesis generation
7.     std::vector<cv::Point2d> sample;
8.     for (int j = 1; j < ransac_n_sample; j++)
9.     {
10.         int index = rng.uniform(0, data.size());
11.         sample.push_back(data[index]);
12.     }
13.     cv::Vec4d vvxy;
14.     cv::fitLine(sample, vvxy, CV_DIST_L2, 0, 0.01, 0.01);
15.     cv::Vec3d line = CONVERT_LINE(vvxy);

16.     // Step 2: Hypothesis evaluation
17.     int score = 0;
18.     for (size_t j = 0; j < data.size(); j++)
19.     {
20.         double error = fabs(line(0) * data[j].x + line(1) * data[j].y + line(2));
21.         if (error < ransac_thresh) score++;
22.     }

23.     if (score > best_score)
24.     {
25.         best_score = score;
26.         best_line = line;
27.     }
28. }

29. // Estimate a line using least-squares method (for reference)
30. cv::Vec4d vvxy;
31. cv::fitLine(data, vvxy, CV_DIST_L2, 0, 0.01, 0.01);
32. cv::Vec3d lsm_line = CONVERT_LINE(vvxy);
```

## Line Fitting Result

```
* The Truth: 0.707, 0.707, -240.000
* Estimate (RANSAC): 0.712, 0.702, -242.170 (Score: 434)
* Estimate (LSM): 0.748, 0.664, -314.997
```

# Summary: Overview

- **What is 3D Vision?**
- **Single-view Geometry**
  - **Camera Projection Model**
    - A Pinhole Camera Model (**Simple 2D-3D Geometry**)
    - Homogenous Coordinates
    - Geometric Distortion and Rectification
  - **General 2D-3D Geometry**
    - Camera Calibration
    - Absolute Camera Pose Estimation (PnP Problem)
- **Two-view Geometry**
  - **Planar 2D-2D Geometry (Projective Geometry)**
    - Planar Homography
  - **General 2D-2D Geometry (Epipolar Geometry)**
    - Fundamental/Essential Matrix
    - Relative Camera Pose Estimation
    - Relative Point Localization (Triangulation)
- **Multi-view Geometry**
  - Bundle Adjustment (Non-linear Optimization)
  - Applications: Structure-from-motion, Visual SLAM, and Visual Odometry
- **Correspondence Problem**
  - Random Sample Consensus (RANSAC)

Slides and example codes are available:  
[https://github.com/sunglok/3dv\\_tutorial](https://github.com/sunglok/3dv_tutorial)

$$\therefore \mathbf{x} = \mathbf{P}\mathbf{X} \quad (\mathbf{P} = \mathbf{K}[\mathbf{R} | \mathbf{t}])$$
$$\mathbf{x} = \mathbf{K}\hat{\mathbf{x}}$$

$$\therefore \mathbf{x}' = \mathbf{H}\mathbf{x}$$
$$\hat{\mathbf{x}}' = \hat{\mathbf{H}}\hat{\mathbf{x}} \quad (\hat{\mathbf{H}} = \mathbf{R} + \frac{1}{d}\mathbf{t}\mathbf{n}^\top)$$
$$\therefore \mathbf{x}'^\top \mathbf{F} \mathbf{x} = 0$$
$$\hat{\mathbf{x}}'^\top \mathbf{E} \hat{\mathbf{x}} = 0 \quad (\mathbf{E} = [\mathbf{t}]_\times \mathbf{R})$$

$$\sum_i^n \sum_j^m \left\| \mathbf{x}_i^j - \mathbf{P}_j \mathbf{X}_i \right\|_{\Sigma}^2$$

# Summary: Examples

Slides and example codes are available:  
[https://github.com/sunglok/3dv\\_tutorial](https://github.com/sunglok/3dv_tutorial)

- Single-view Geometry
  - Camera Projection Model
    - Simple Camera Calibration and Object Localization
    - Image Formation: `image_formation.cpp`
    - Geometric Distortion Correction: `distortion_correction.cpp`
  - General 2D-3D Geometry
    - Camera Calibration: `camera_calibration.cpp`
    - Pose Estimation and Its Application to Augmented Reality: `pose_estimation.cpp`
- Two-view Geometry
  - Planar 2D-2D Geometry (Projective Geometry)
    - Perspective Distortion Correction: `perspective_correction.cpp`
    - Planar Image Stitching: `image_stitching.cpp`
    - 2D Video Stabilization: `video_stabilization.cpp`
  - General 2D-2D Geometry (Epipolar Geometry)
    - Monocular Visual Odometry (Epipolar Version): `visual_odometry_epipolar.cpp`
    - Triangulation (Two-view Reconstruction): `triangulation.cpp`
- Multi-view Geometry
  - Bundle Adjustment: `bundle_adjustment.cpp`
  - Sparse and Dense 3D Reconstruction using VisualSfM
- Correspondence Problem
  - Line Fitting with RANSAC: `ransac_line.cpp`